

AGILE: Lightweight and Efficient Asynchronous GPU-SSD Integration

Zhuoping Yang*, Jinming Zhuang*, Xingzhen Chen*, Alex K. Jones†, and Peipei Zhou*

Brown University*; Syracuse University†;

zhuoping_yang@brown.edu
peipei_zhou@brown.edu

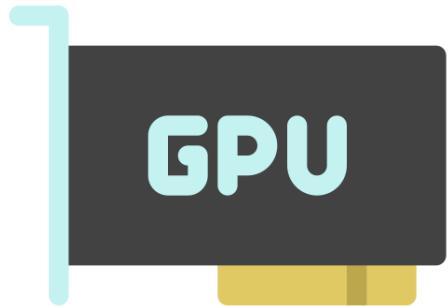
AGILE is open source!



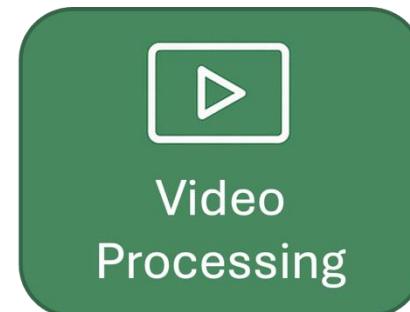
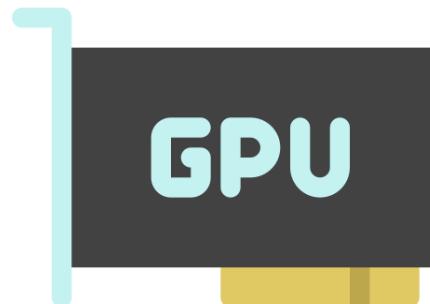
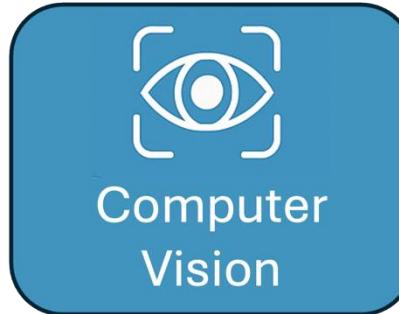
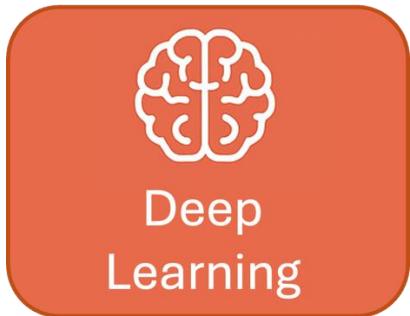
<https://peipeizhou-eecs.github.io/>

<https://github.com/arc-research-lab/AGILE>

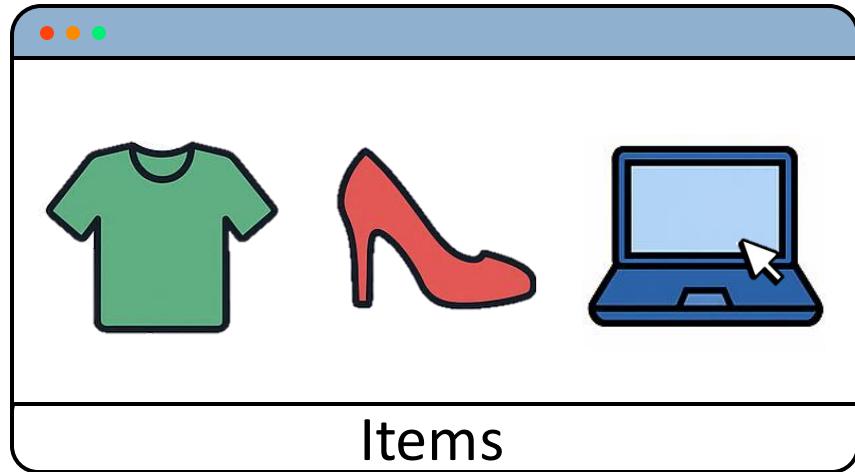


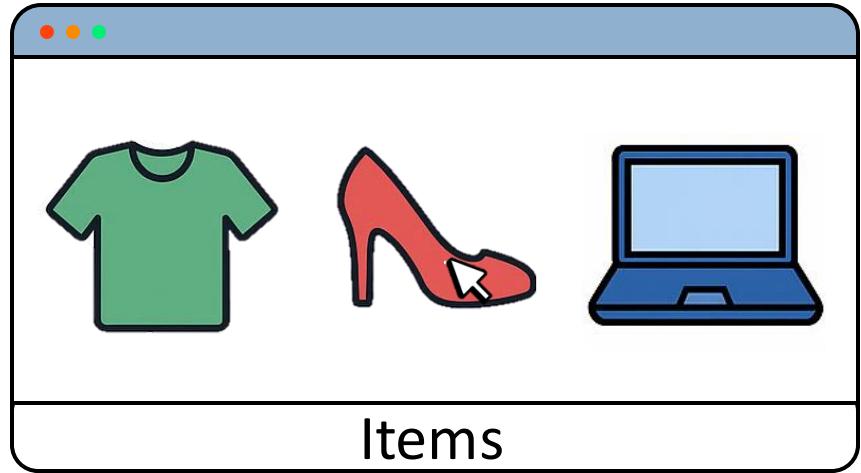


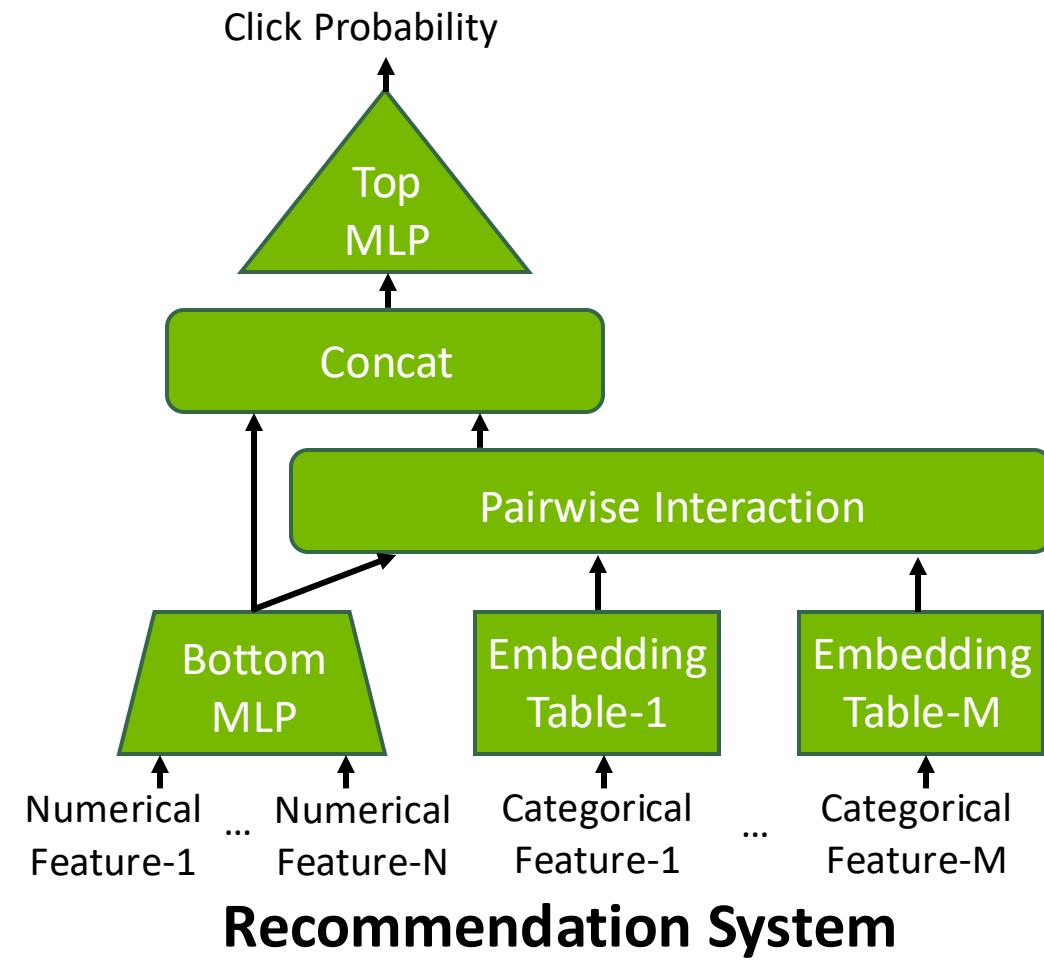
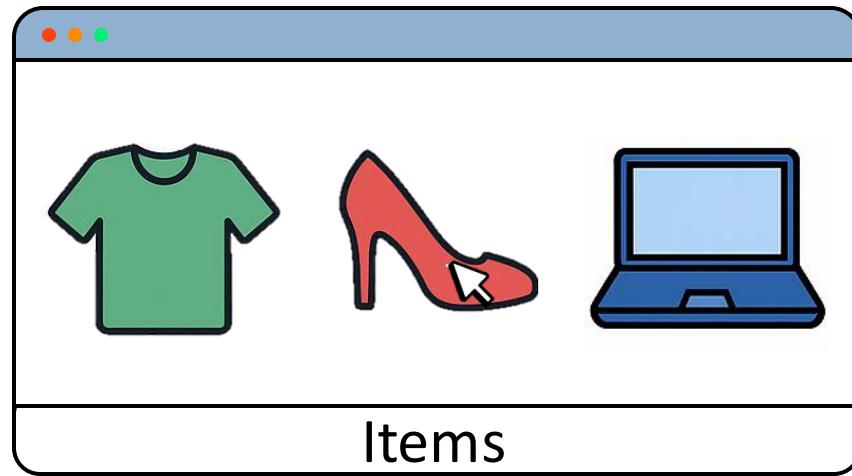
GPGPU: The engine behind modern applications.

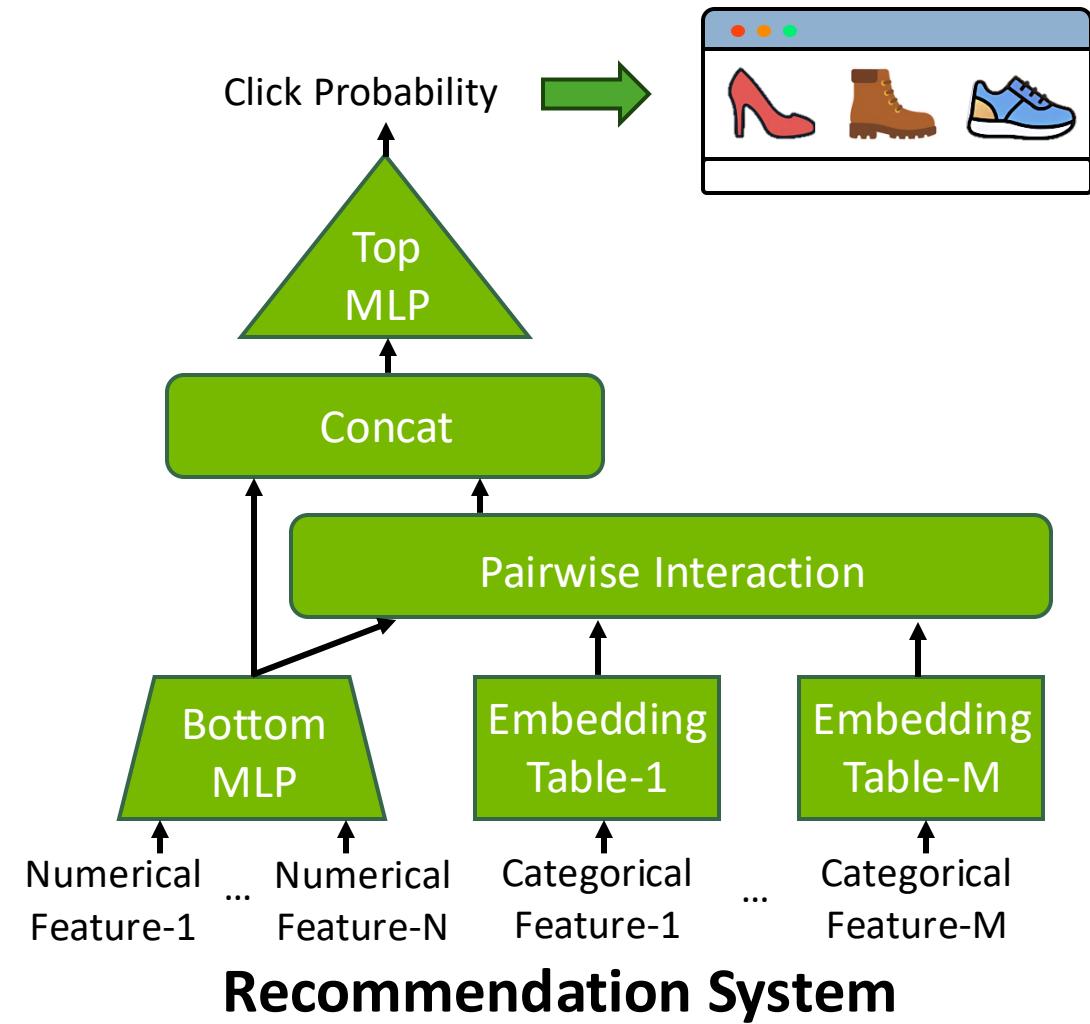
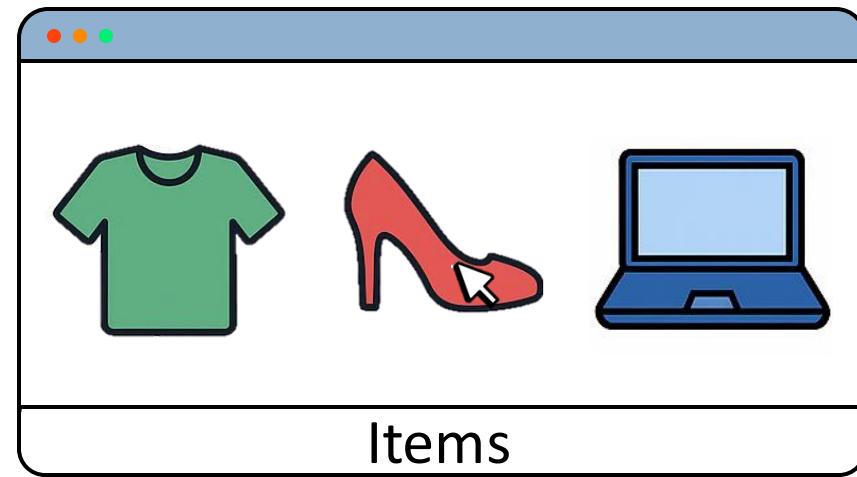


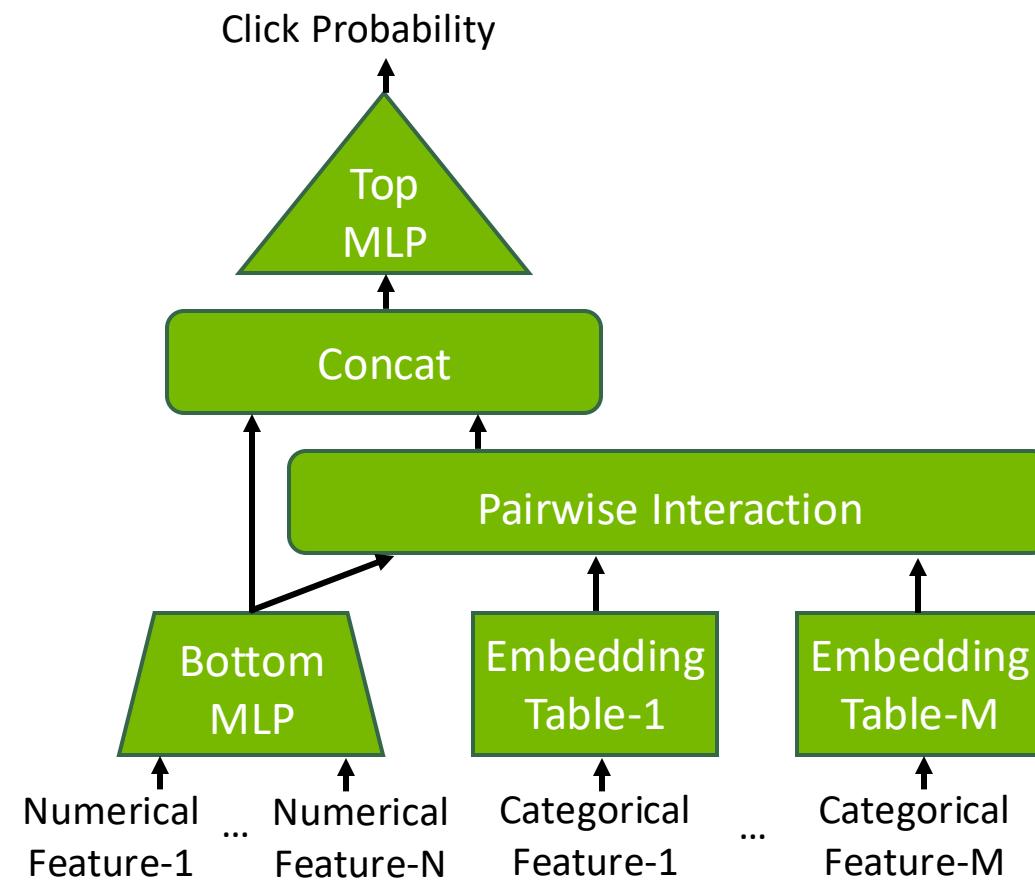
GPGPU: The engine behind modern applications.



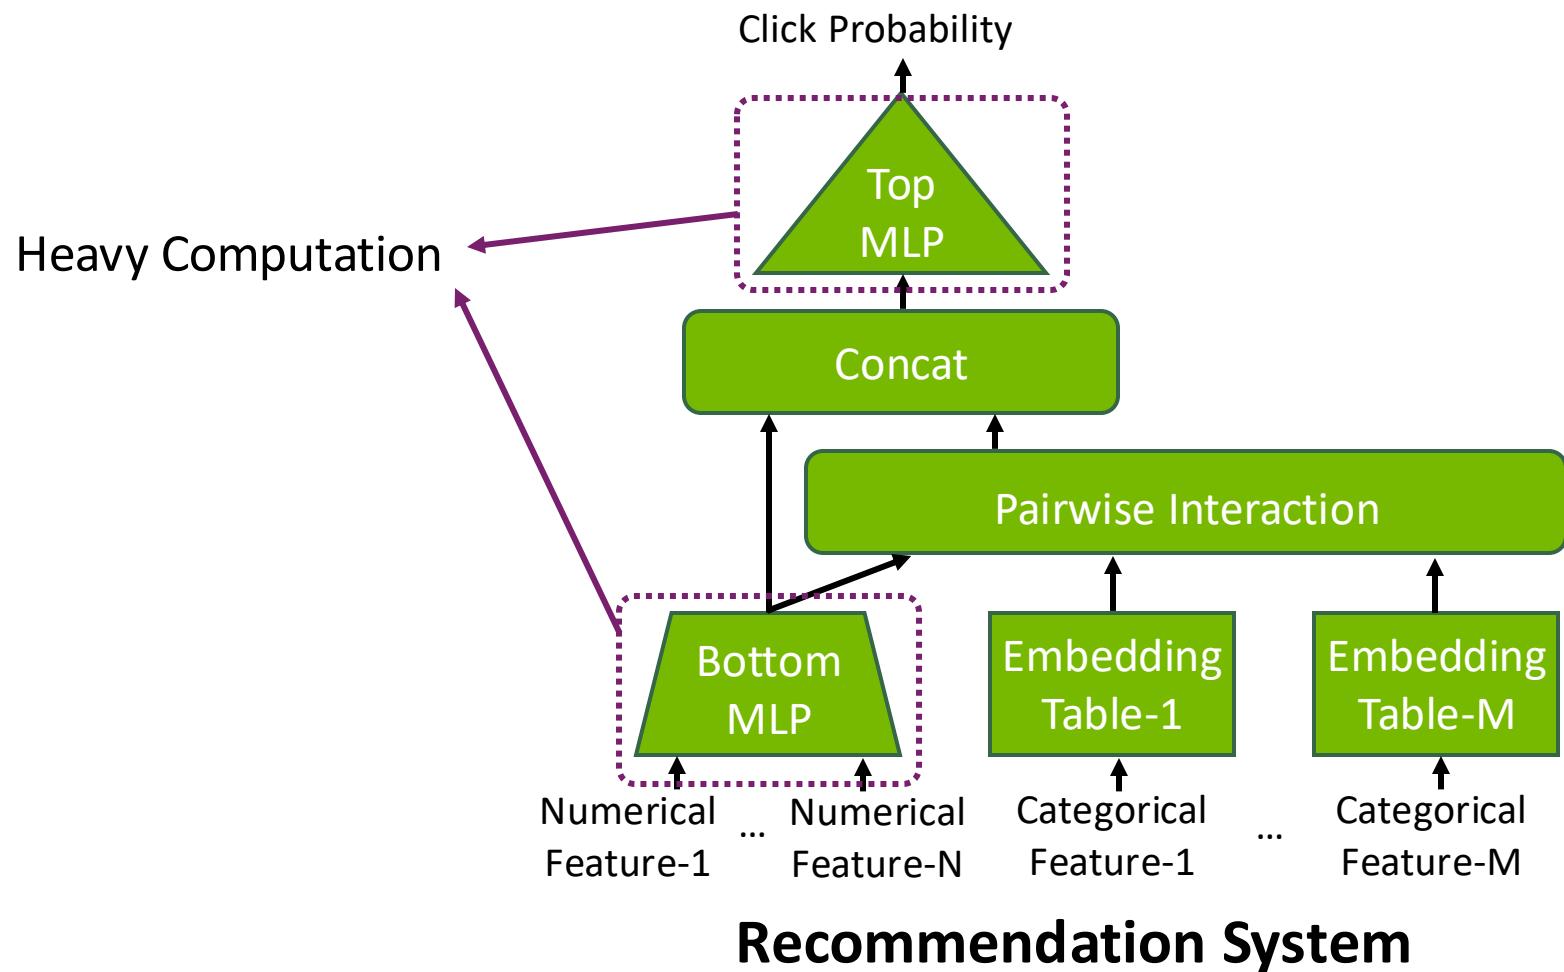


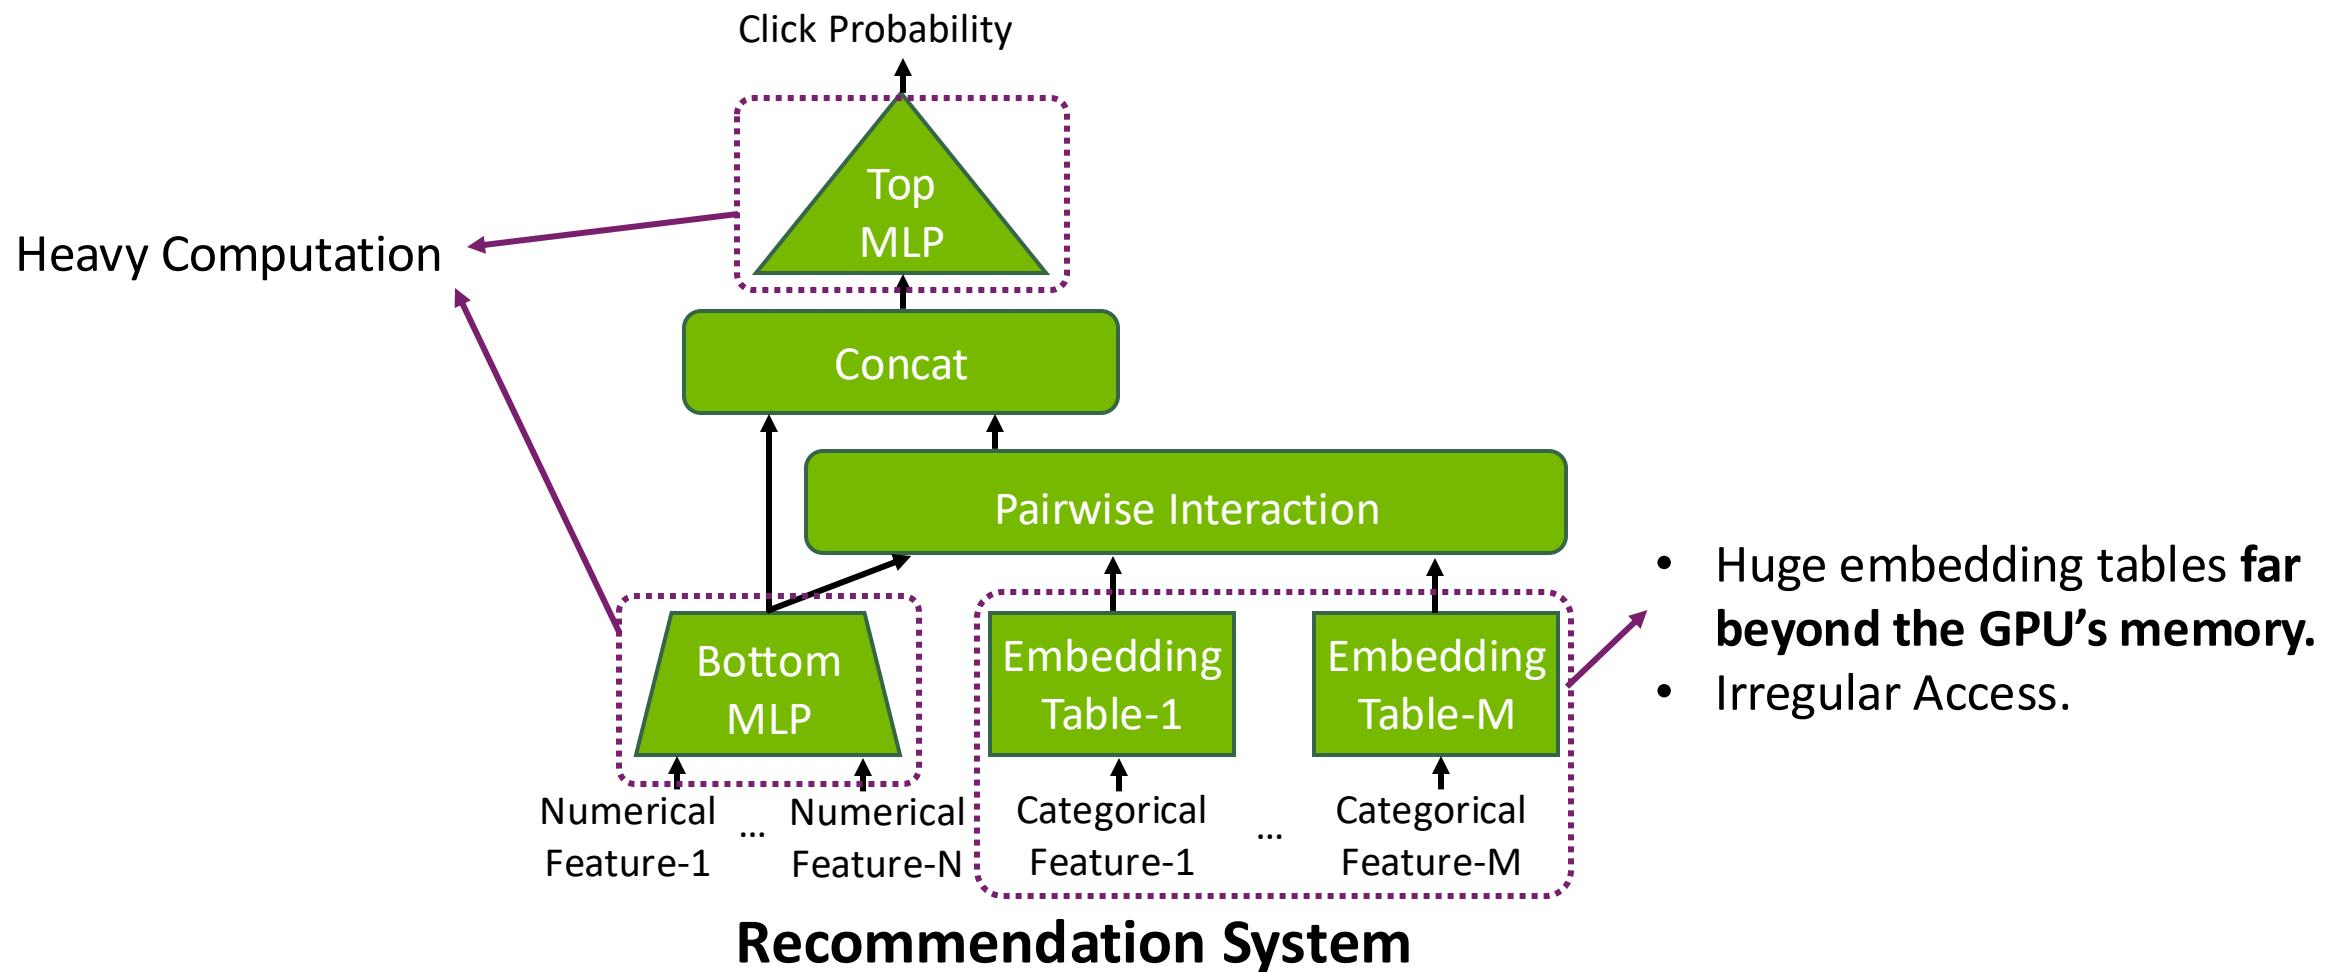


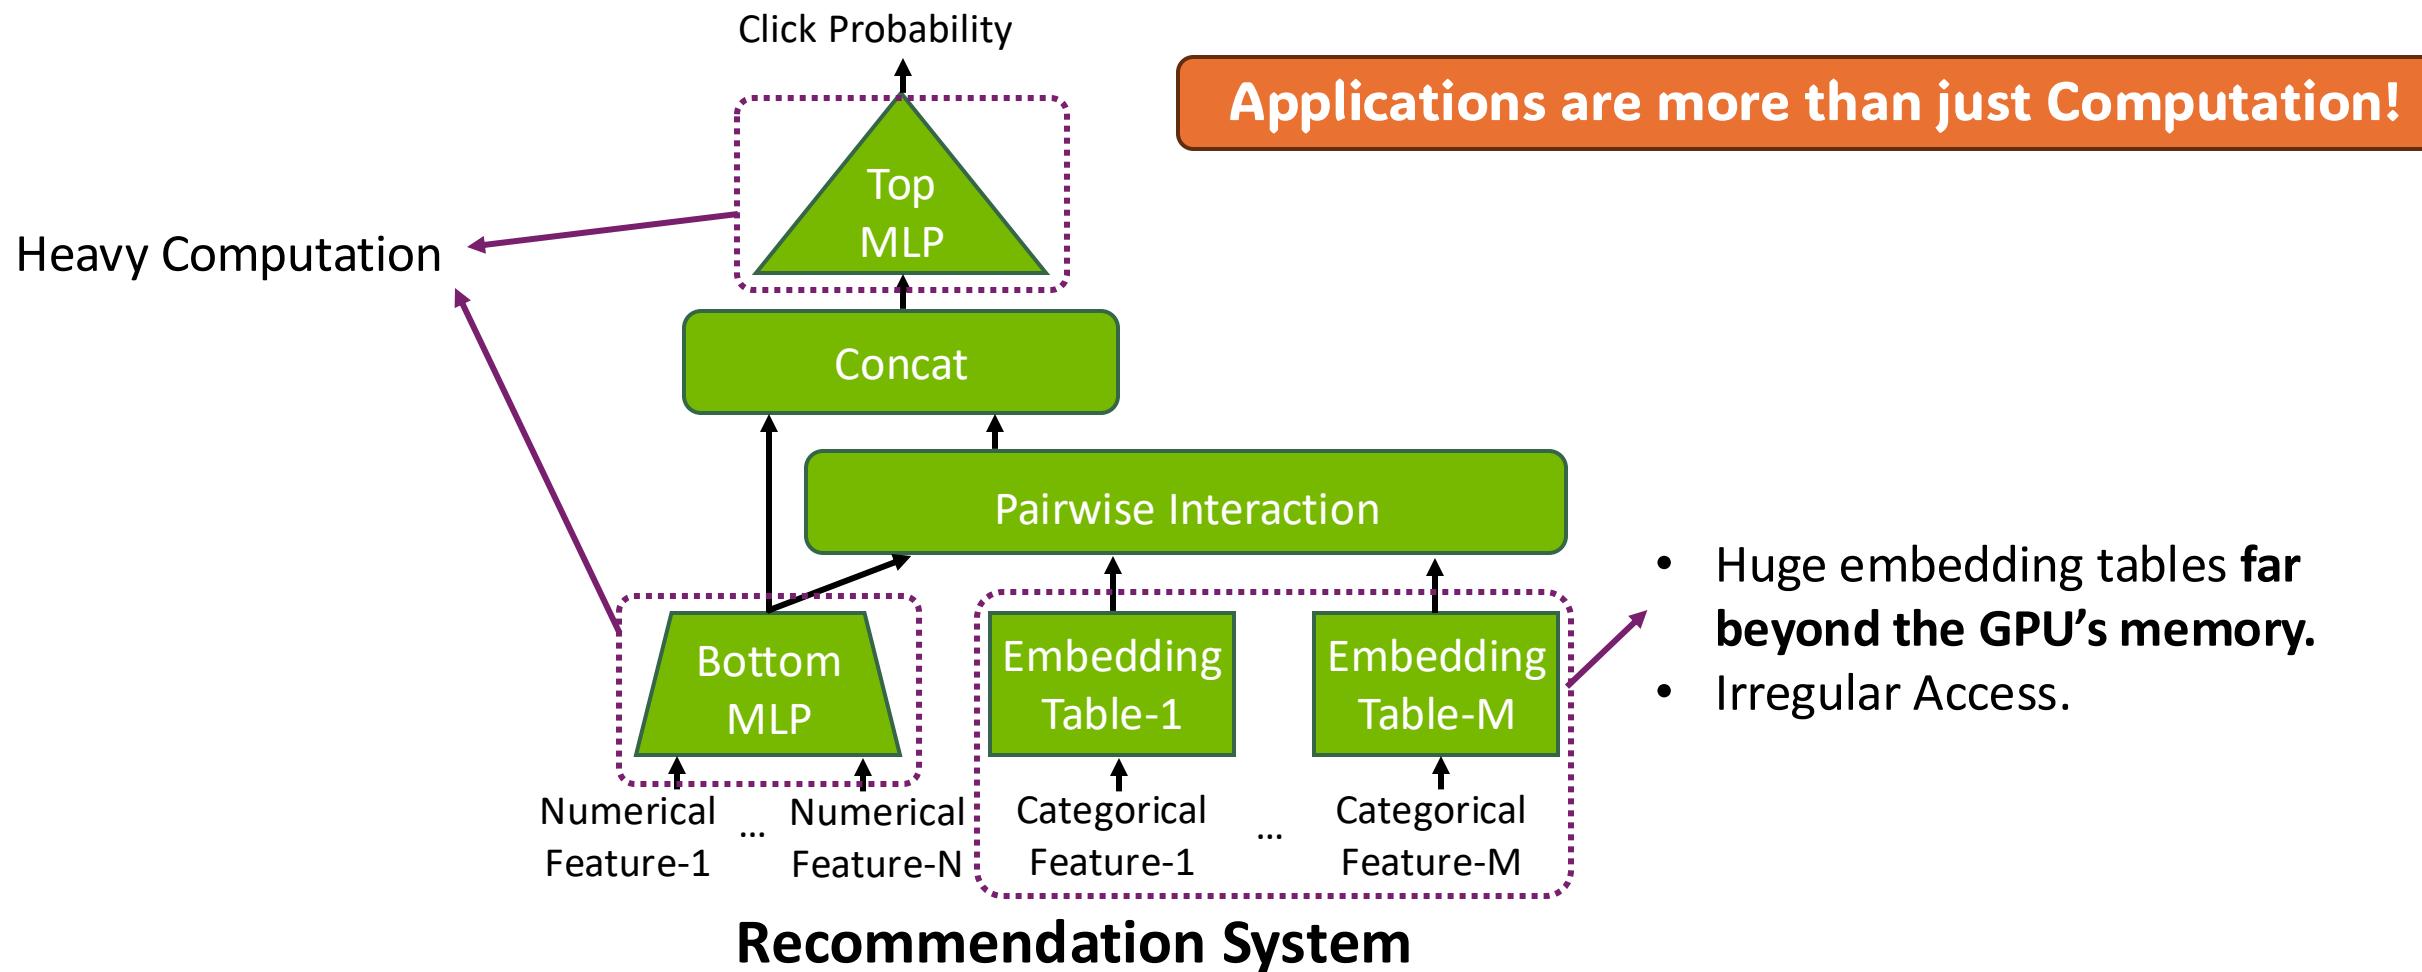


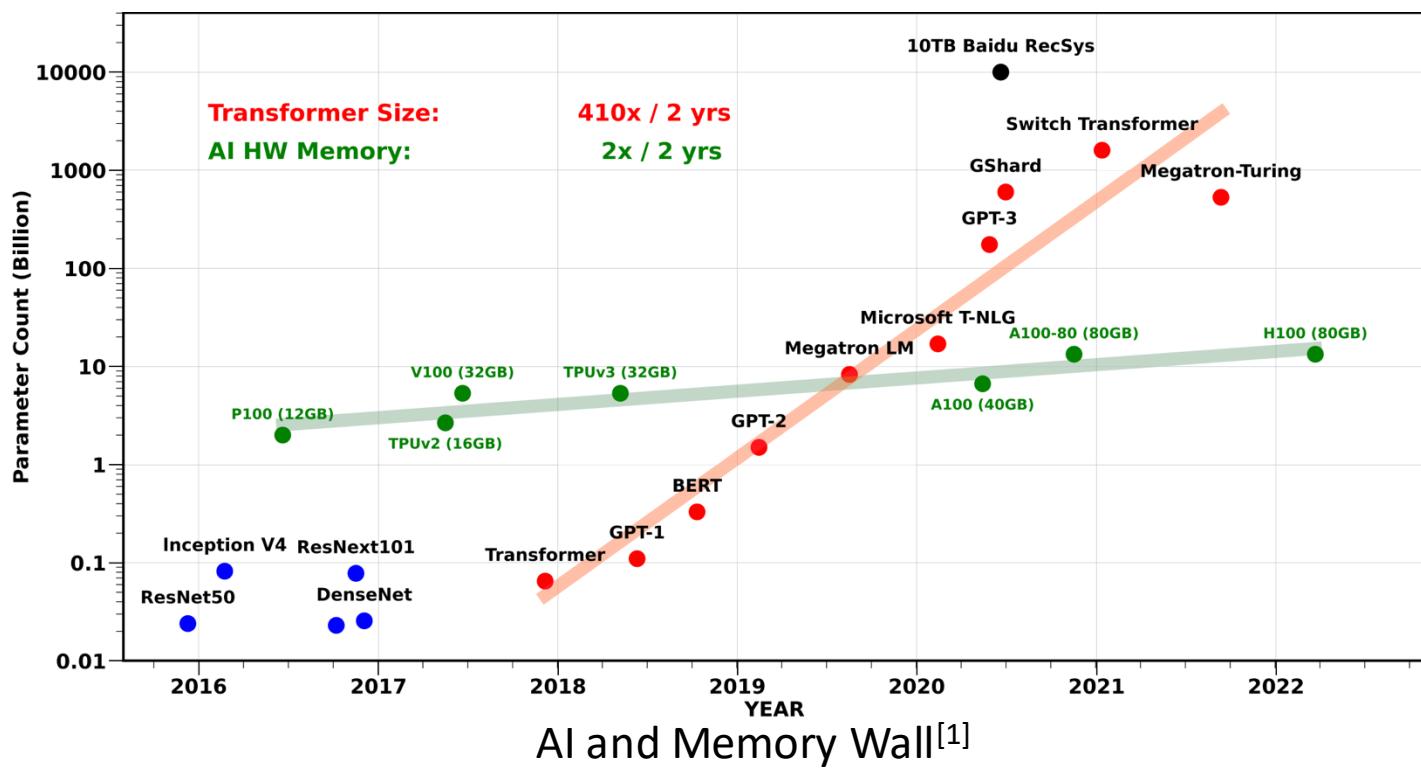


Recommendation System



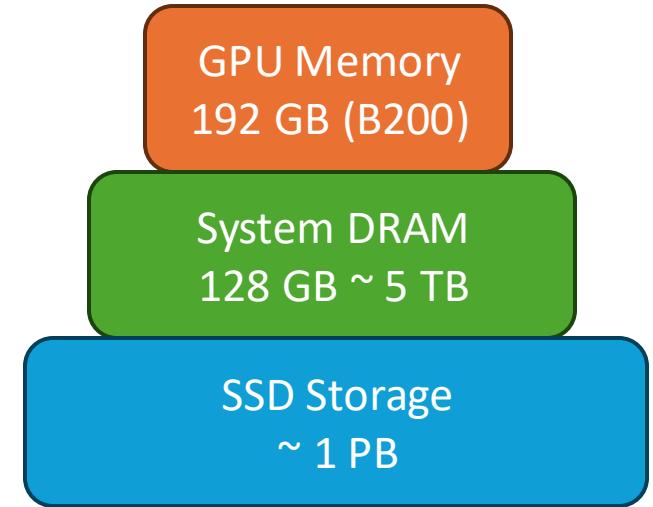
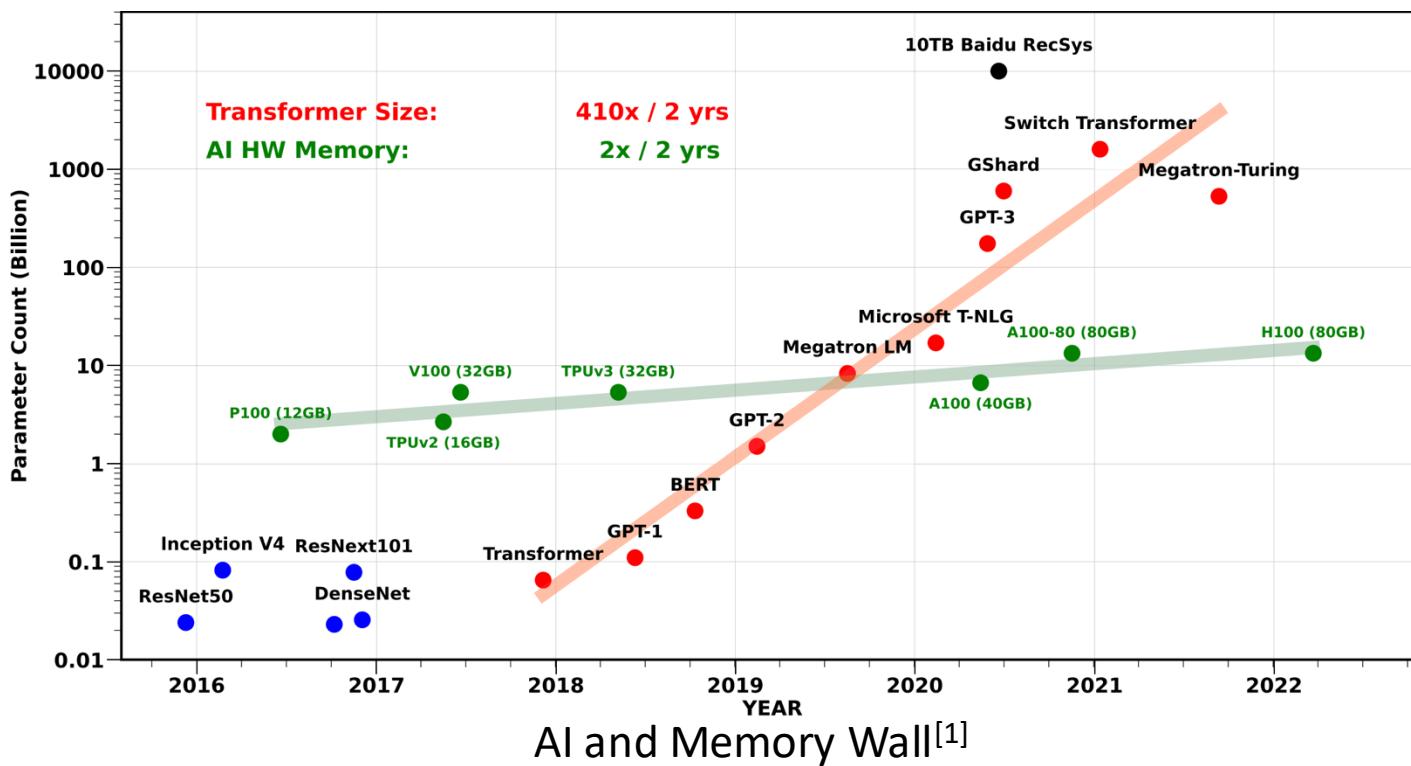






GPUs are increasingly constrained by the memory wall

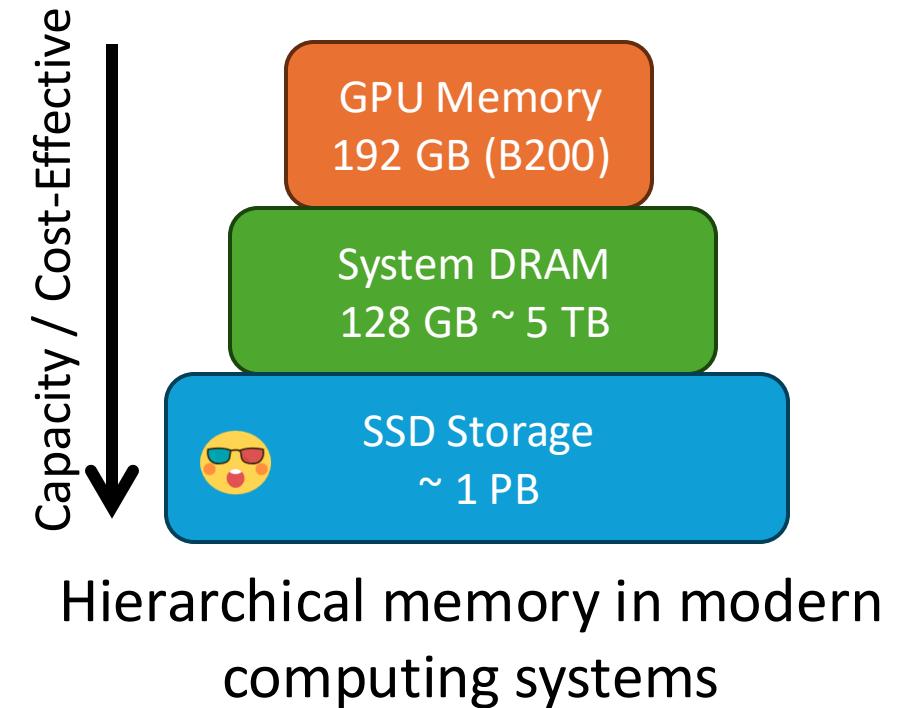
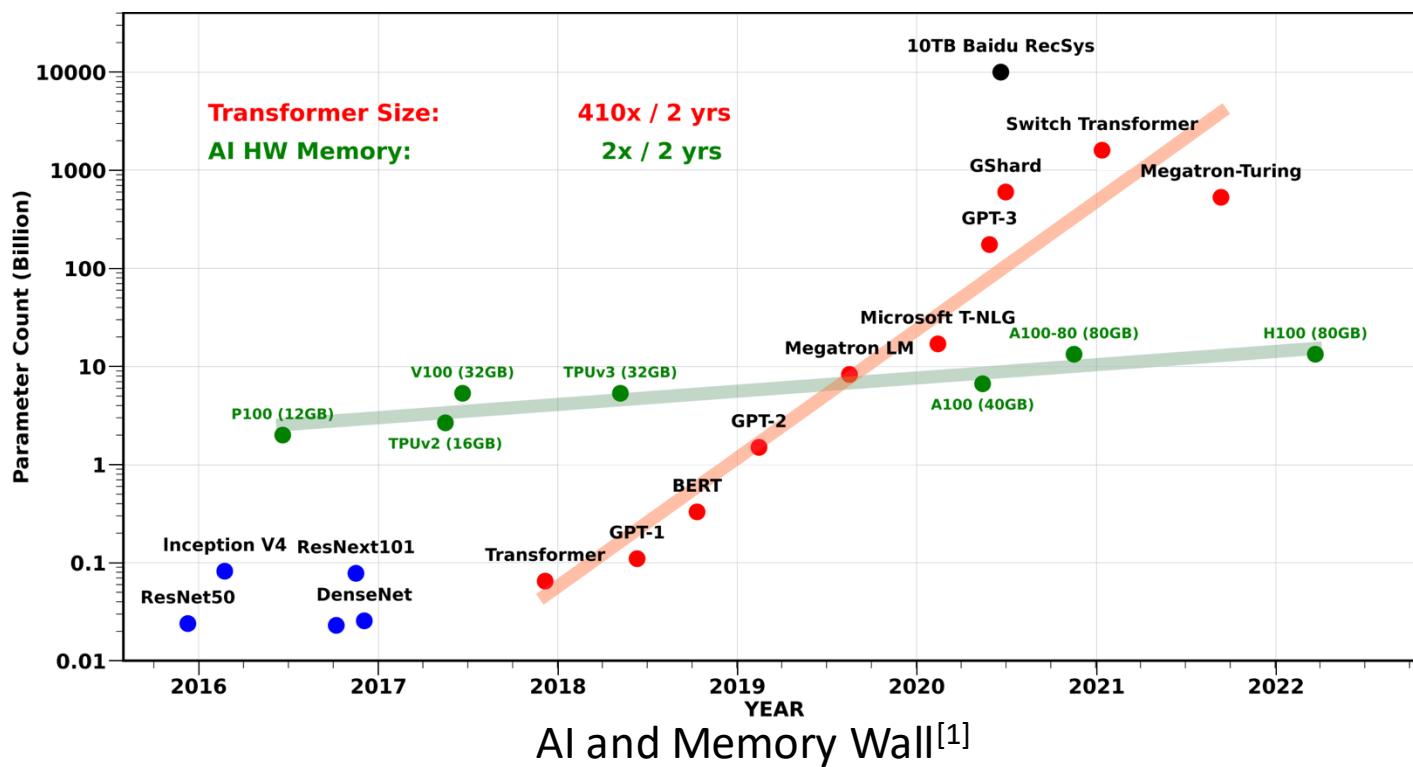
[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.



Hierarchical memory in modern computing systems

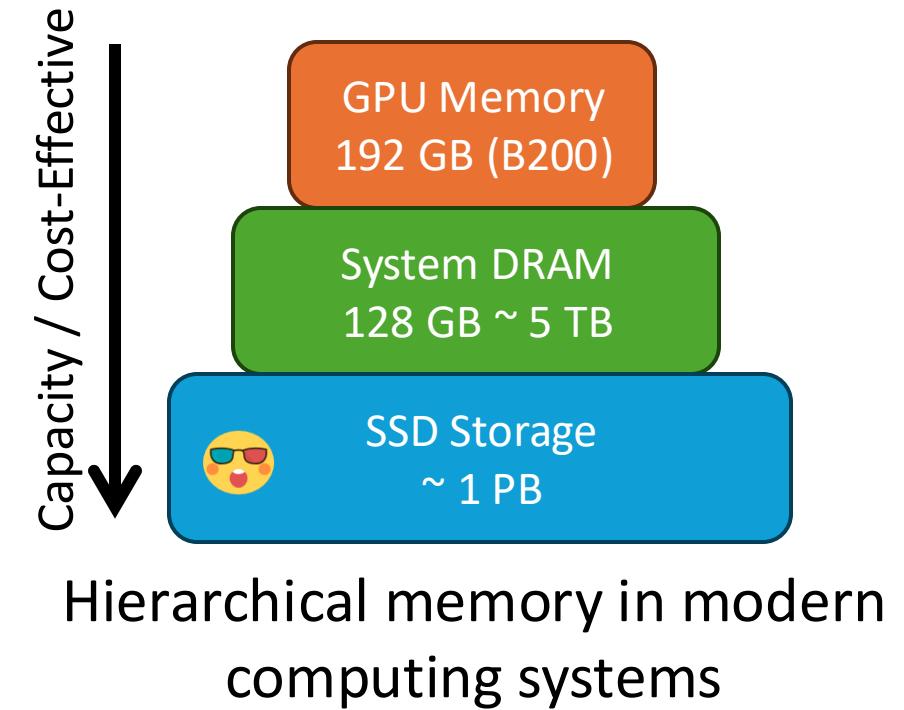
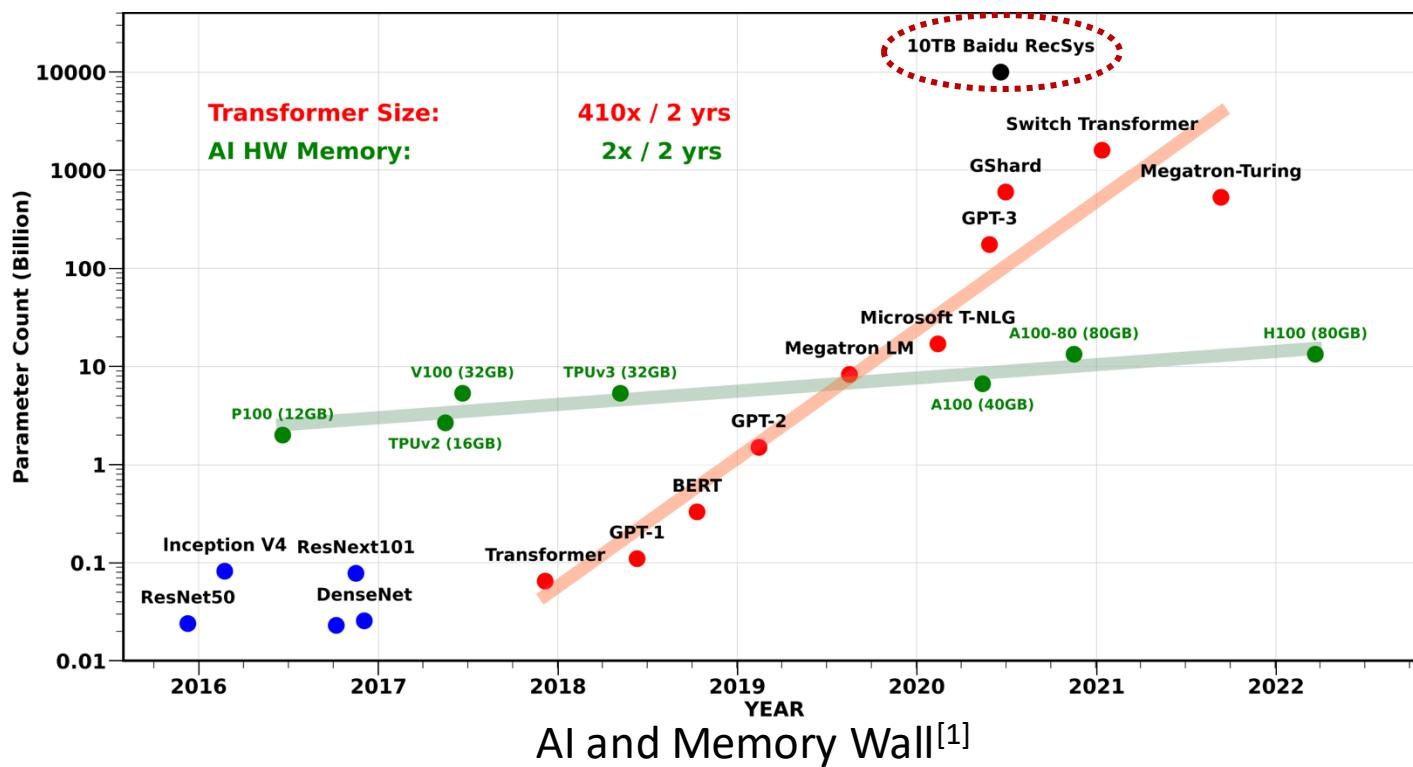
GPUs are increasingly constrained by the memory wall

[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.



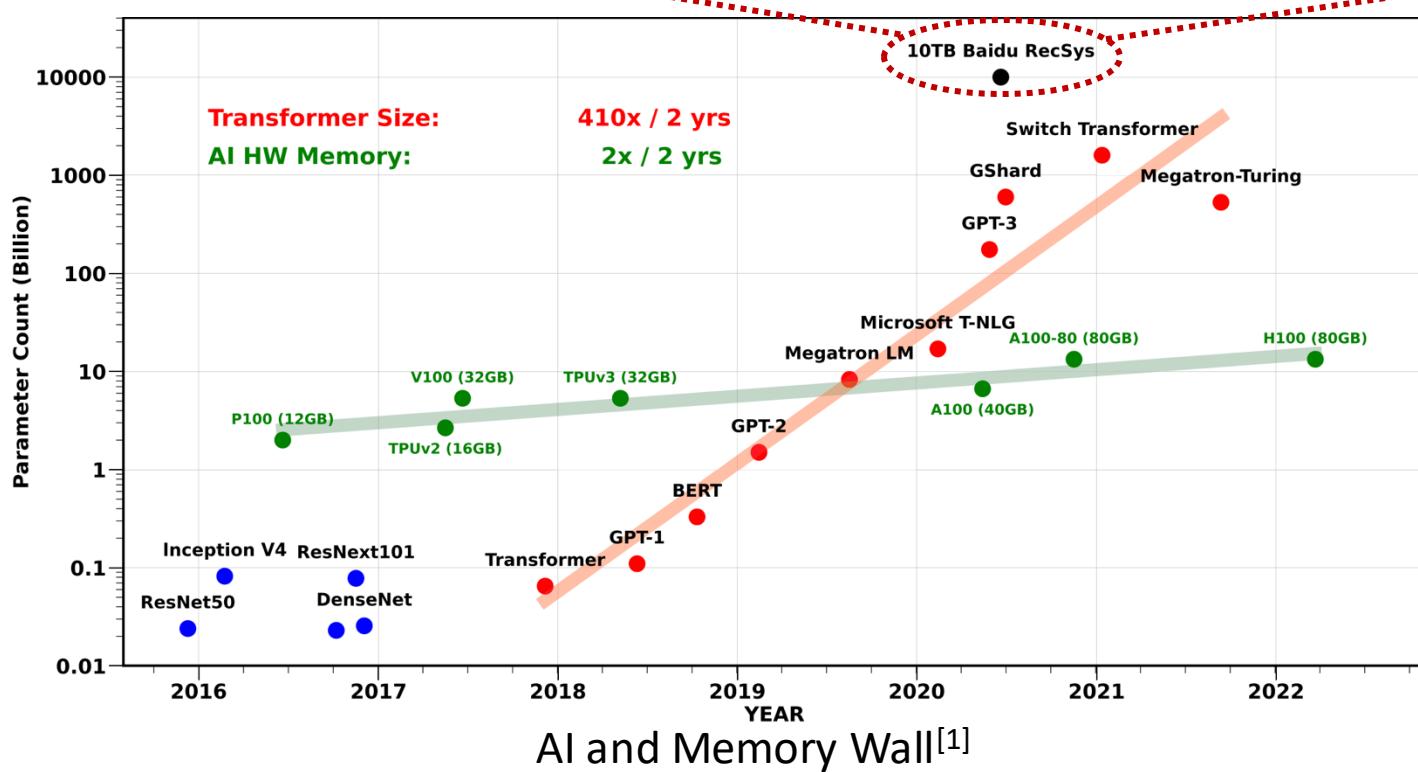
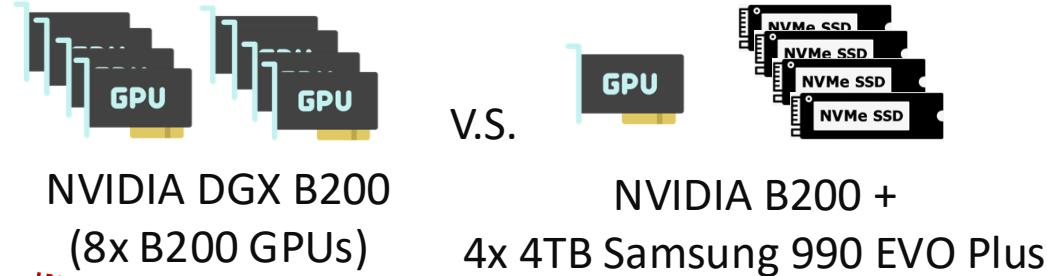
GPUs are increasingly constrained by the memory wall

[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.

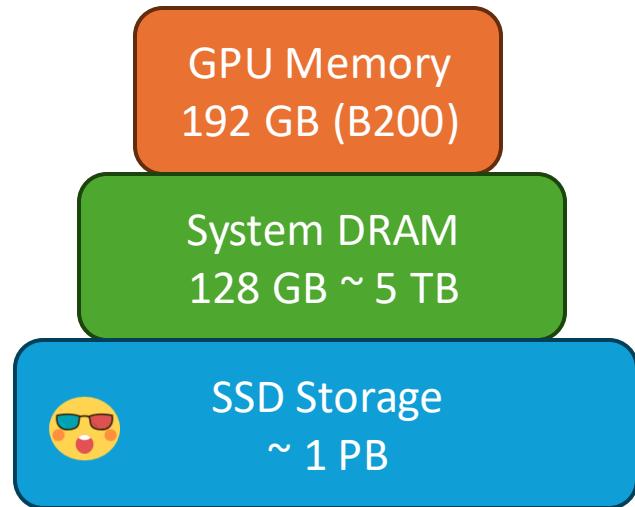


GPUs are increasingly constrained by the memory wall

[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.



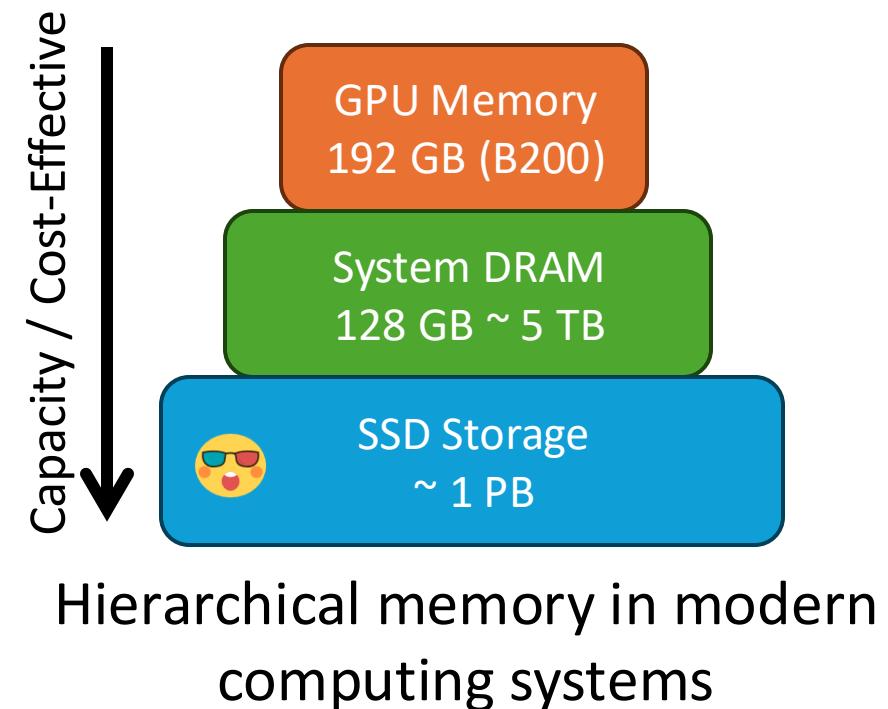
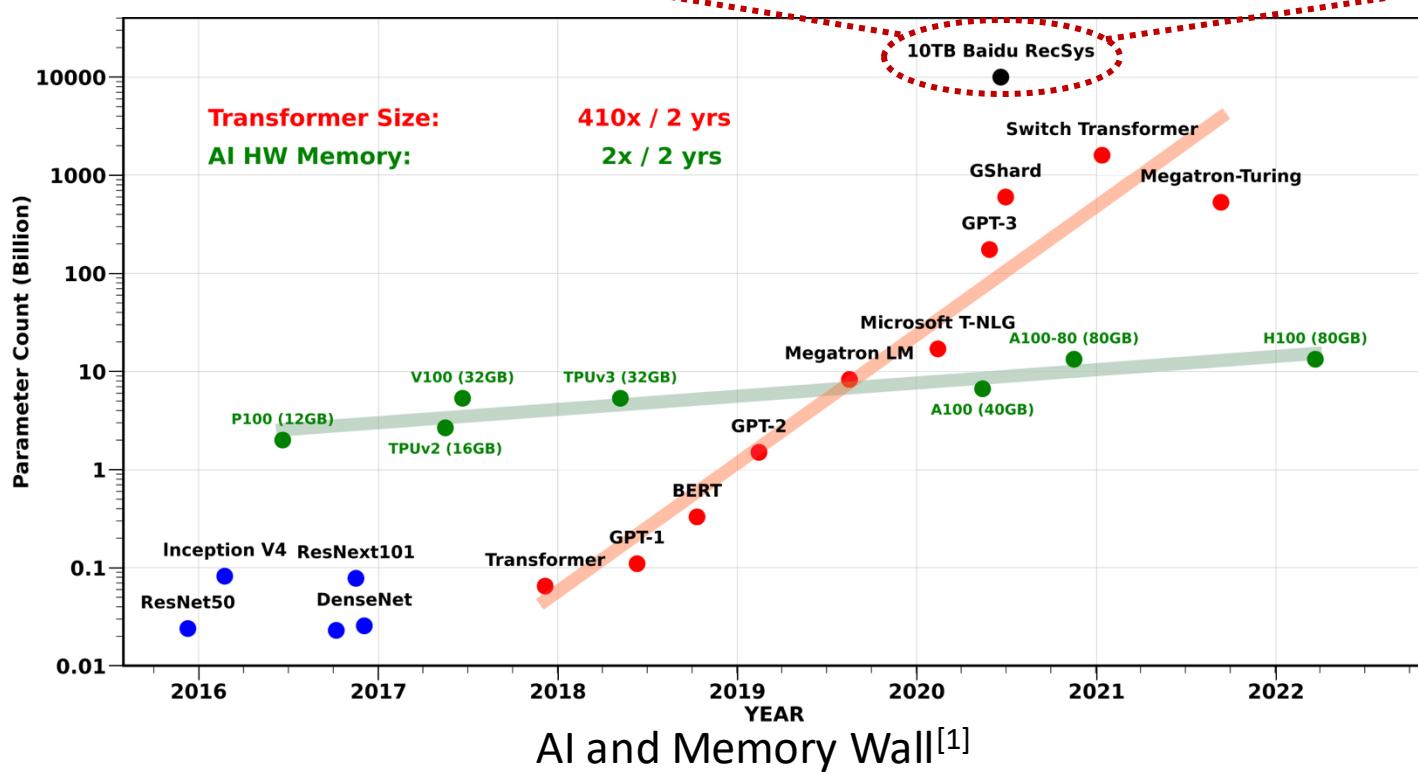
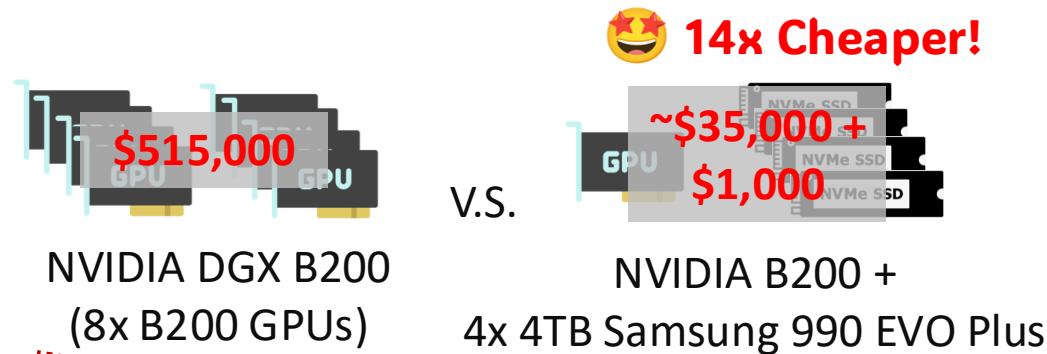
Capacity / Cost-Effective



Hierarchical memory in modern computing systems

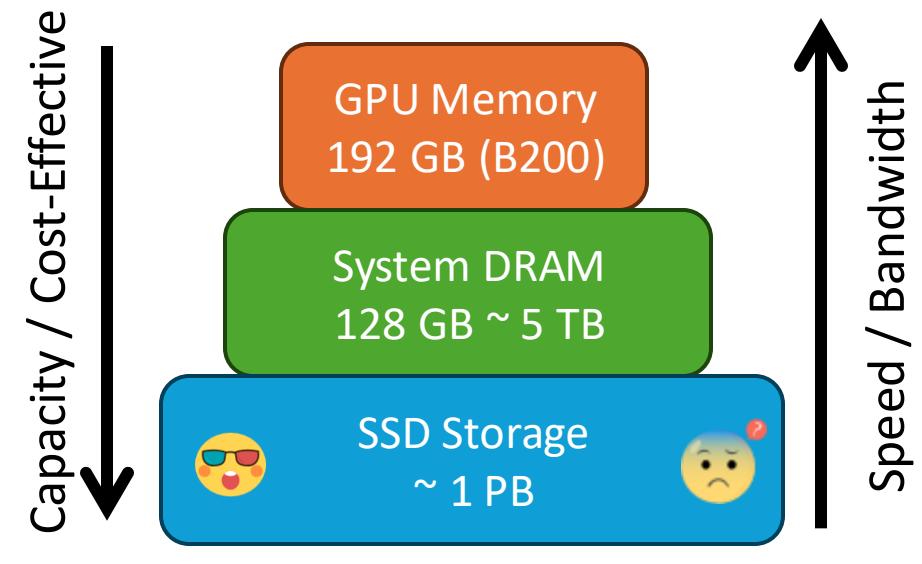
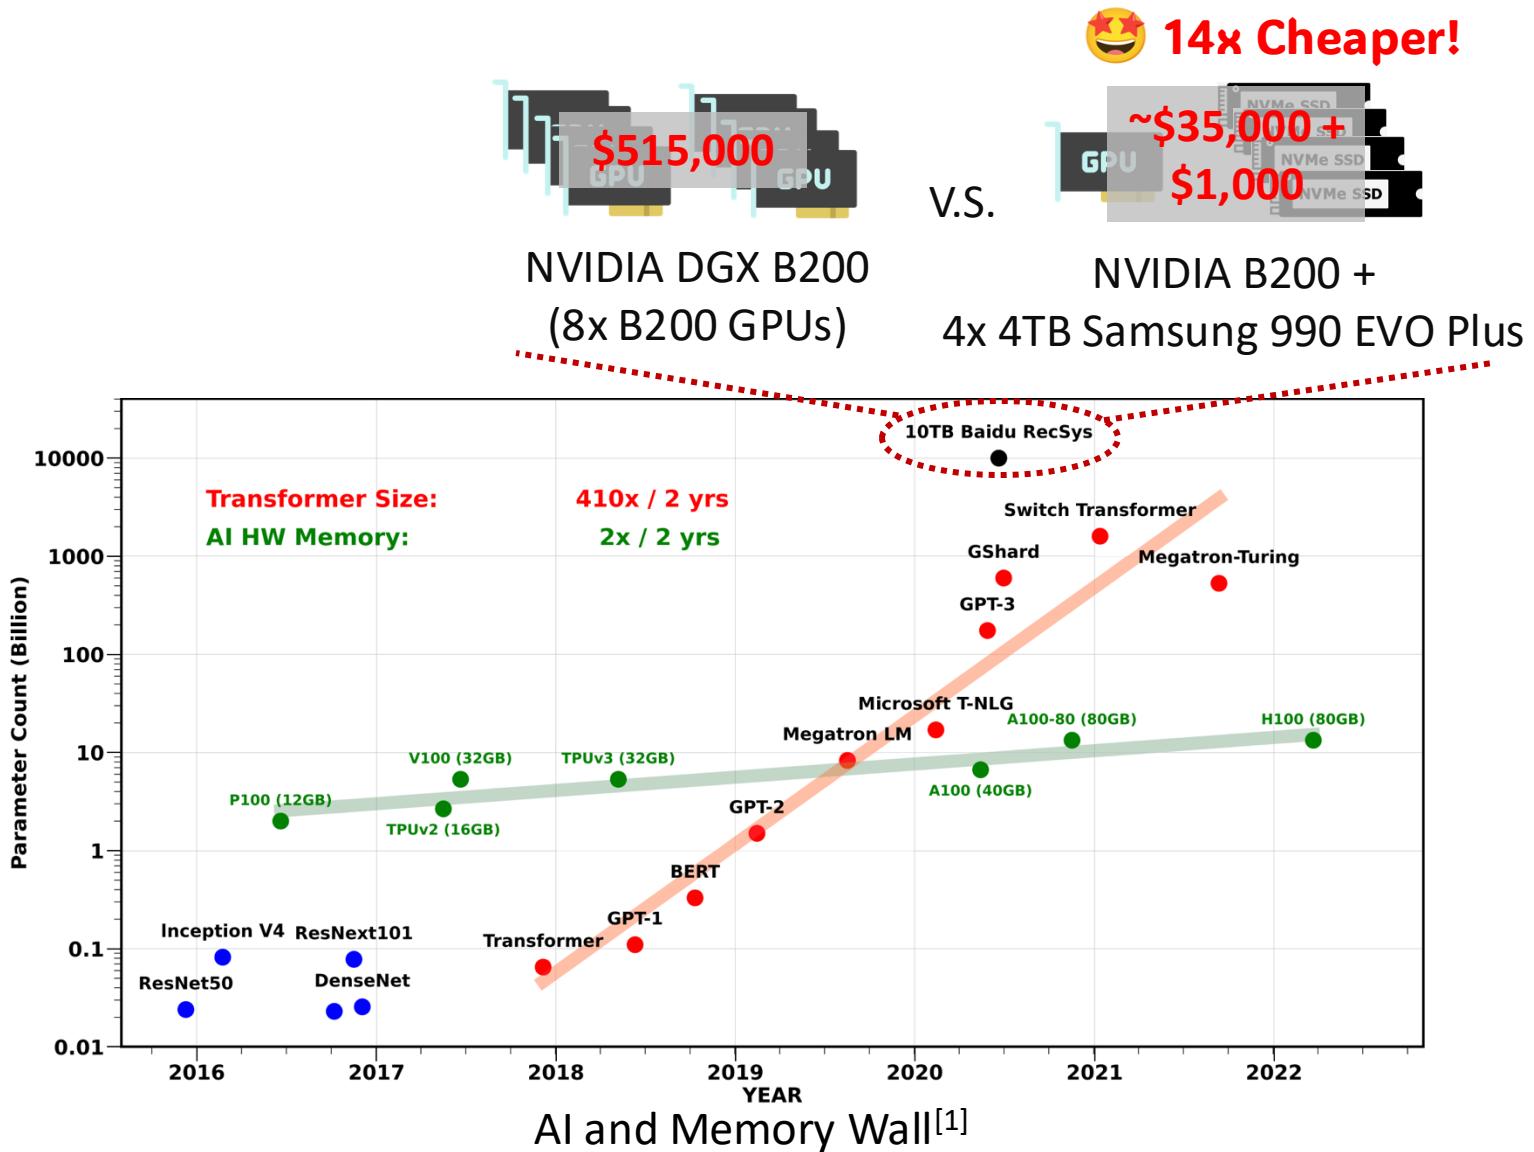
GPUs are increasingly constrained by the memory wall

[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.



GPUs are increasingly constrained by the memory wall

[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.

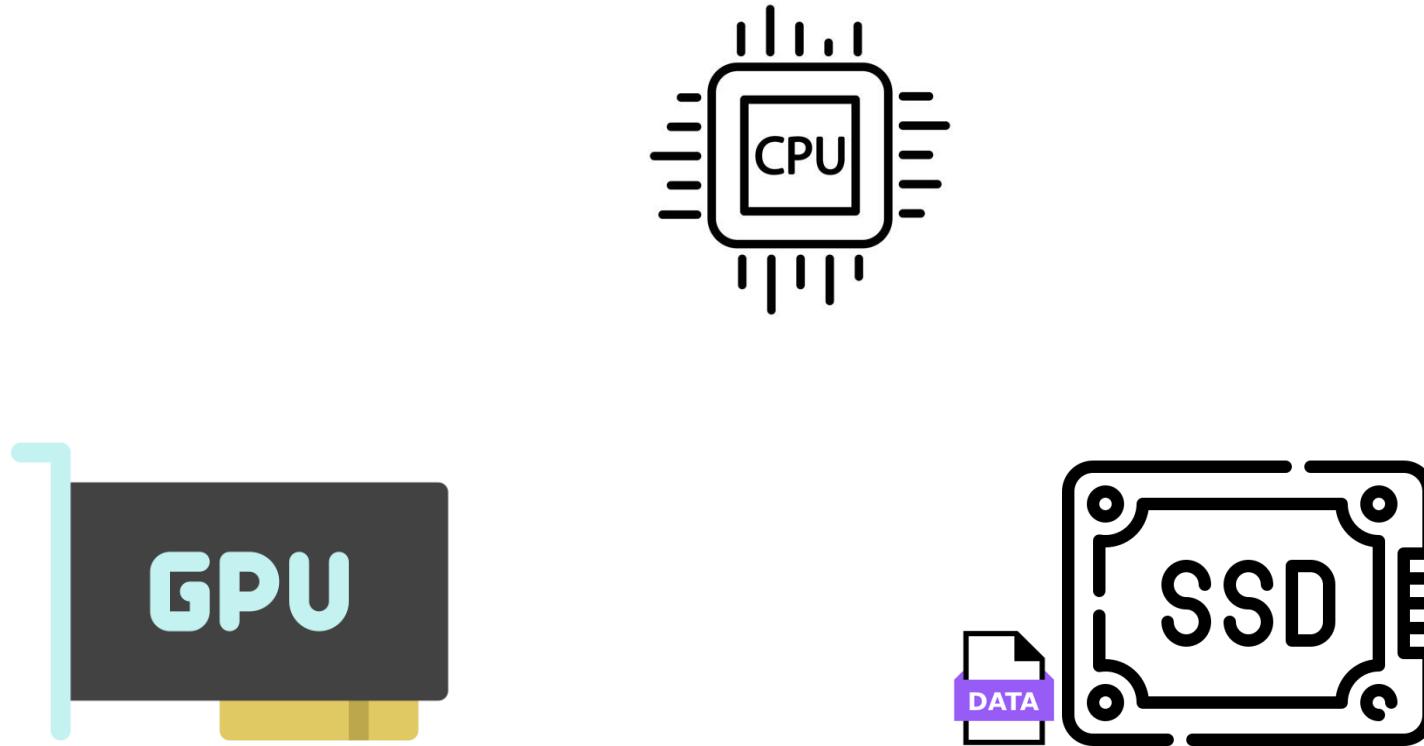


Hierarchical memory in modern computing systems

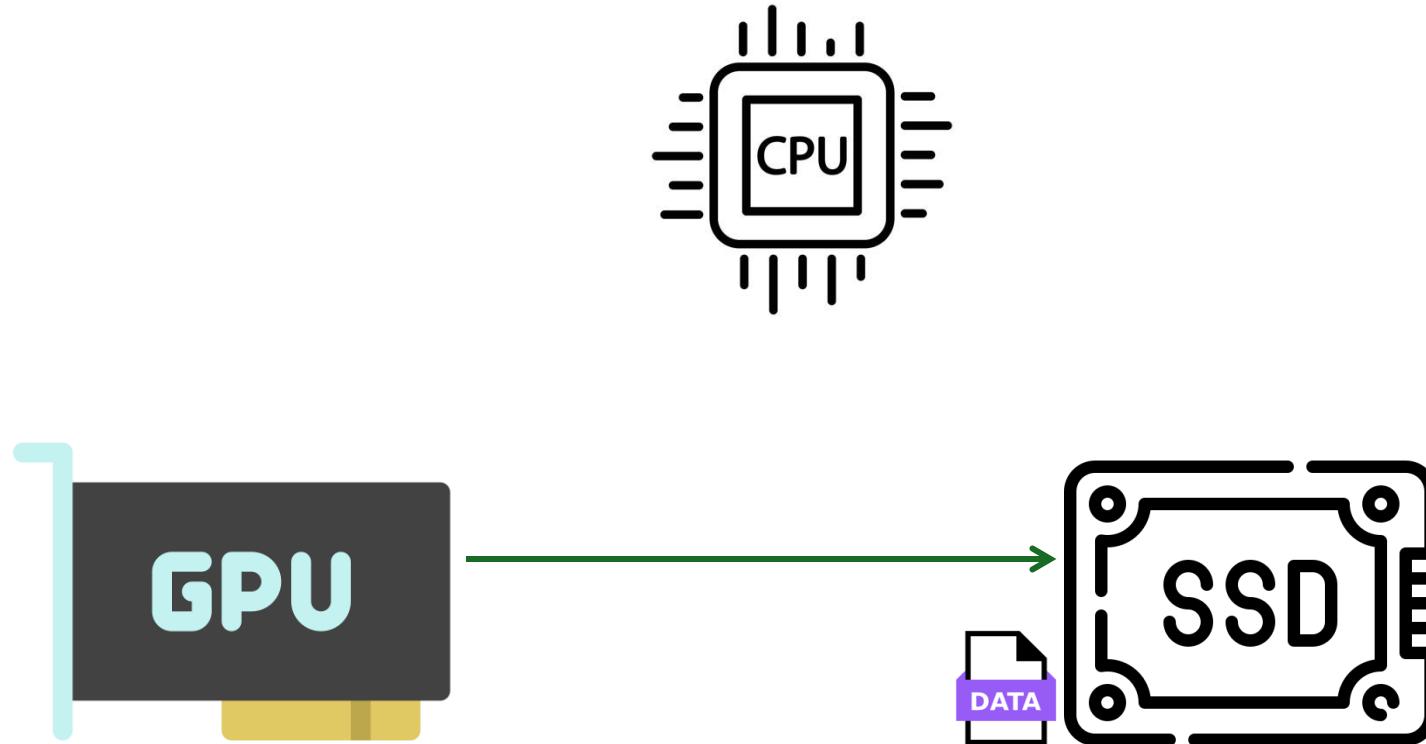
GPUs are increasingly constrained by the memory wall

[1] Gholami, Amir, et al. "Ai and memory wall." *IEEE Micro* 44.3 (2024): 33-39.

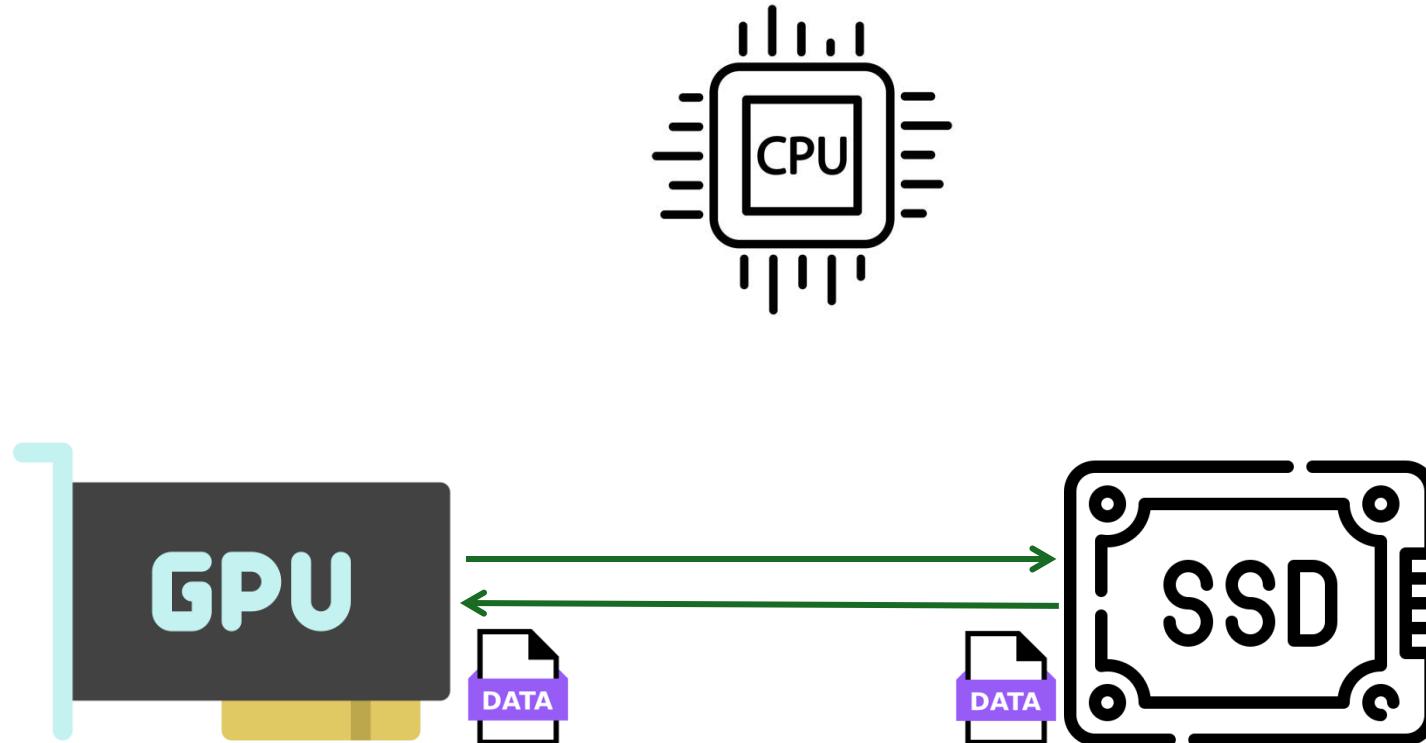
BaM^[1] (A GPU-Centric Storage Access Model)



BaM^[1] (A GPU-Centric Storage Access Model)



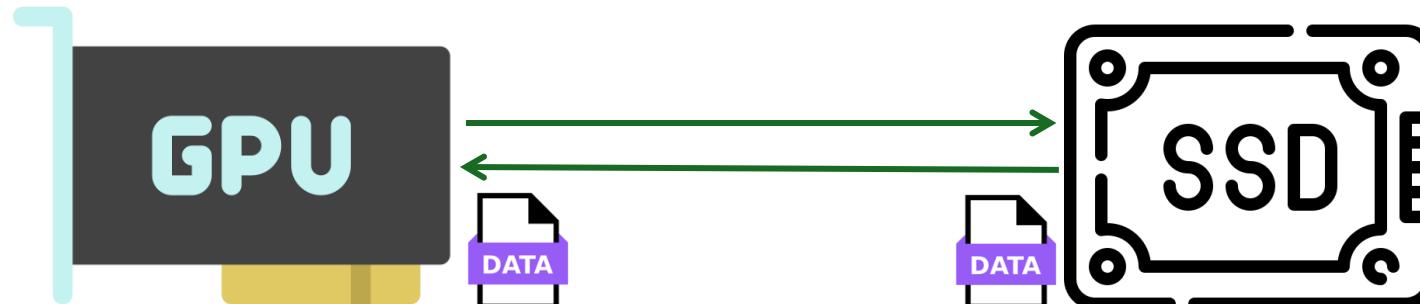
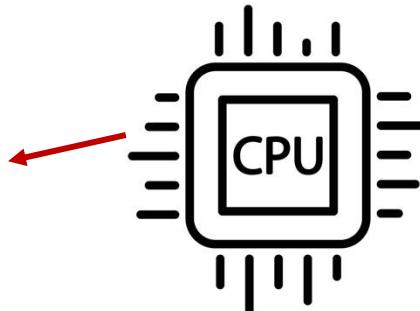
BaM^[1] (A GPU-Centric Storage Access Model)



BaM^[1] (A GPU-Centric Storage Access Model)

CPU is bypassed completely in both the control plane and the data plane.

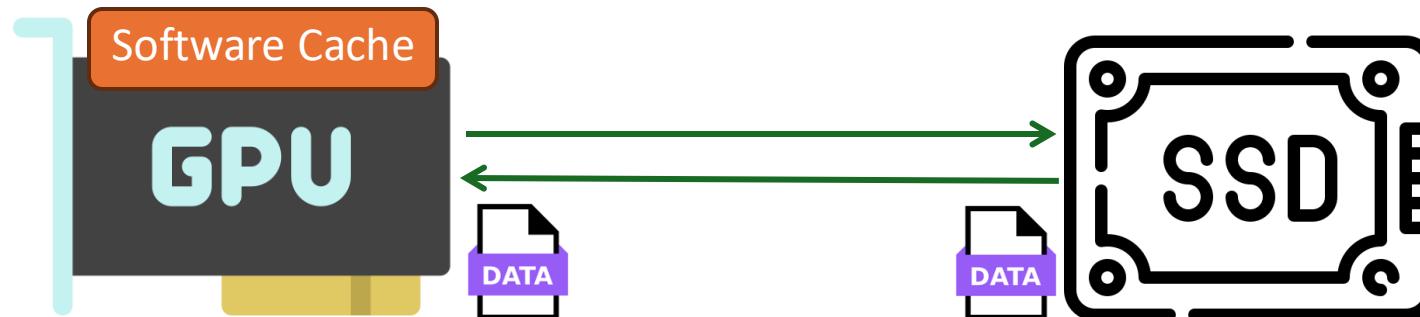
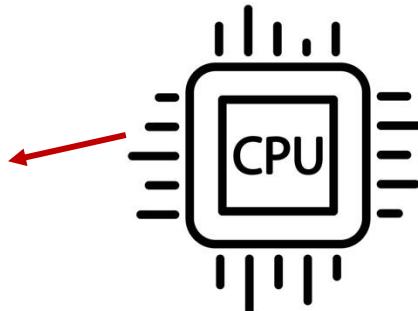
- No extra copies or synchronization overheads



BaM^[1] (A GPU-Centric Storage Access Model)

CPU is bypassed completely in both the control plane and the data plane.

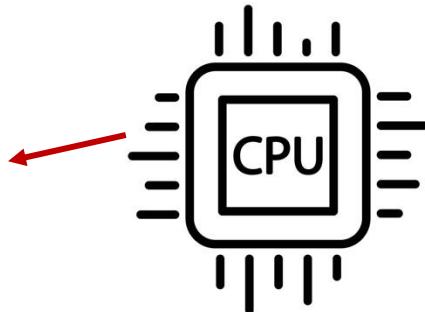
- No extra copies or synchronization overheads



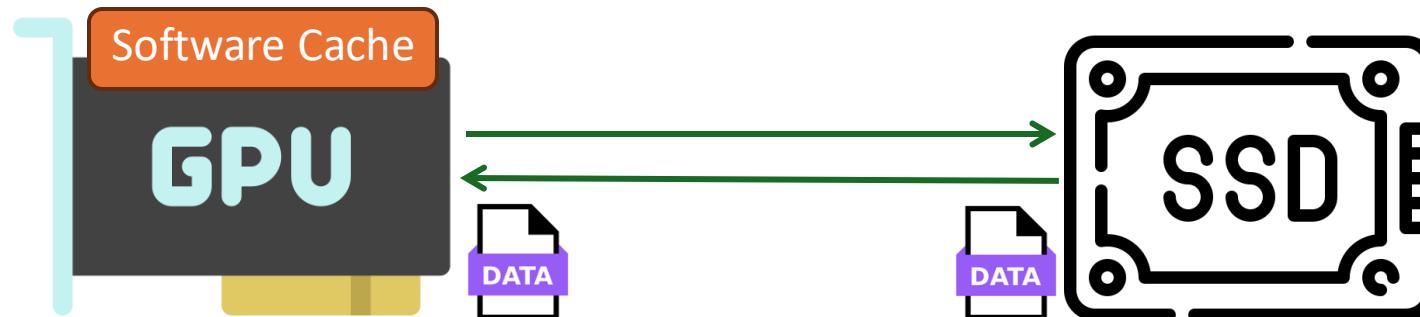
BaM^[1] (A GPU-Centric Storage Access Model)

CPU is bypassed completely in both the control plane and the data plane.

- No extra copies or synchronization overheads



Can we further improve the GPU-centric storage access model?



[1] Zaid Qureshi, Vikram Sharma Mailthody, et al. "GPU-initiated on-demand high-throughput storage access in the BaM system architecture."

Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS). 2023.

Programming Models for GPU-Centric I/O

```
1 def kernel_block_access
2     a = read_block...
3     b = read_block...
4     compute(a, b)
5     ...
```

Synchronous I/O Model
(BaM)

Programming Models for GPU-Centric I/O

```
1 def kernel_block_access
2     a = read_block(..)
3     b = read_block(..)
4     compute(a, b)
5     ...
```

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

Synchronous I/O Model
(BaM)

Programming Models for GPU-Centric I/O

```
1 def kernel_block_access
2     a = read_block(..)
3     b = read_block(..)
4     compute(a, b)
5     ...
```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```
1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

Asynchronous I/O Model

Programming Models for GPU-Centric I/O

```

1 def kernel_block_access
2     a = read_block(..)
3     b = read_block(..)
4     compute(a, b)
5 ...

```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```

1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4 ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8 ...

```

Asynchronous I/O Model

- Return quickly after the request is issued.

Programming Models for GPU-Centric I/O

```

1 def kernel_block_access
2     a = read_block(..)
3     b = read_block(..)
4     compute(a, b)
5 ...

```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```

1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8 ...

```

Asynchronous I/O Model

- Return quickly after the request is issued.
- Switch to other tasks during the slow transfer.

Programming Models for GPU-Centric I/O

```

1 def kernel_block_access
2     a = read_block(..)
3     b = read_block(..)
4     compute(a, b)
5 ...

```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```

1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8 ...

```

Asynchronous I/O Model

- Return quickly after the request is issued.
- Switch to other tasks during the slow transfer.
- Continue computation once the requested data is ready.

Programming Models for GPU-Centric I/O

Async I/O enables overlapping the slow transfer with other tasks.

```

1 def kernel_block_access
2   a = read_block(..)
3   b = read_block(..)
4   compute(a, b)
5 ...

```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```

1 def kernel_async_access
2   a = read_async(..)
3   b = read_async(..)
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8 ...

```

Asynchronous I/O Model

- Return quickly after the request is issued.
- Switch to other tasks during the slow transfer.
- Continue computation once the requested data is ready.

Programming Models for GPU-Centric I/O

Async I/O enables overlapping the slow transfer with other tasks.

```

1 def kernel_block_access
2   a = read_block(..)
3   b = read_block(..)
4   compute(a, b)
5 ...

```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```

1 def kernel_async_access
2   a = read_async(..)
3   b = read_async(..)
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8 ...

```

Asynchronous I/O Model

- Return quickly after the request is issued.
- Switch to other tasks during the slow transfer.
- Continue computation once the requested data is ready.



How can we support an async GPU-centric storage access model efficiently?

Programming Models for GPU-Centric I/O

Async I/O enables overlapping the slow transfer with other tasks.

```

1 def kernel_block_access
2   a = read_block(..)
3   b = read_block(..)
4   compute(a, b)
5 ...

```

Synchronous I/O Model
(BaM)

- Return after the slow transfer is finished.
- GPU thread must wait for the slow transfer.

```

1 def kernel_async_access
2   a = read_async(..)
3   b = read_async(..)
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8 ...

```

Asynchronous I/O Model

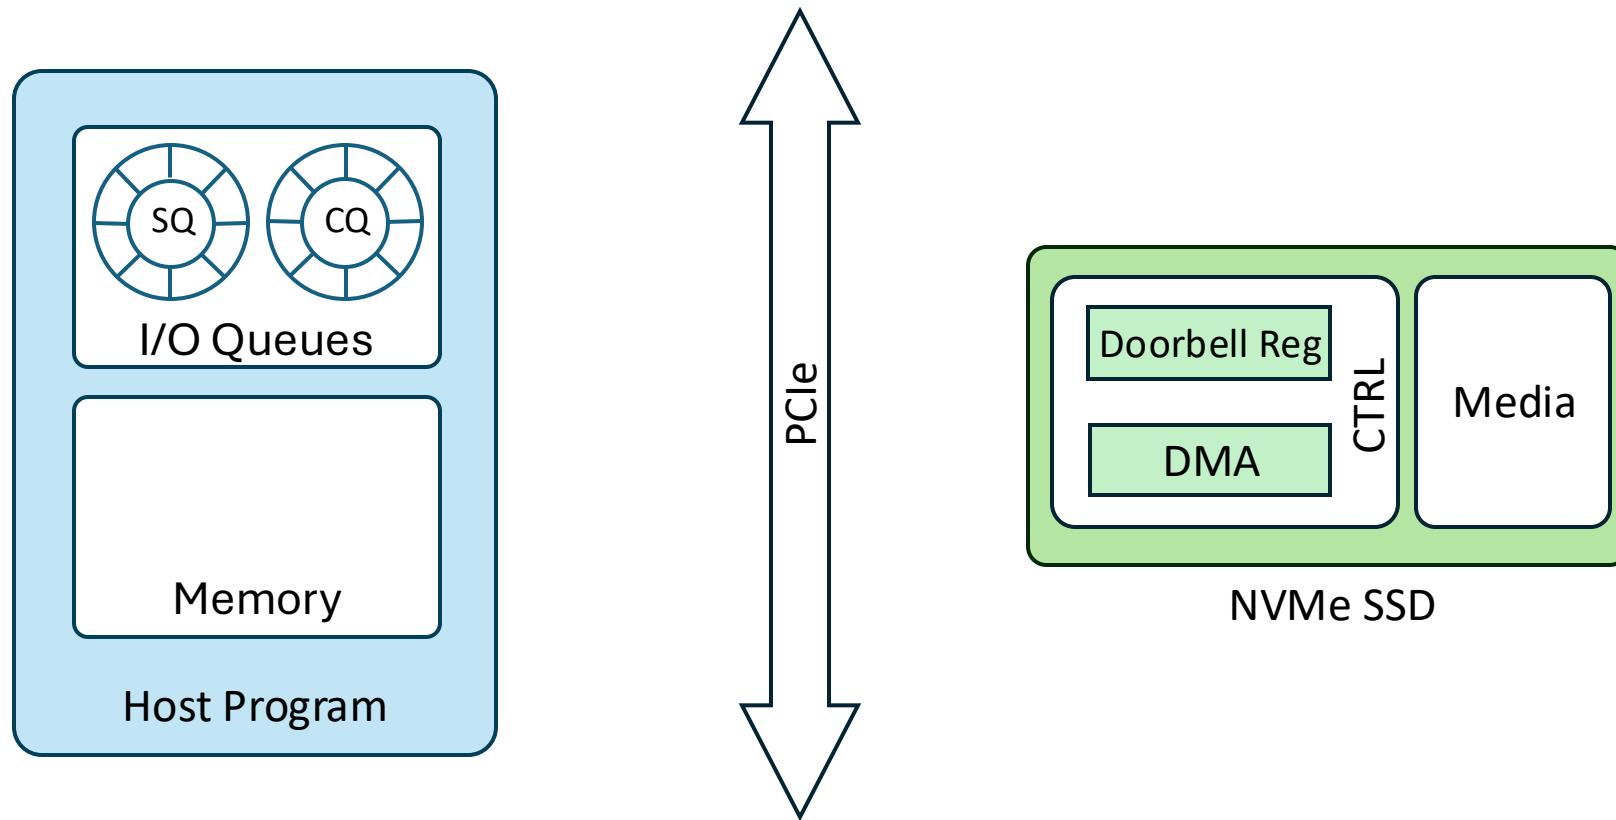
- Return quickly after the request is issued.
- Switch to other tasks during the slow transfer.
- Continue computation once the requested data is ready.



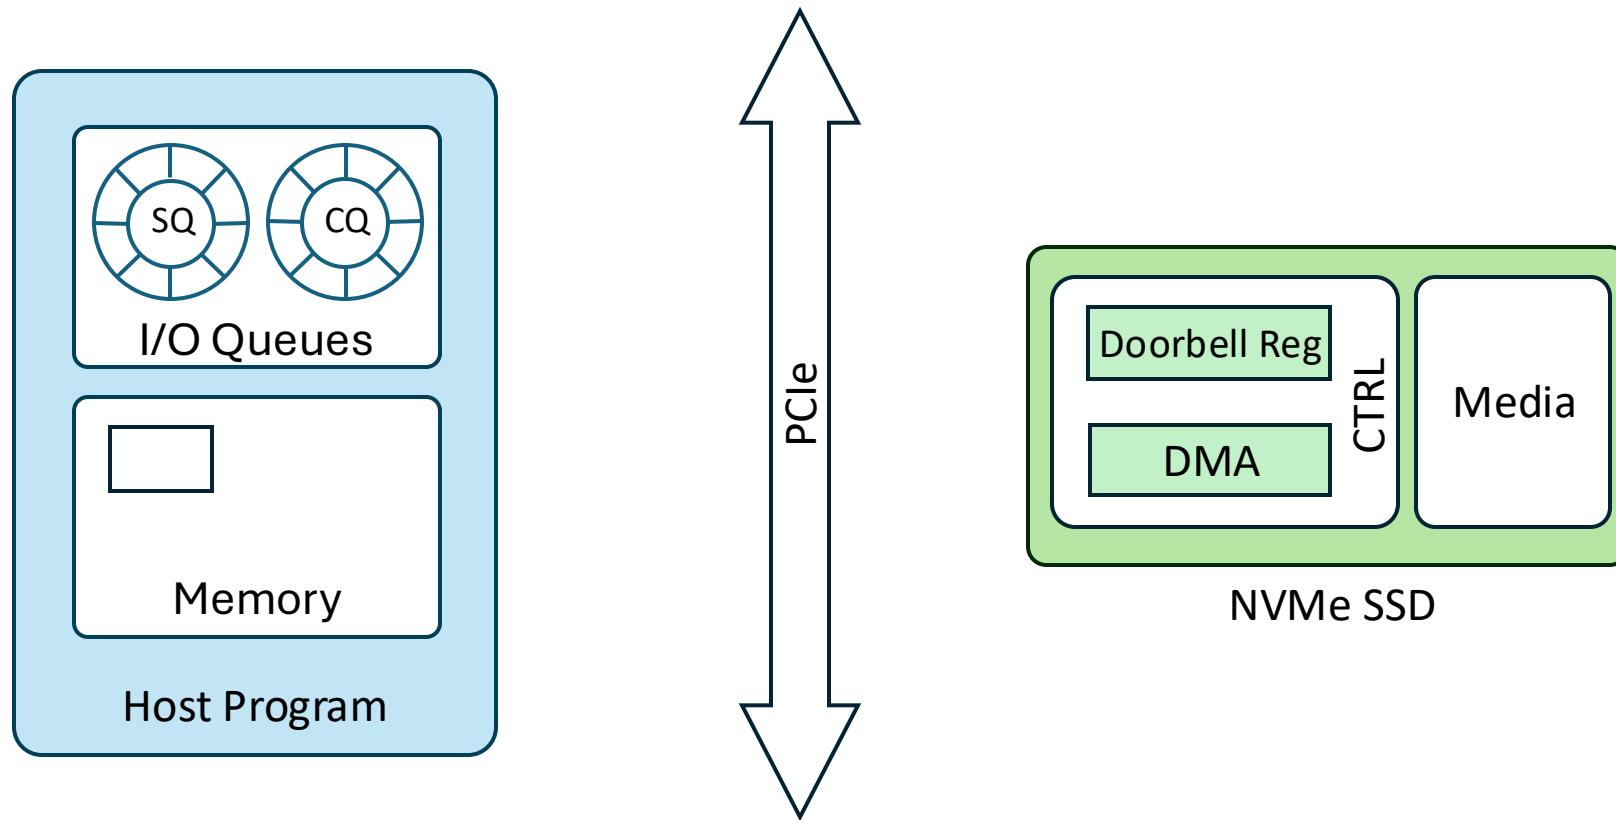
How can we support an async GPU-centric storage access model efficiently?

➤ Use AGILE! (this work) 😊

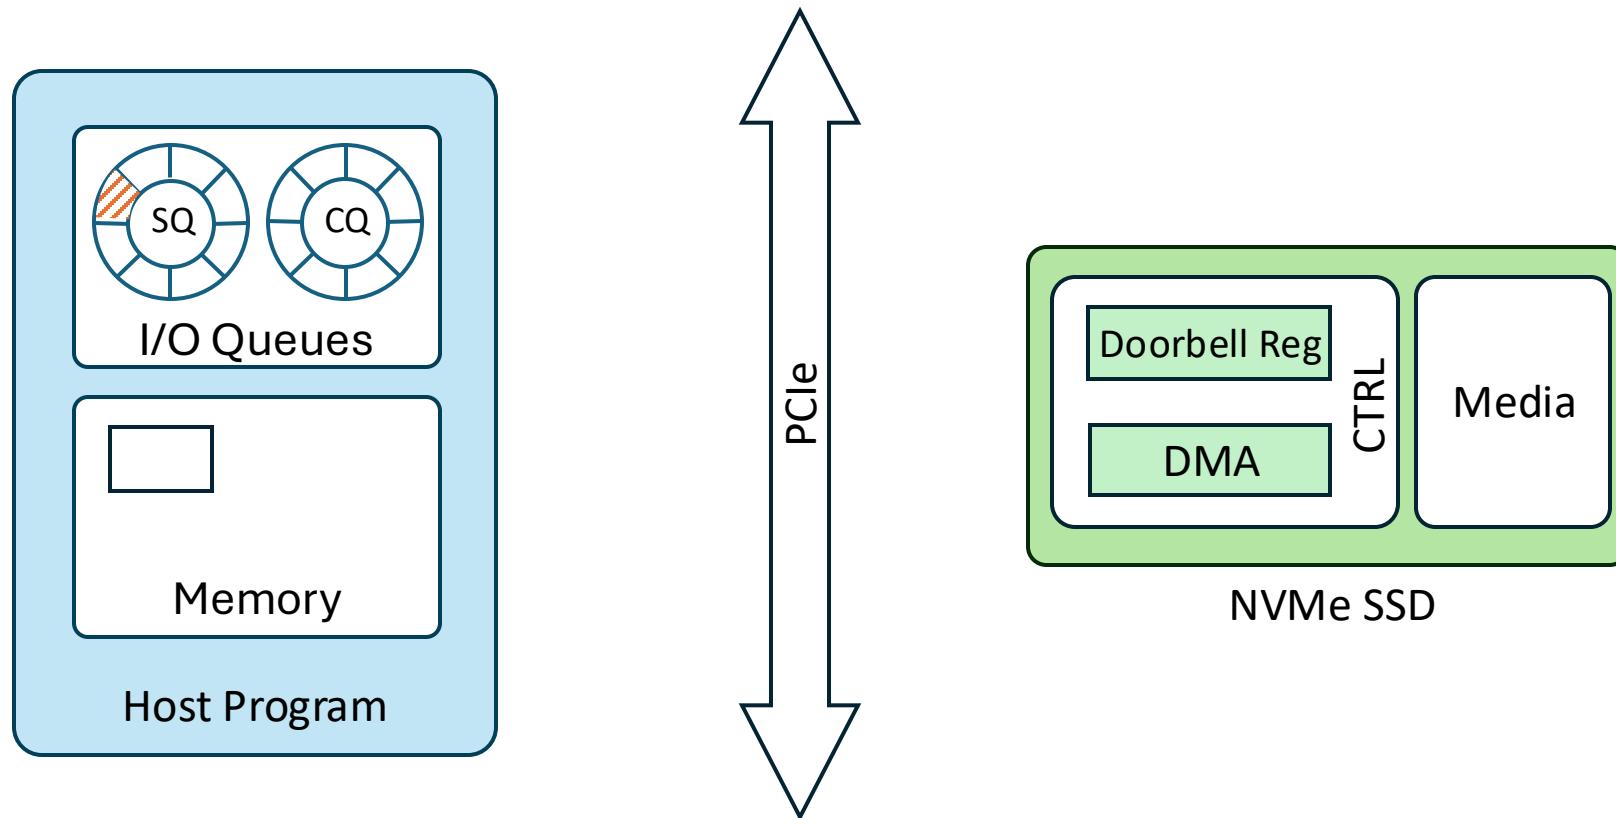
Background of NVMe Protocol



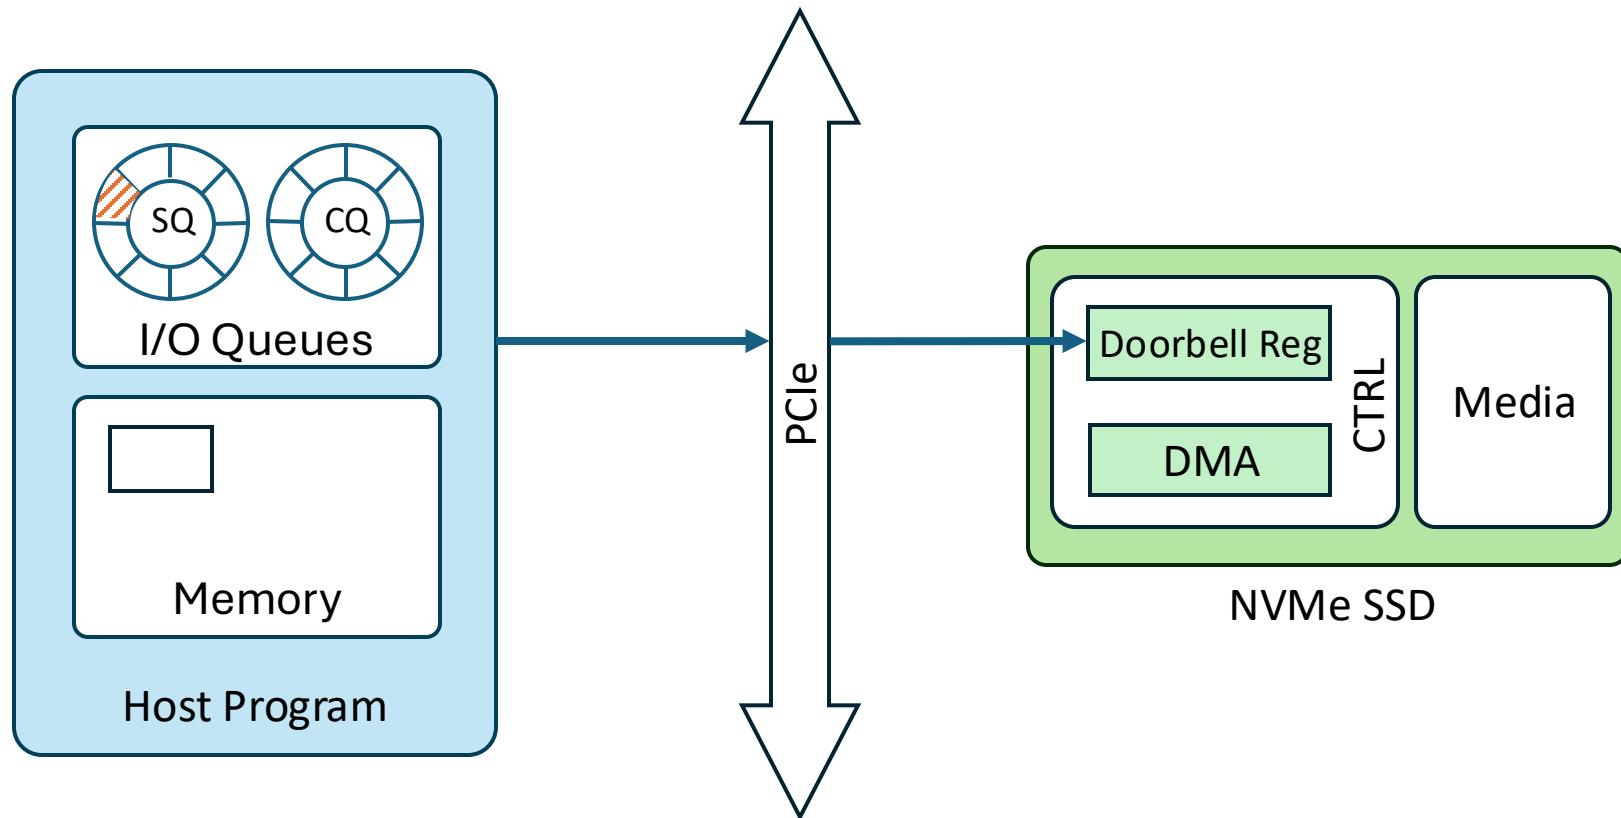
Background of NVMe Protocol



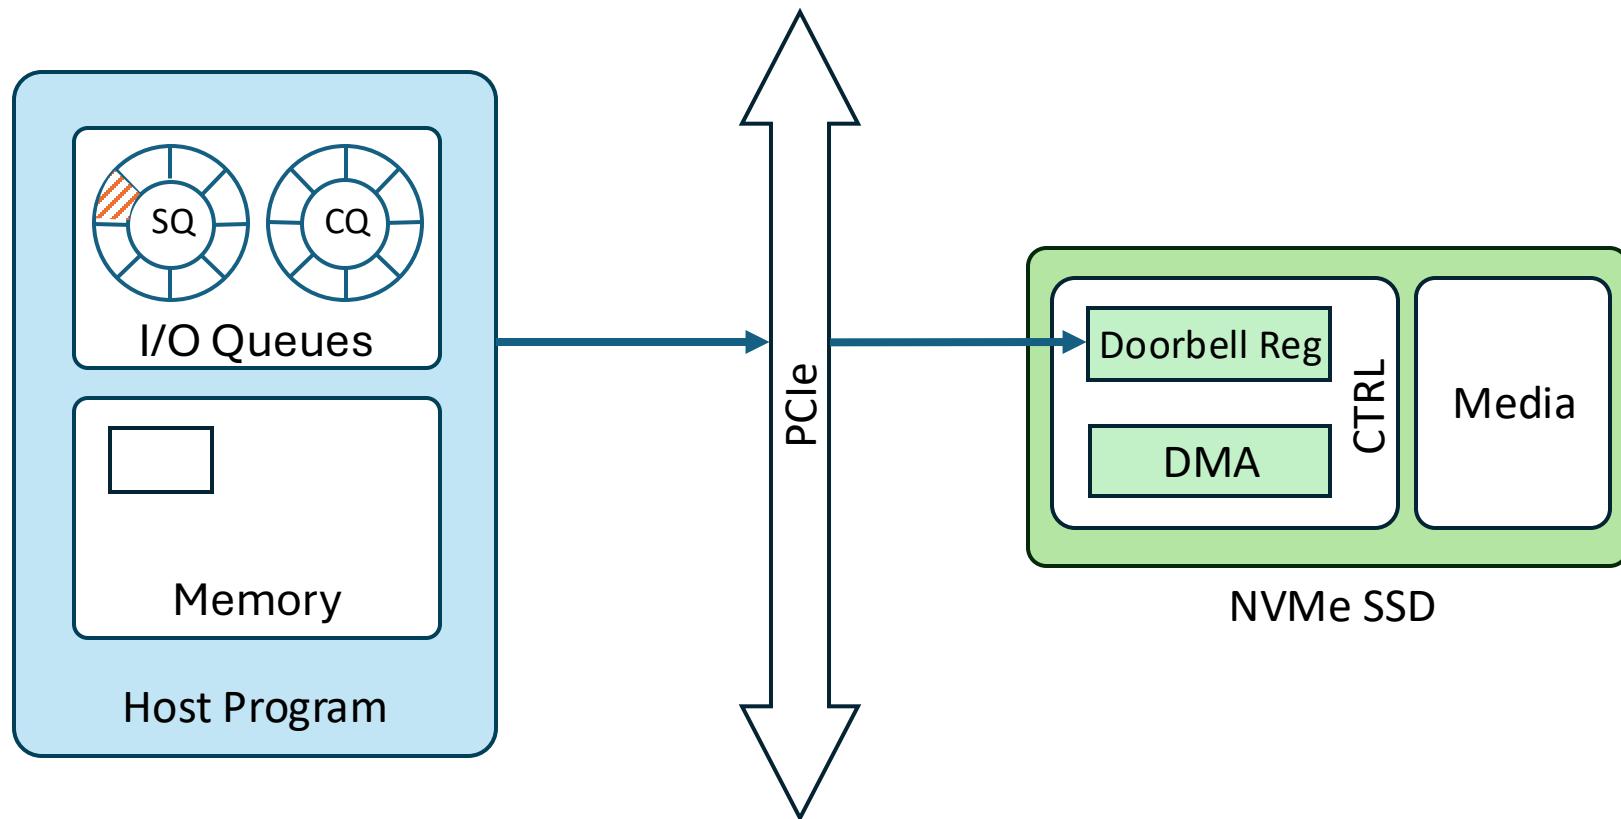
Background of NVMe Protocol



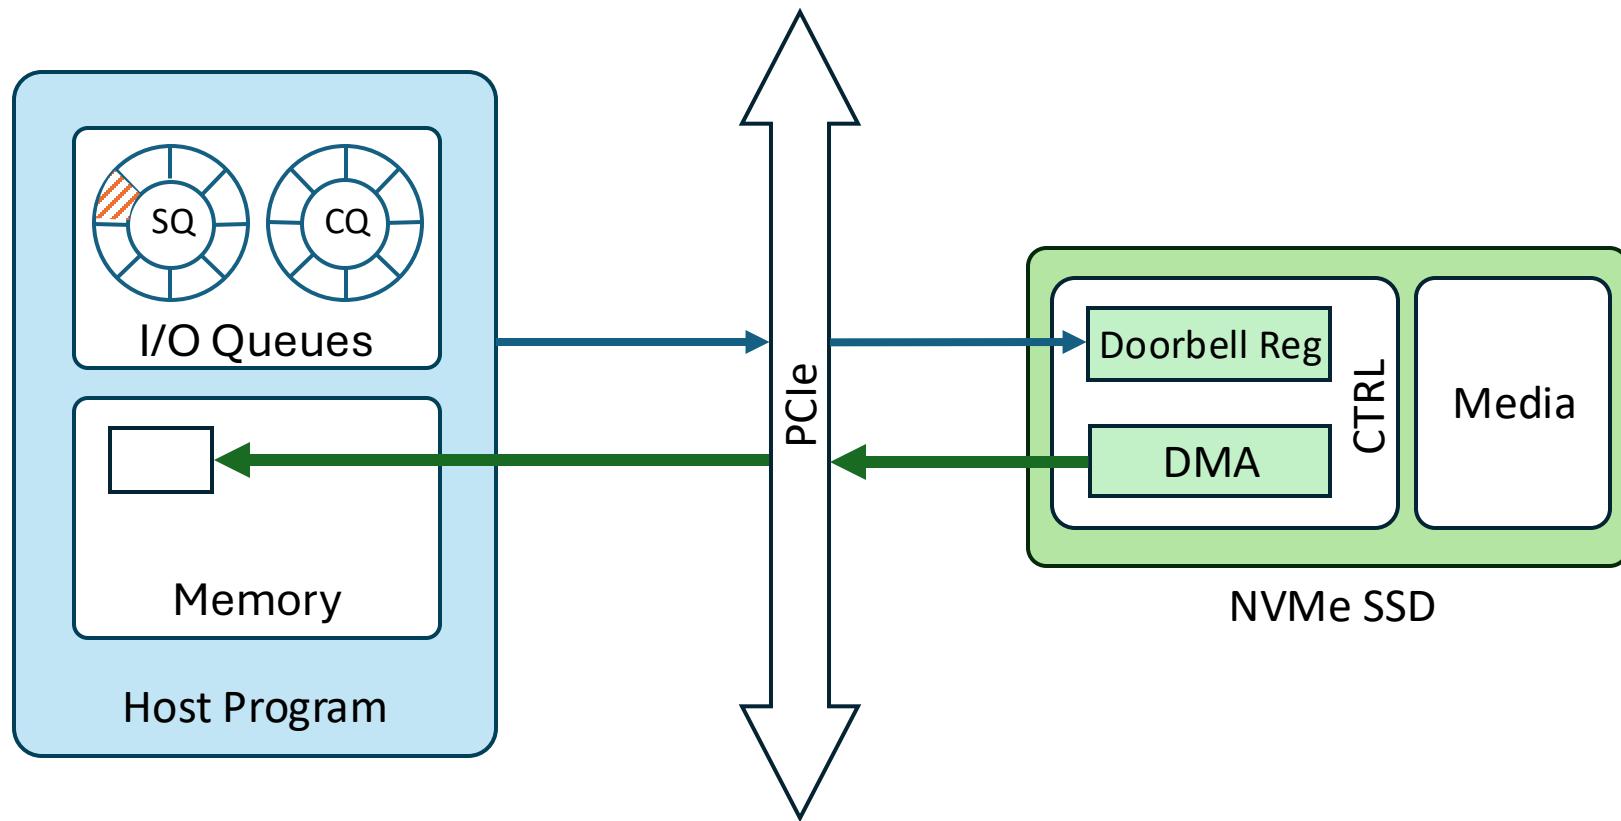
Background of NVMe Protocol



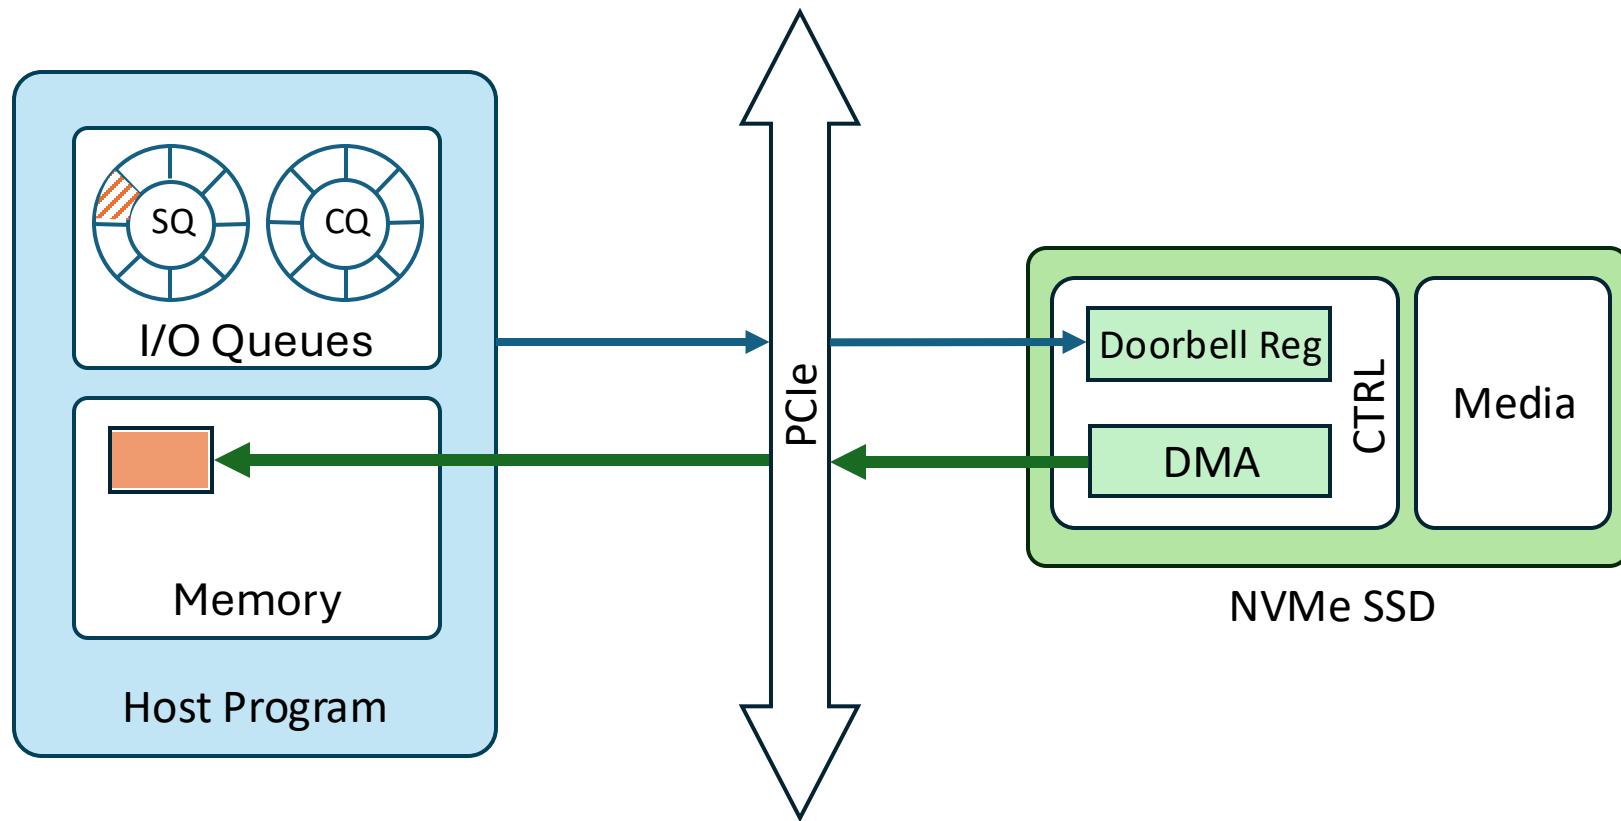
Background of NVMe Protocol



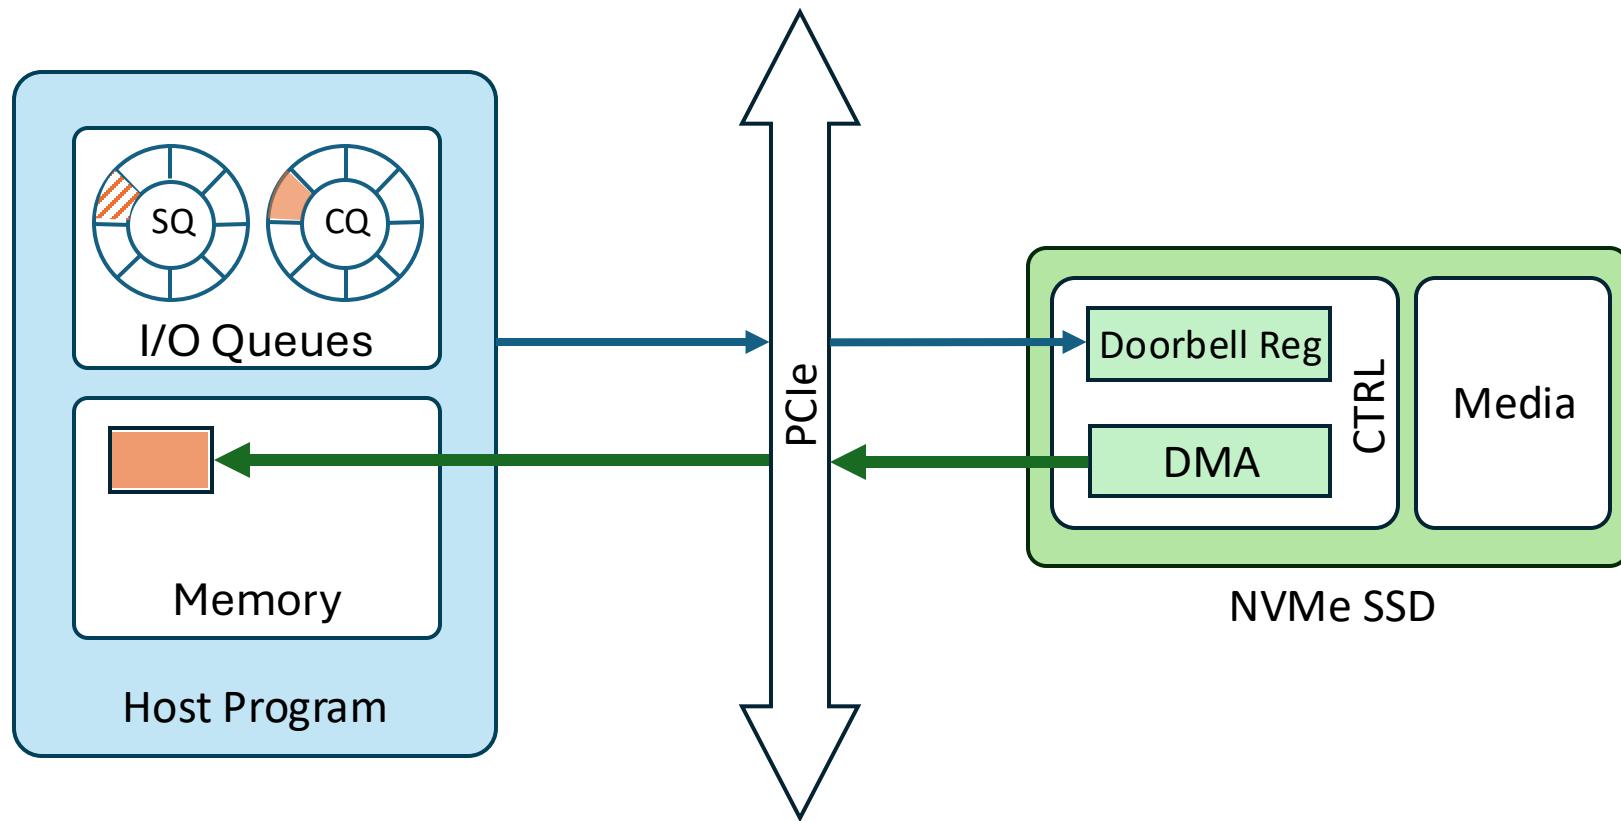
Background of NVMe Protocol



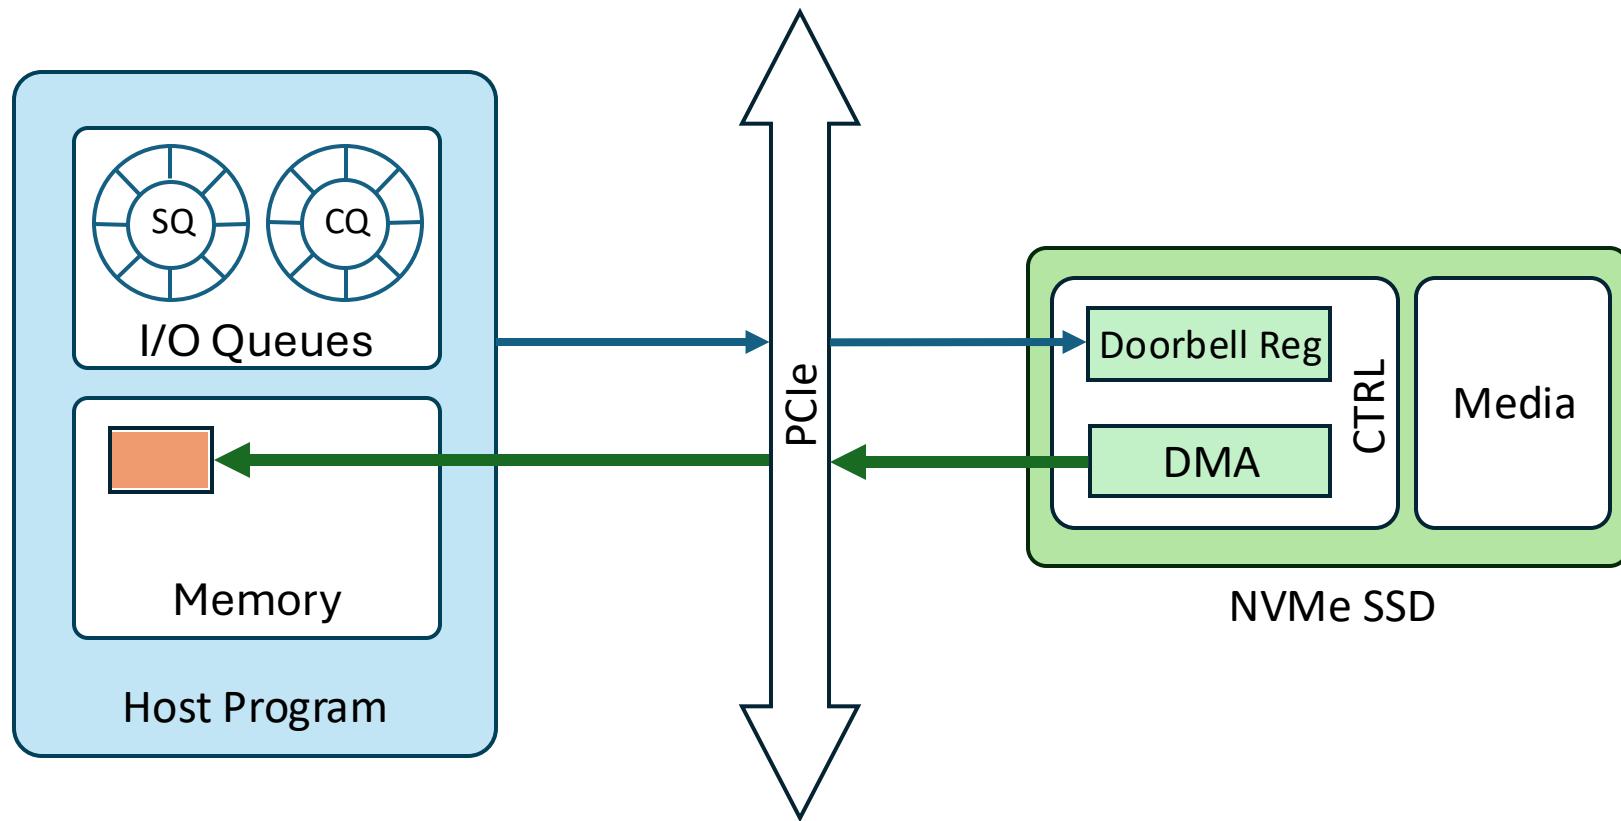
Background of NVMe Protocol



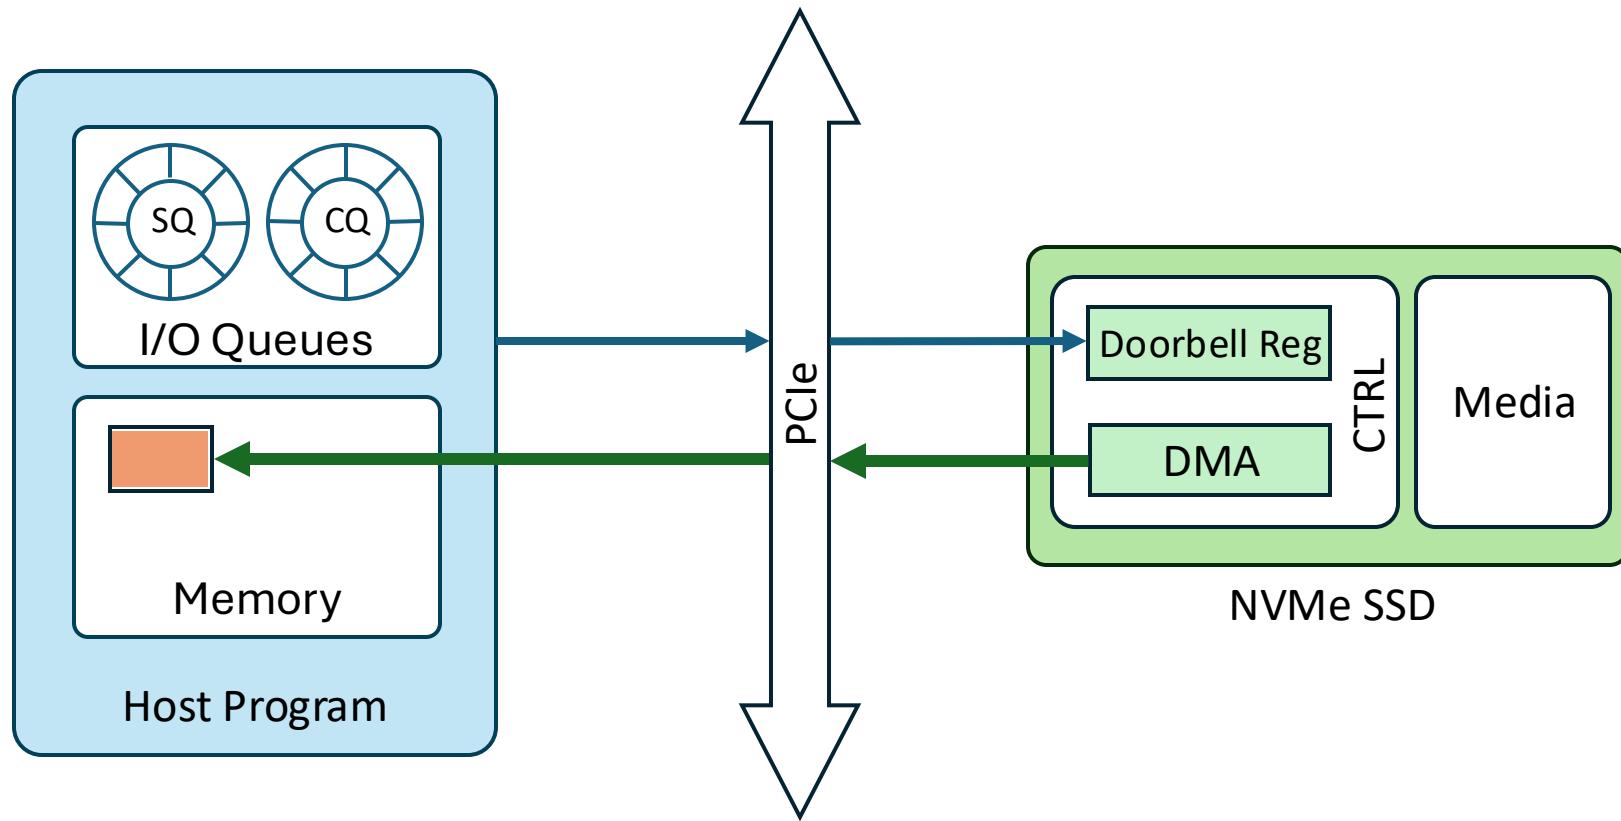
Background of NVMe Protocol



Background of NVMe Protocol

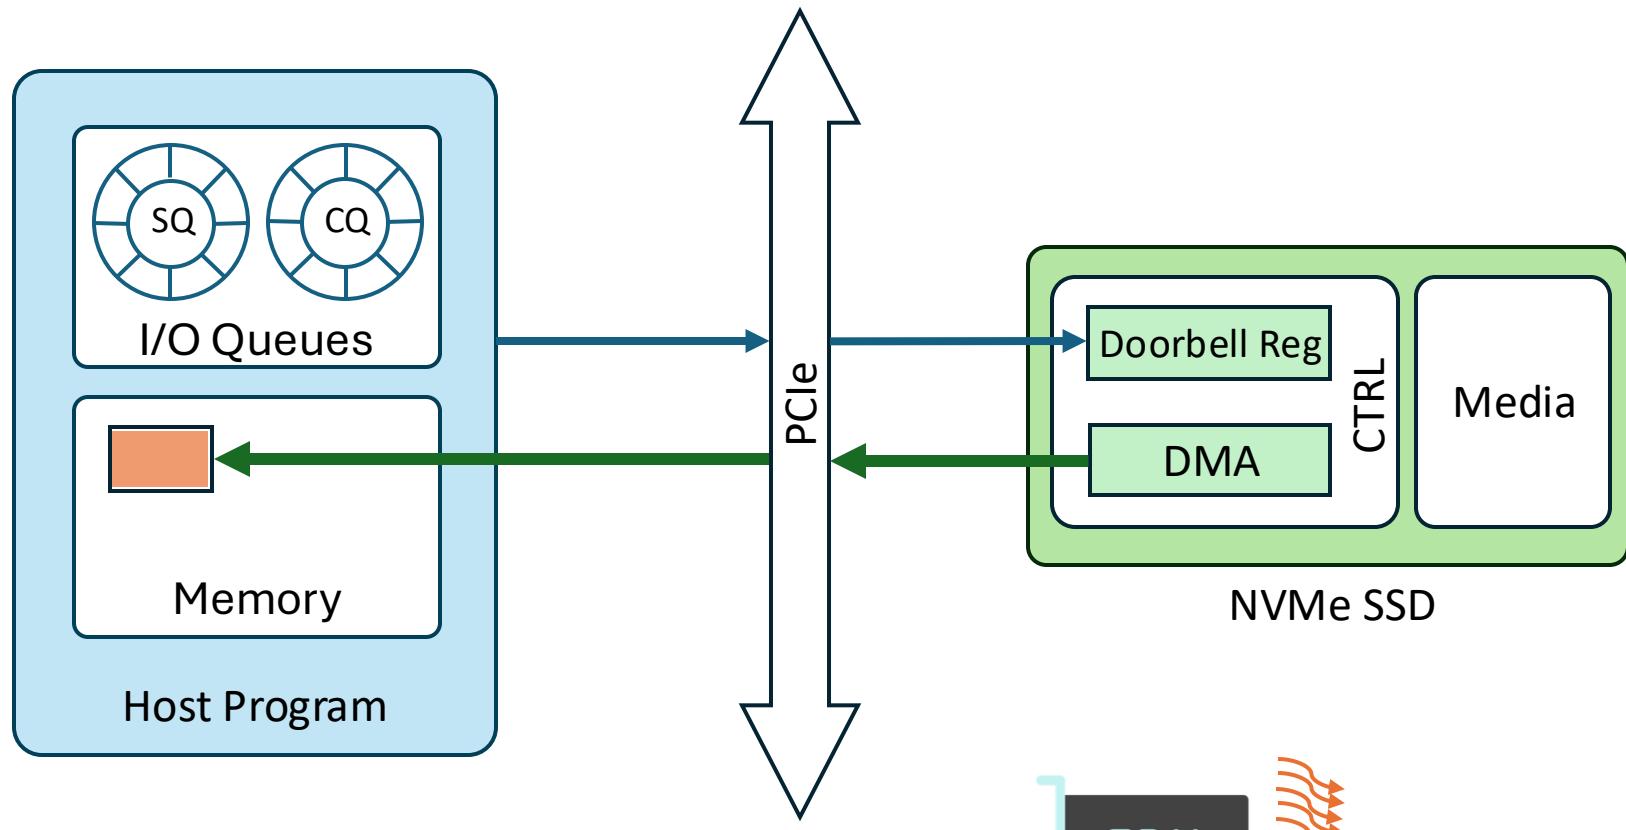


Background of NVMe Protocol



The NVMe protocol is fundamentally asynchronous

Background of NVMe Protocol



The NVMe protocol is fundamentally asynchronous

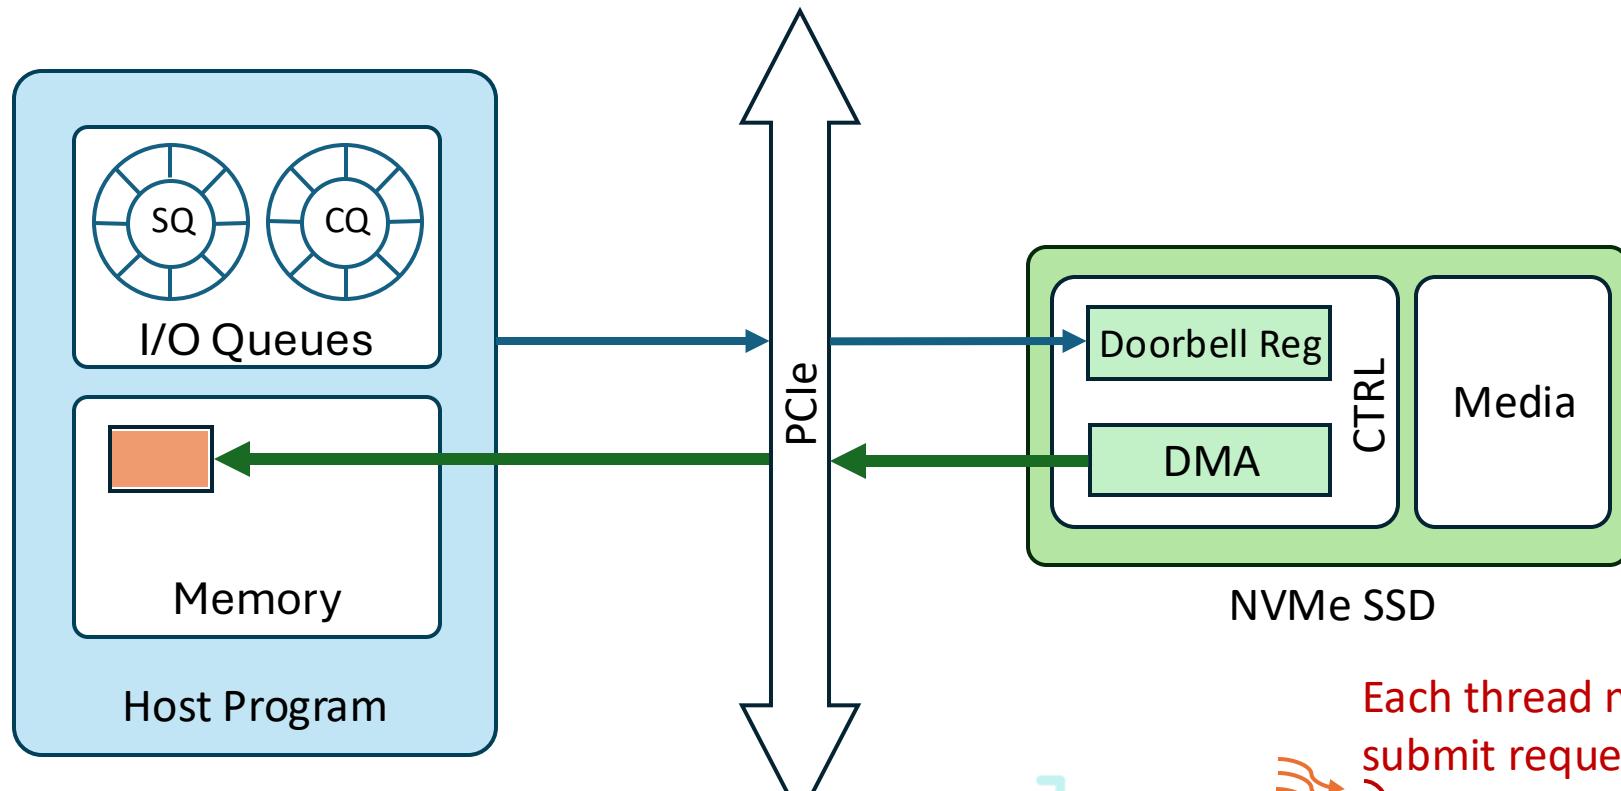


Massive parallel
threads

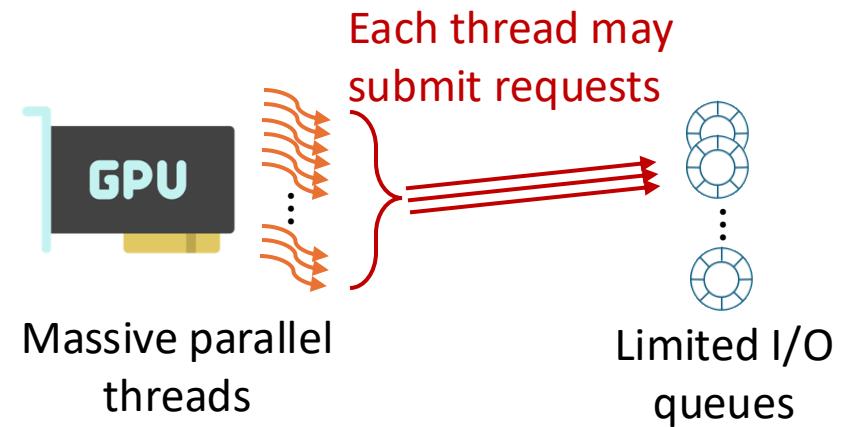


Limited I/O
queues

Background of NVMe Protocol



The NVMe protocol is fundamentally asynchronous



Deadlock Risks in Async GPU-Centric NVMe I/O

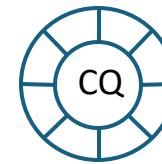
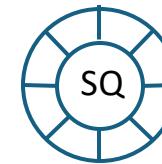
Async GPU-centric NVMe I/O can easily lead to a deadlock.

Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

- A deadlock example in GPU-centric NVMe read:

```
1 def kernel_async_access
2     a = read_async...
3     b = read_async...
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

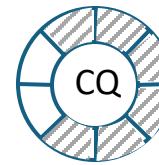
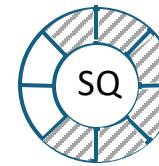


Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

- A deadlock example in GPU-centric NVMe read:

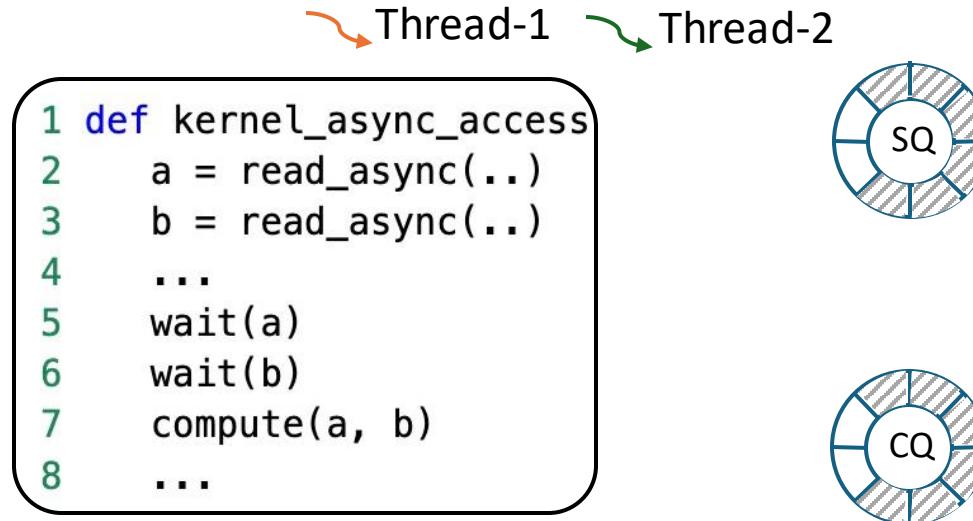
```
1 def kernel_async_access
2     a = read_async...
3     b = read_async...
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```



Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

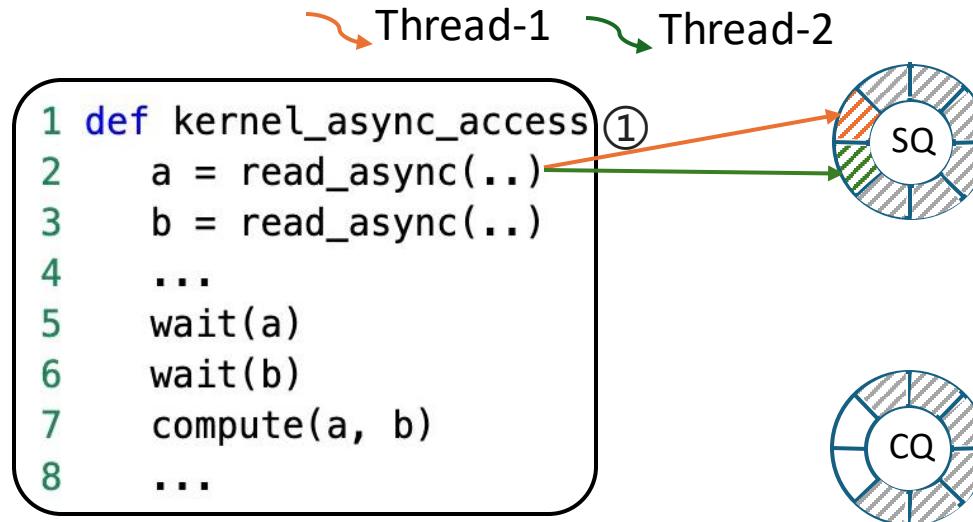
- A deadlock example in GPU-centric NVMe read:



Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

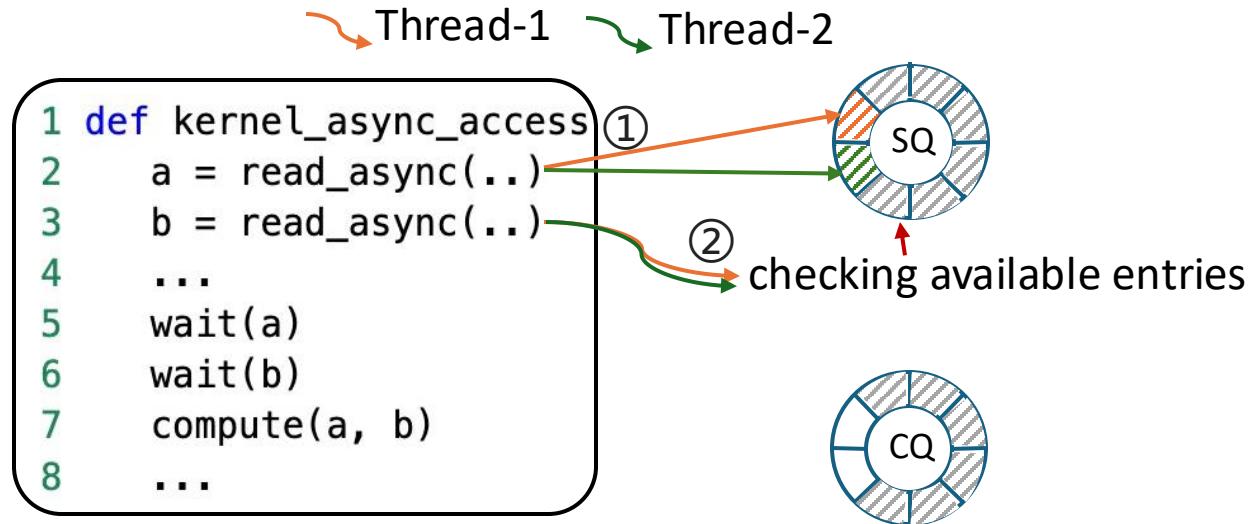
- A deadlock example in GPU-centric NVMe read:



Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

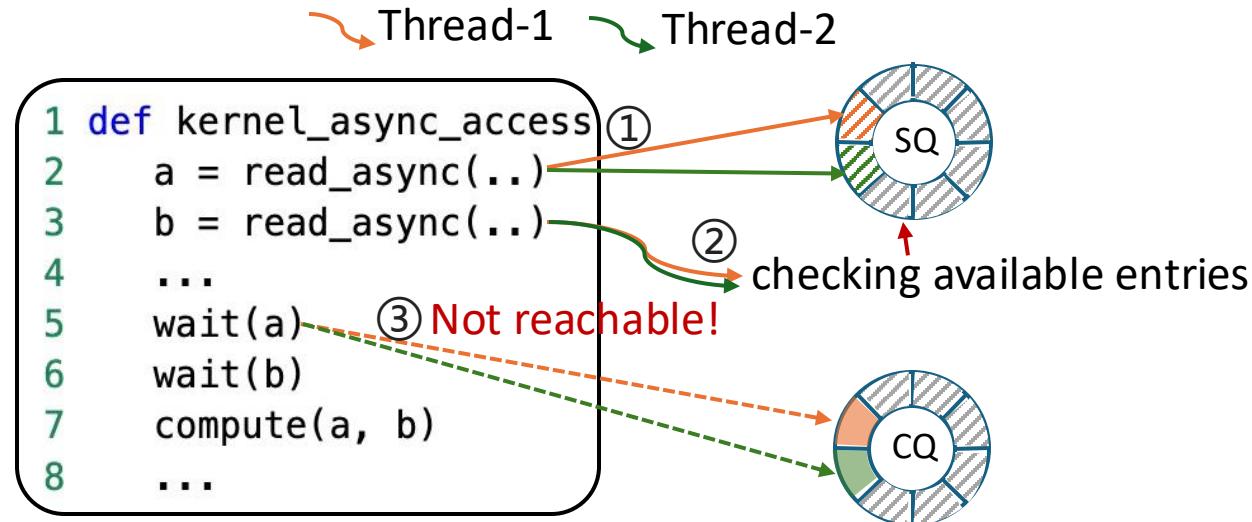
- A deadlock example in GPU-centric NVMe read:



Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

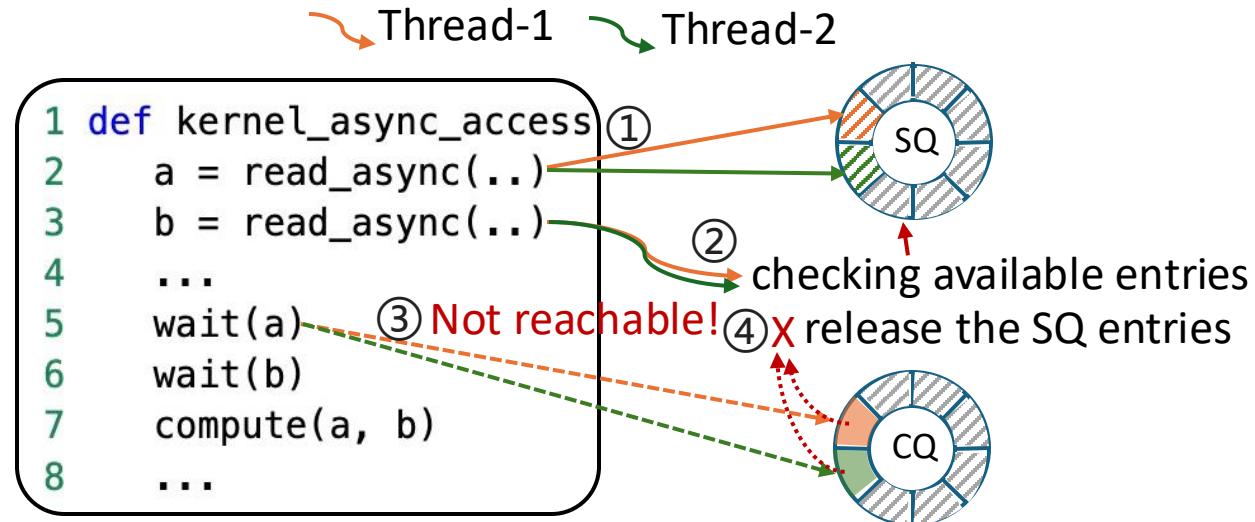
- A deadlock example in GPU-centric NVMe read:



Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

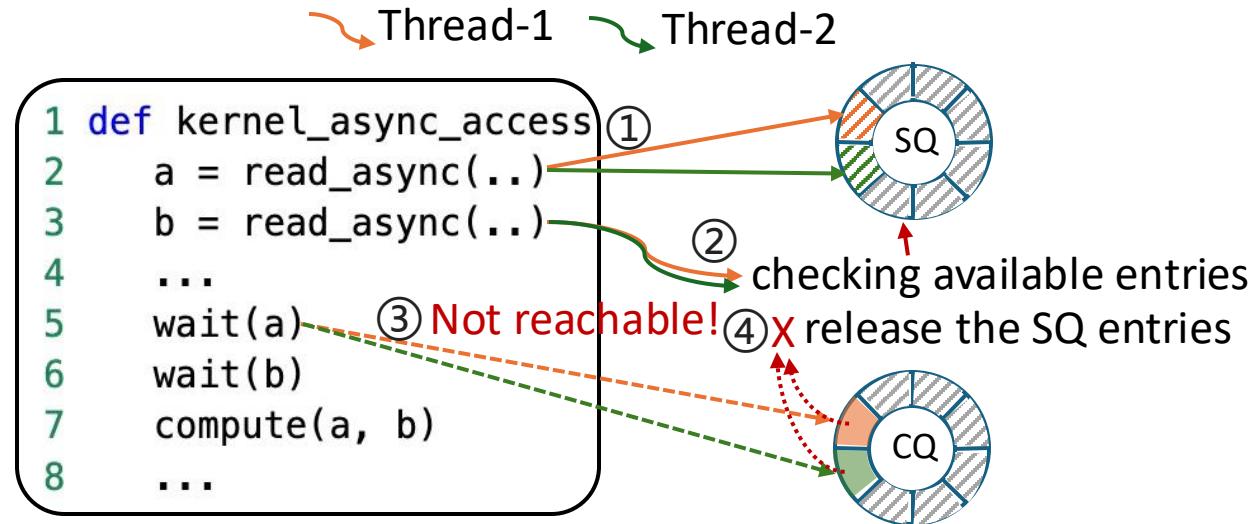
- A deadlock example in GPU-centric NVMe read:



Deadlock Risks in Async GPU-Centric NVMe I/O

Async GPU-centric NVMe I/O can easily lead to a deadlock.

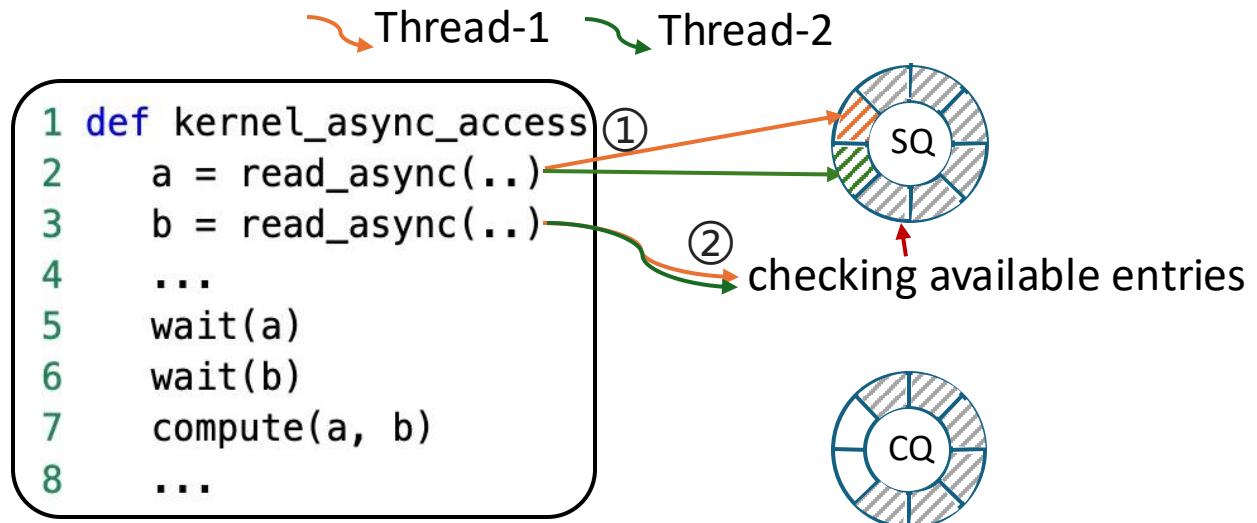
- A deadlock example in GPU-centric NVMe read:



- The root cause of this deadlock is allowing user threads to **hold multiple locks** that can block other user threads. (I/O queues, software cache lines)

AGILE Polling Service

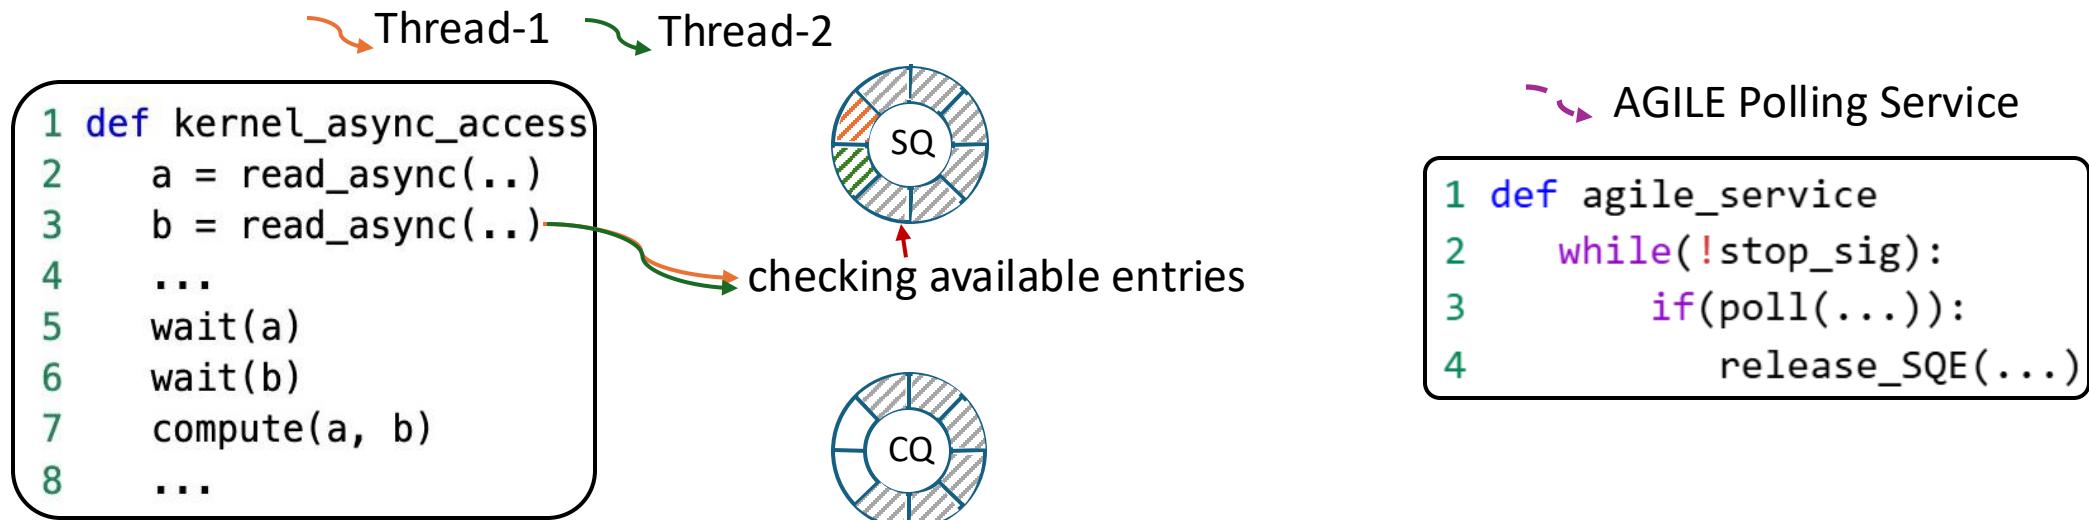
To avoid deadlock from asynchronous NVMe I/O:



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

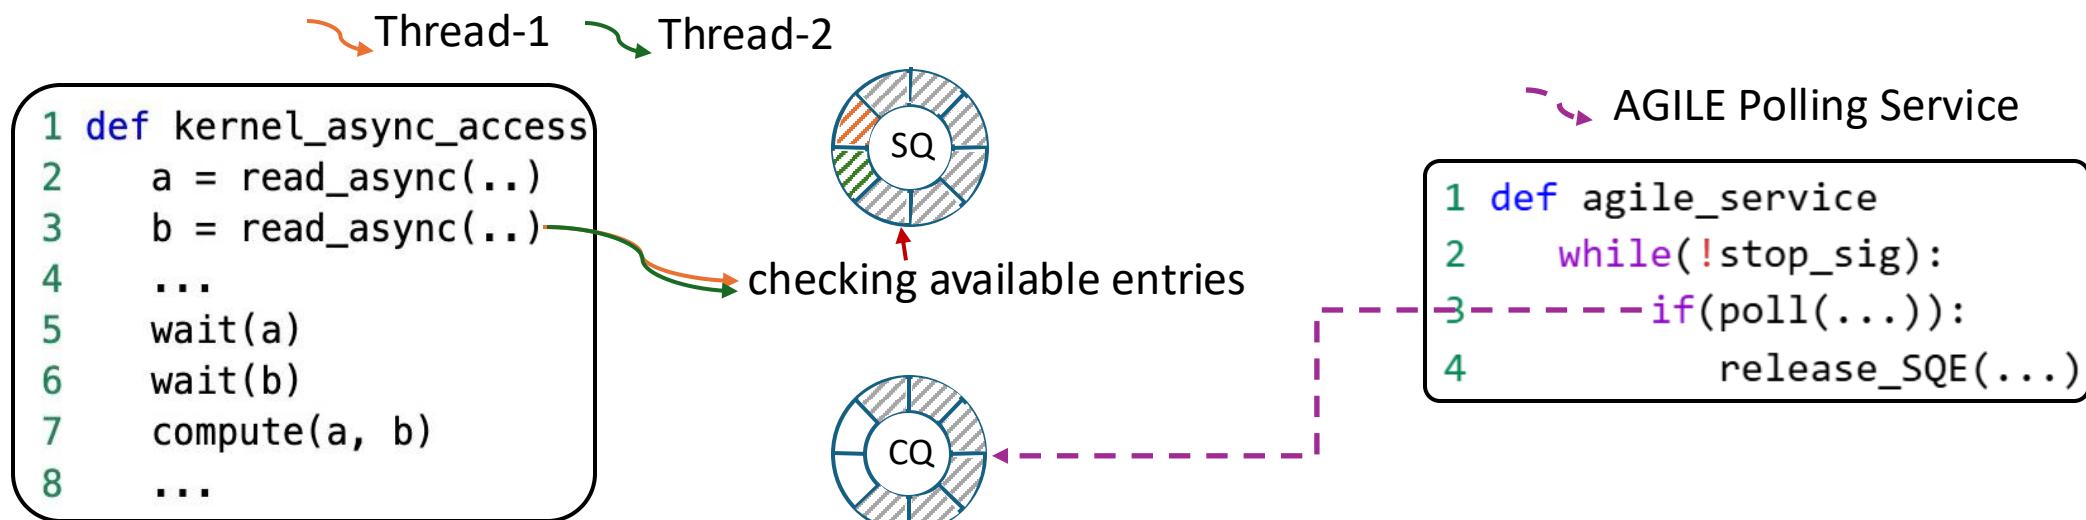
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

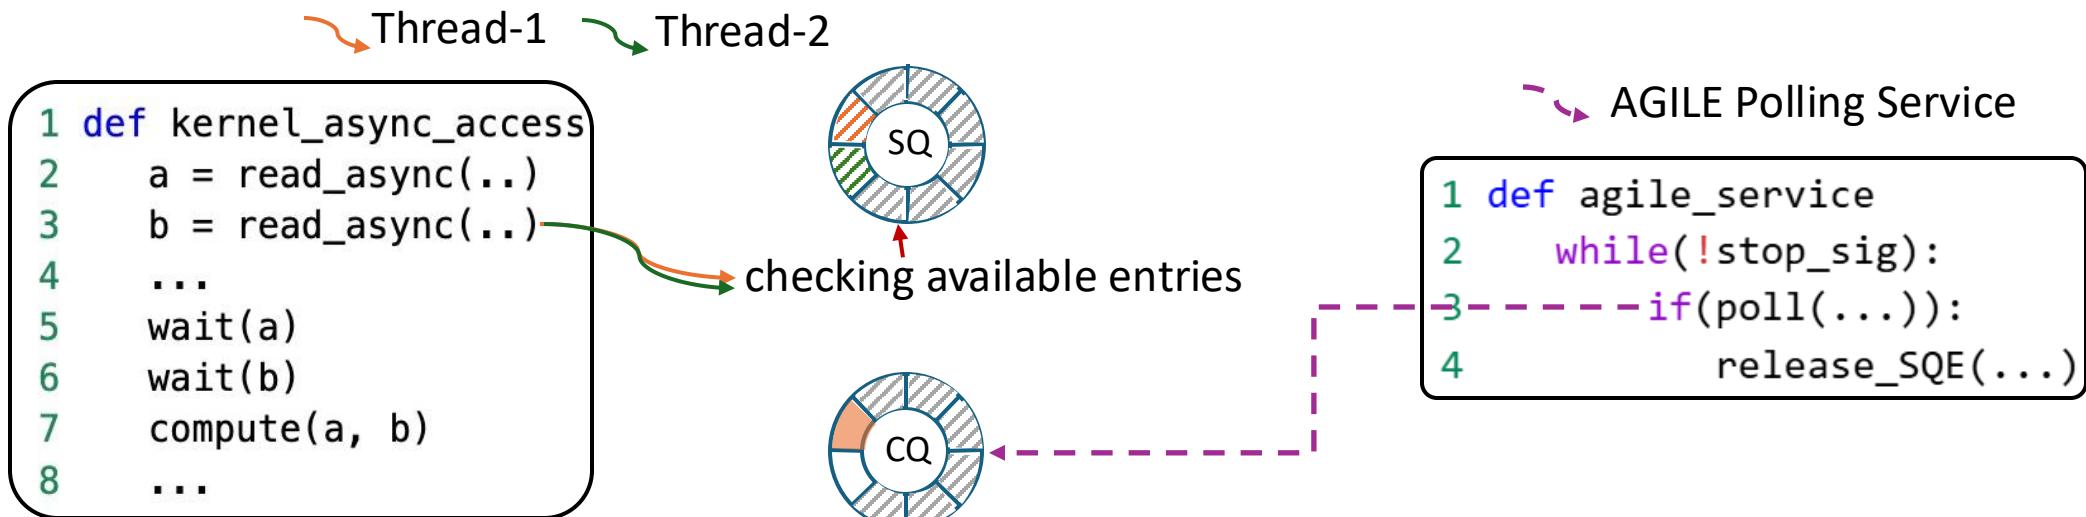
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

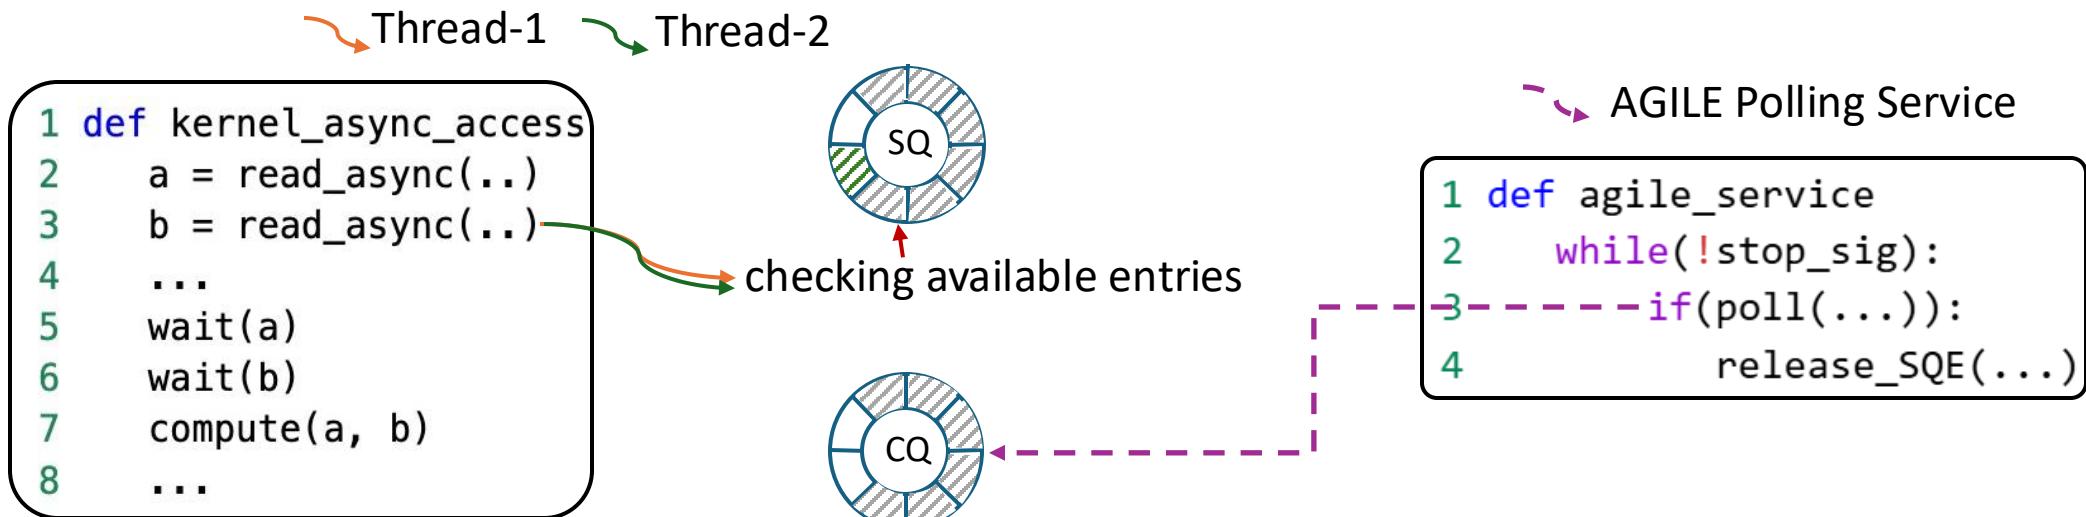
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

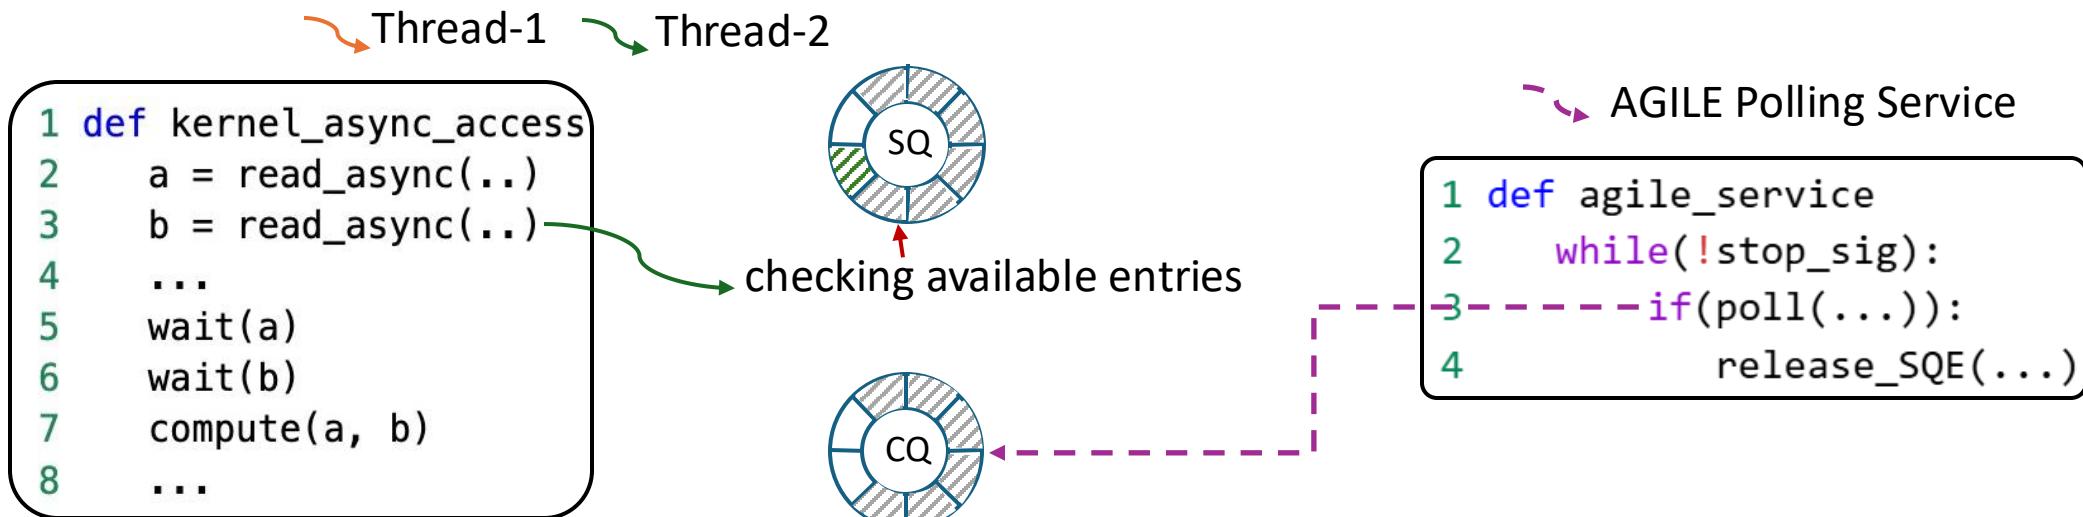
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

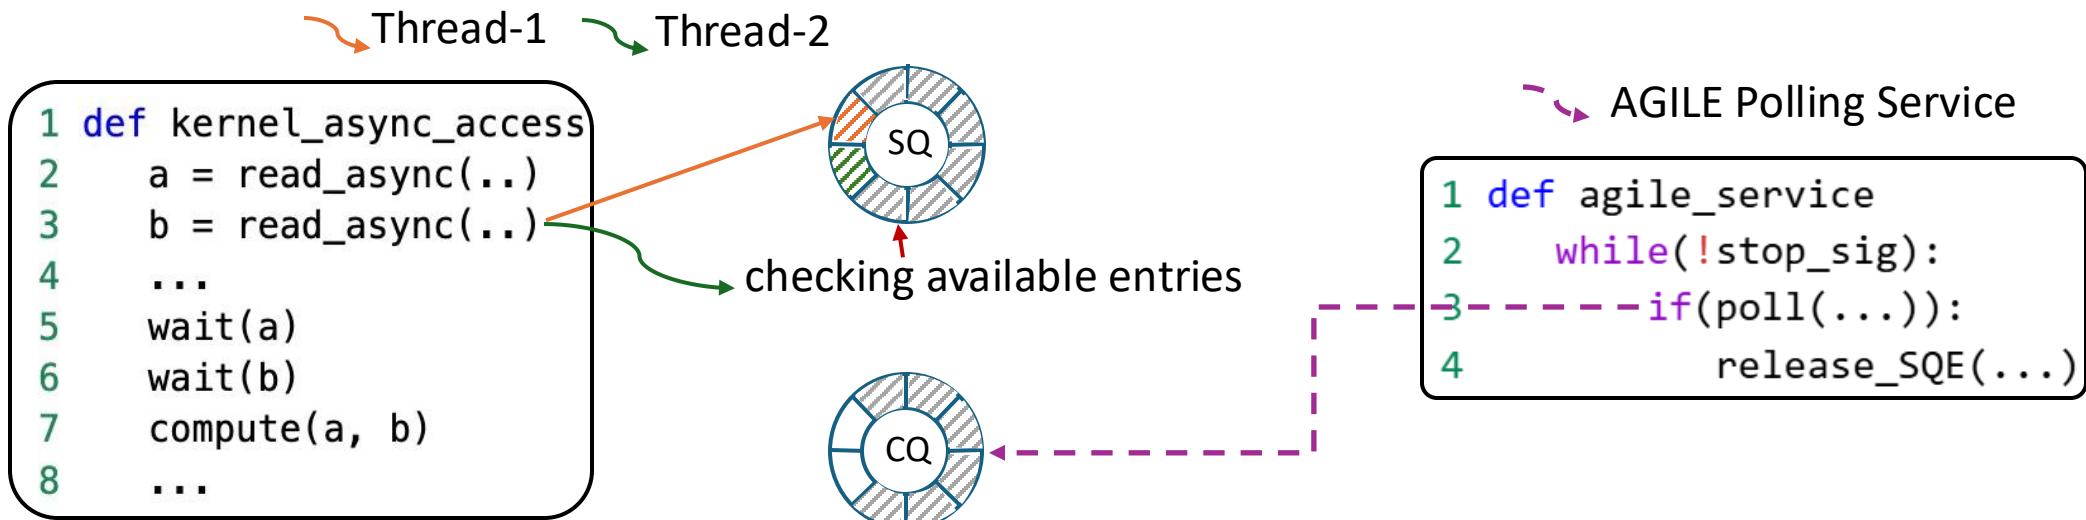
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

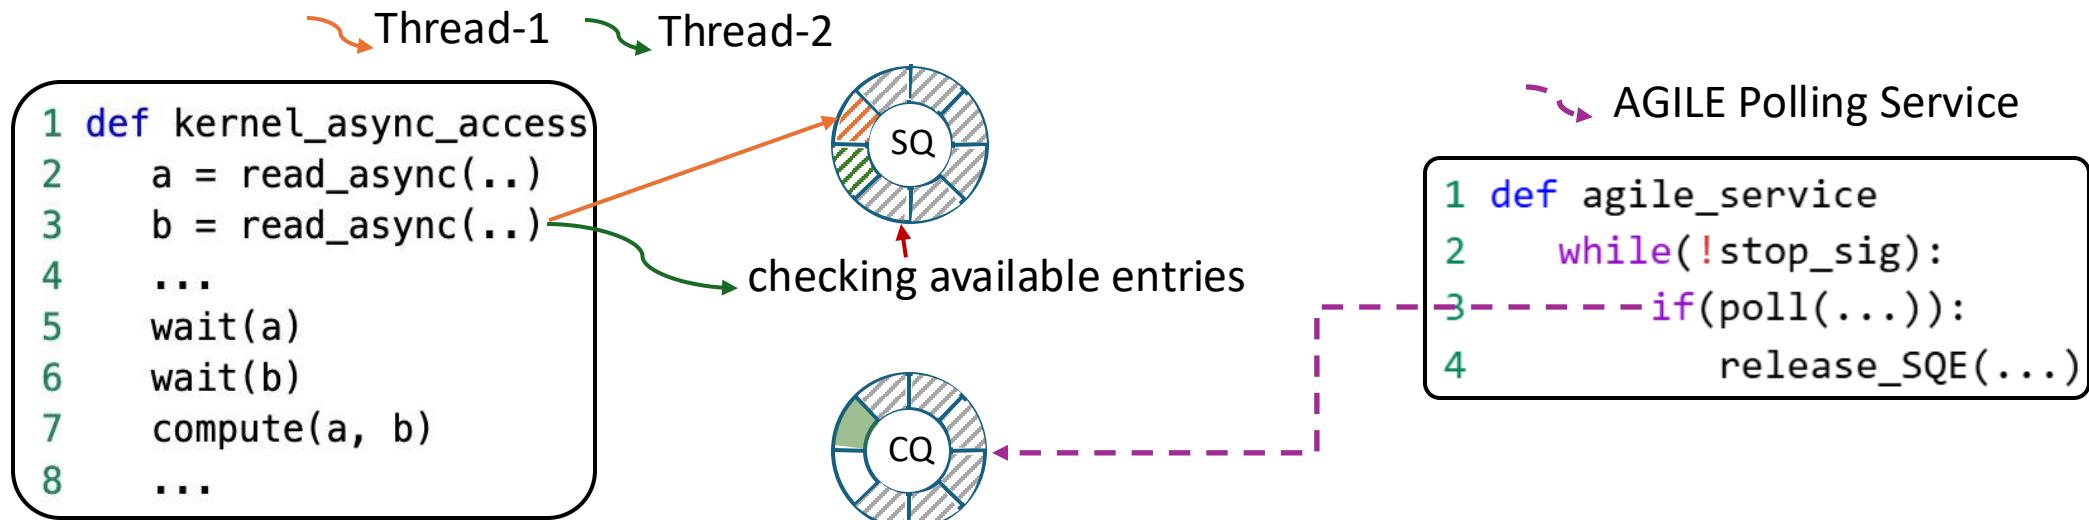
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

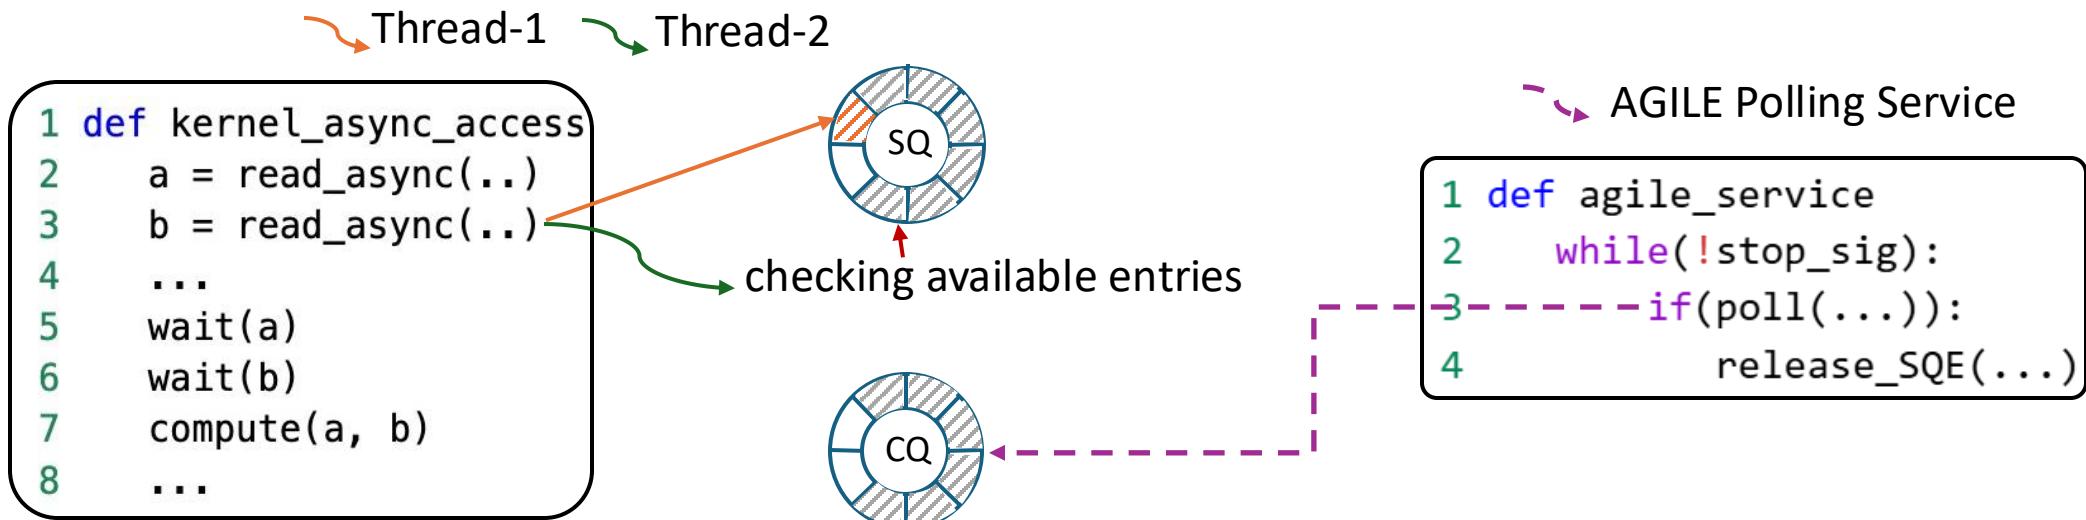
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

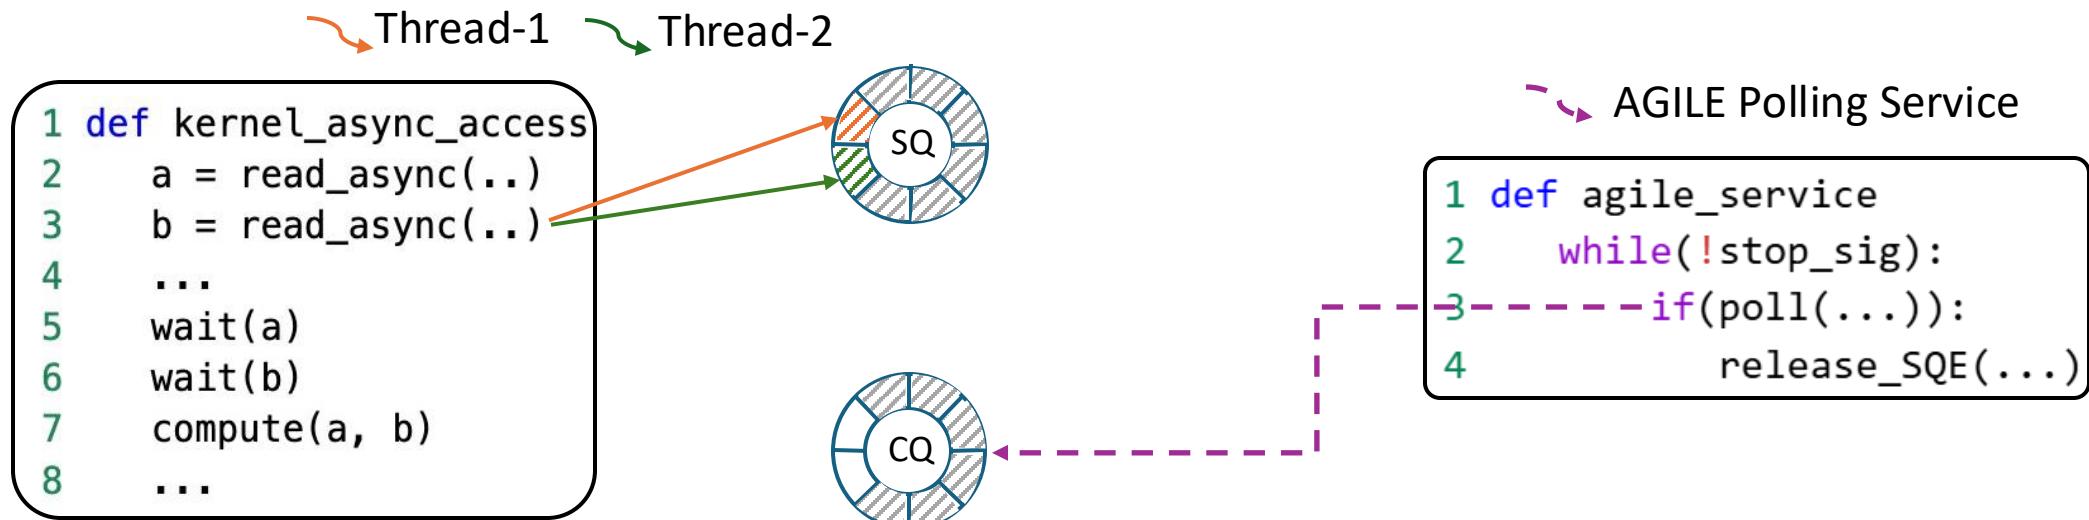
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

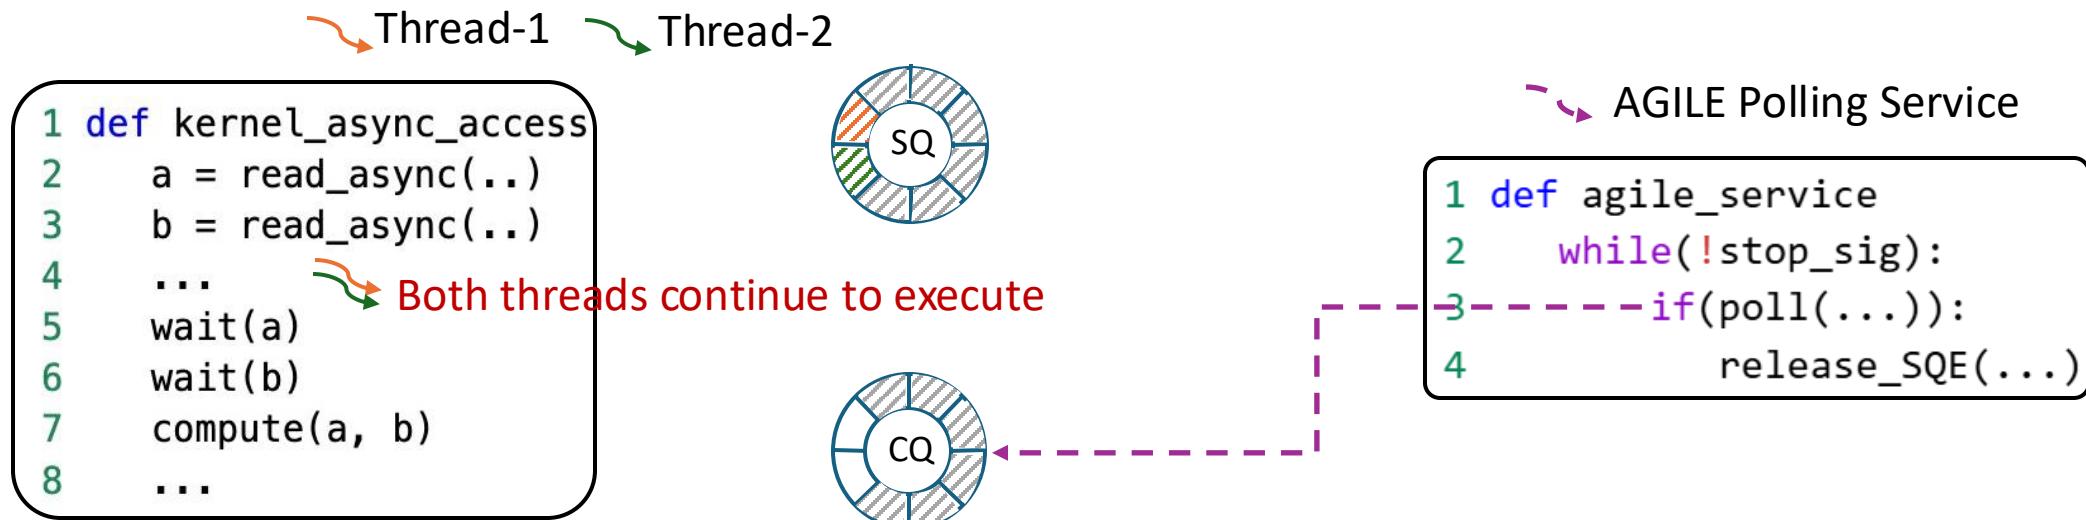
- A warp-centric AGILE polling service running in the background



AGILE Polling Service

To avoid deadlock from asynchronous NVMe I/O:

- A warp-centric AGILE polling service running in the background

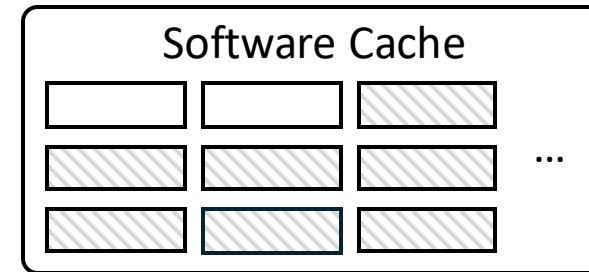


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

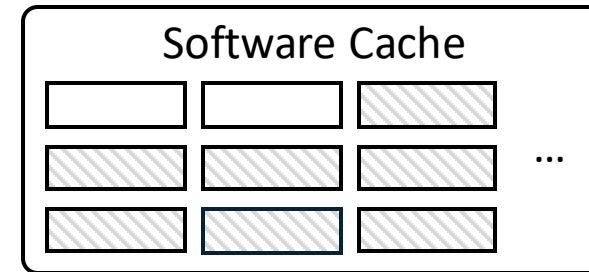


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2 → a = read_async...
3   b = read_async...
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8   ...
```

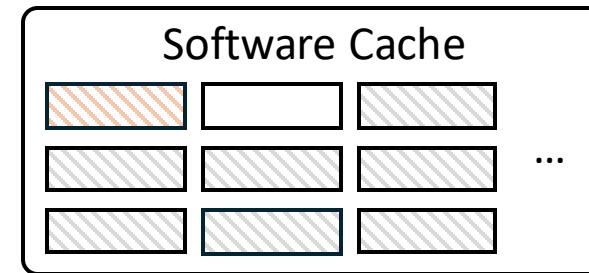


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2 → a = read_async...
3   b = read_async...
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8   ...
```

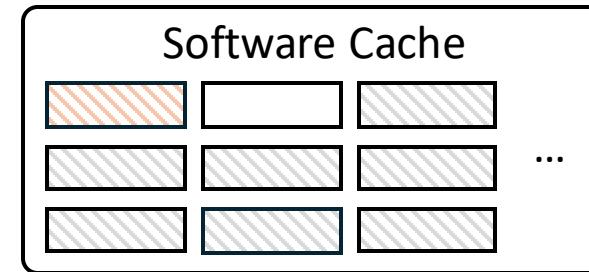


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3 → b = read_async(...)
4 ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8 ...
```

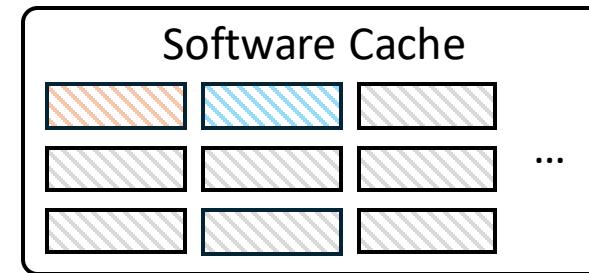


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3 → b = read_async(...)
4 ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8 ...
```

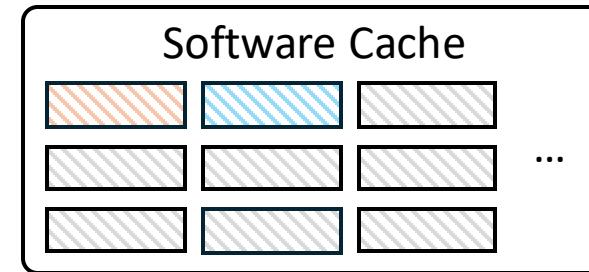


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     → ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

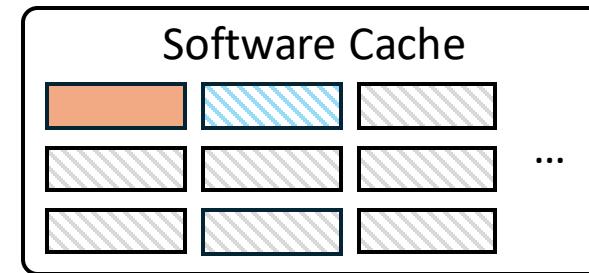


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     → ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

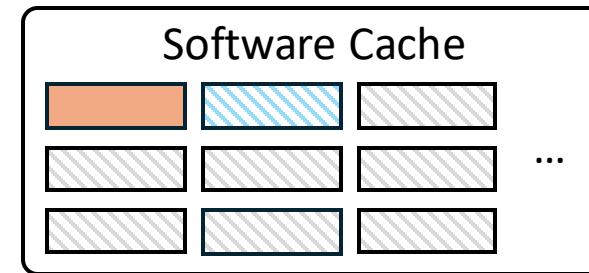


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5 →  wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

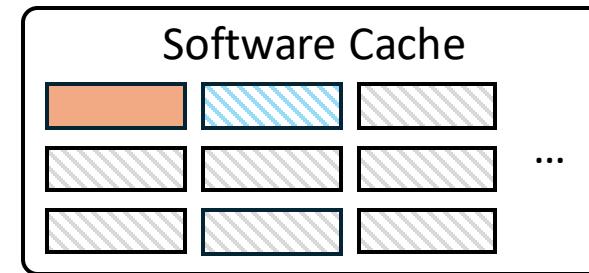


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5 →  wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

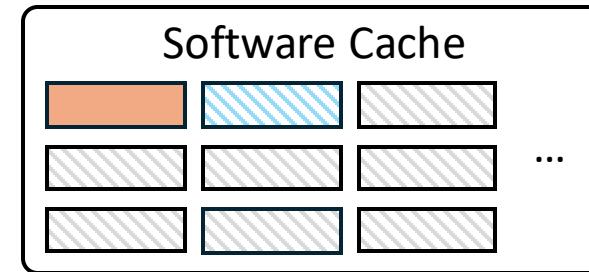


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6 →  wait(b)
7     compute(a, b)
8     ...
```

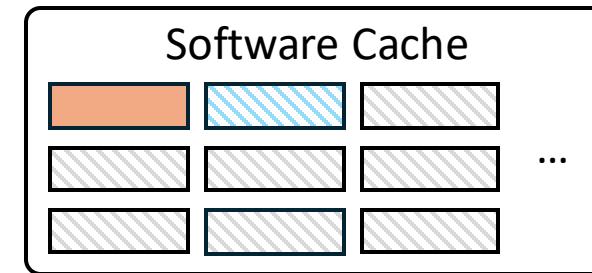


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)      Other threads request new
6 →  wait(b) →  software cache lines
7     compute(a, b)
8     ...
```



AGILE Software Cache

To eliminate deadlock from the software cache:

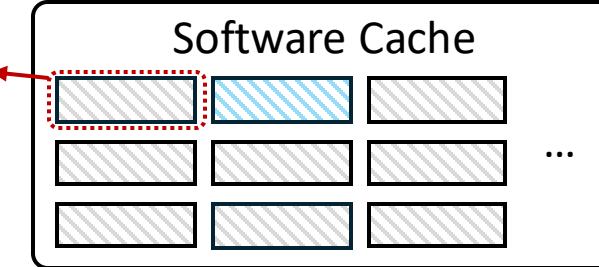
- AGILE does not allow user threads to avoid cache line eviction.

```

1 def kernel_async_access
2     a = read_async...
3     b = read_async...
4     ...
5     wait(a)
6     → wait(b)    Other threads request new
7     compute(a, b)
8     ...

```

a is evicted at line 6 to
avoid blocking others

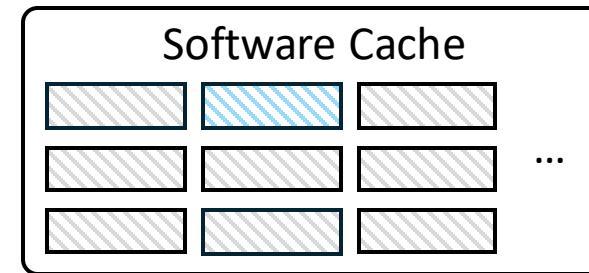


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6 →  wait(b)
7     compute(a, b)
8     ...
```

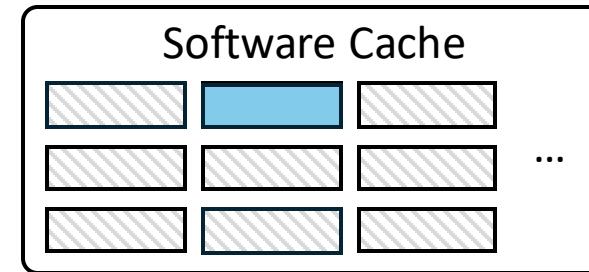


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6 →  wait(b)
7     compute(a, b)
8     ...
```

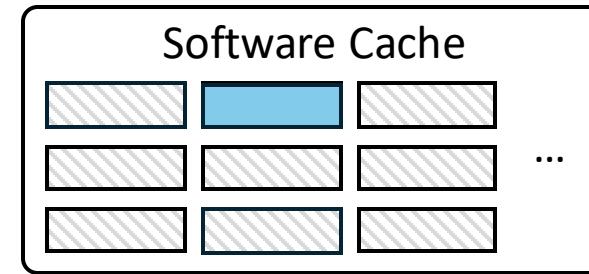


AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6     wait(b)
7     → compute(a, b)
8     ...
```



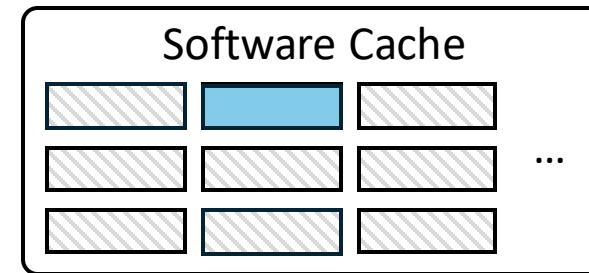
AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6     wait(b)
7     → compute(a, b)
8     ...
```

Unavailable again



AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

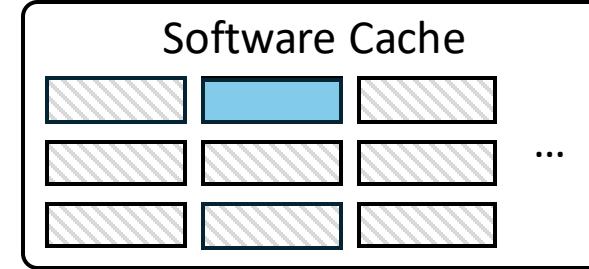
```

1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6     wait(b)
7     → compute(a, b)
8     ...

```

Unavailable again

- Use an extra synchronous read when accessing 'a'



AGILE Software Cache

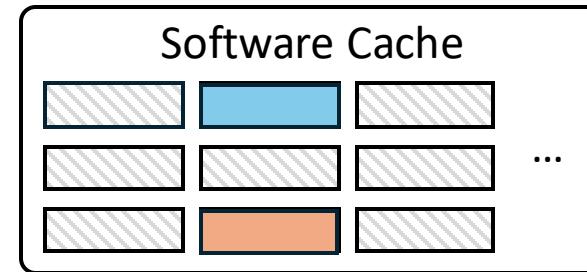
To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.

```
1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6     wait(b)
7     → compute(a, b)
8     ...
```

Unavailable again

- Use an extra synchronous read when accessing 'a'



AGILE Software Cache

To eliminate deadlock from the software cache:

- AGILE does not allow user threads to avoid cache line eviction.
- Avoid Deadlocks at the cost of additional I/Os.

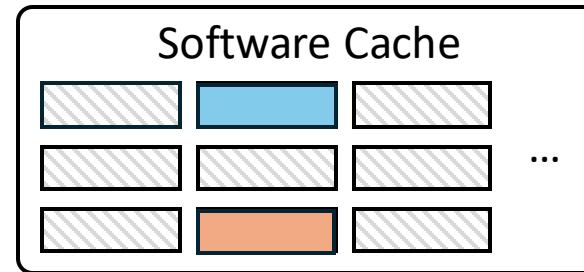
```

1 def kernel_async_access
2     a = read_async(...)
3     b = read_async(...)
4     ...
5     wait(a)
6     wait(b)
7     → compute(a, b)
8     ...

```

Unavailable again

- Use an extra synchronous read when accessing 'a'



Integrate User-managed Buffers into AGILE

By default, all accesses are proxied by AGILE software cache.

- Requested data may be unavailable again to eliminate deadlocks.

```
1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...
```

Access via AGILE software cache
(maybe require additional reads)

Integrate User-managed Buffers into AGILE

By default, all accesses are proxied by AGILE software cache.

- Requested data may be unavailable again to eliminate deadlocks.

➤ To address this problem, AGILE allows user-managed buffers.

```

1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...

```

Access via AGILE software cache
(maybe require additional reads)

```

1 def kernel_async_access( \
2     a_user_buf, b_user_buf):
3     read_async(a_user_buf)
4     read_async(b_user_buf)
5     ...
6     a_user_buf.wait()
7     b_user_buf.wait()
8     compute(a_user_buf, b_user_buf)
9     ...

```

Integrate User-managed Buffers into AGILE

By default, all accesses are proxied by AGILE software cache.

- Requested data may be unavailable again to eliminate deadlocks.

➤ To address this problem, AGILE allows user-managed buffers.

```

1 def kernel_async_access
2   a = read_async(..)
3   b = read_async(..)
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8   ...

```

Access via AGILE software cache
(maybe require additional reads)

User-managed buffers

```

1 def kernel_async_access( \
2   a_user_buf, b_user_buf):
3   read_async(a_user_buf)
4   read_async(b_user_buf)
5   ...
6   a_user_buf.wait()
7   b_user_buf.wait()
8   compute(a_user_buf, b_user_buf)
9   ...

```

Integrate User-managed Buffers into AGILE

By default, all accesses are proxied by AGILE software cache.

- Requested data may be unavailable again to eliminate deadlocks.

➤ To address this problem, AGILE allows user-managed buffers.

```

1 def kernel_async_access
2     a = read_async(..)
3     b = read_async(..)
4     ...
5     wait(a)
6     wait(b)
7     compute(a, b)
8     ...

```

Access via AGILE software cache
(maybe require additional reads)

User-managed buffers

```

1 def kernel_async_access( \
2     a_user_buf, b_user_buf):
3     read_async(a_user_buf)
4     read_async(b_user_buf)
5     ...
6     a_user_buf.wait()
7     b_user_buf.wait()
8     compute(a_user_buf, b_user_buf)
9     ...

```

Data will not be evicted once ready.

Integrate User-managed Buffers into AGILE

By default, all accesses are proxied by AGILE software cache.

- Requested data may be unavailable again to eliminate deadlocks.

➤ To address this problem, AGILE allows user-managed buffers.

```

1 def kernel_async_access
2   a = read_async(..)
3   b = read_async(..)
4   ...
5   wait(a)
6   wait(b)
7   compute(a, b)
8   ...

```

Access via AGILE software cache
(maybe require additional reads)

User-managed buffers

```

1 def kernel_async_access( \
2   a_user_buf, b_user_buf):
3   read_async(a_user_buf)
4   read_async(b_user_buf)
5   ...
6   a_user_buf.wait()
7   b_user_buf.wait()
8   compute(a_user_buf, b_user_buf)
9   ...

```

Data will not be evicted once ready.

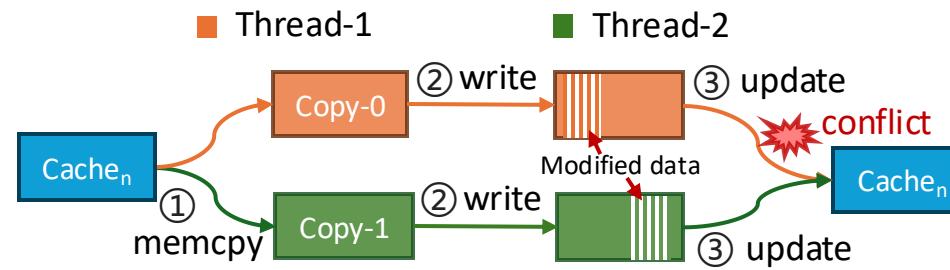
Data is guaranteed to be ready.

Integrate User-managed Buffers into AGILE

When multiple threads request the same data, data conflicts may occur.

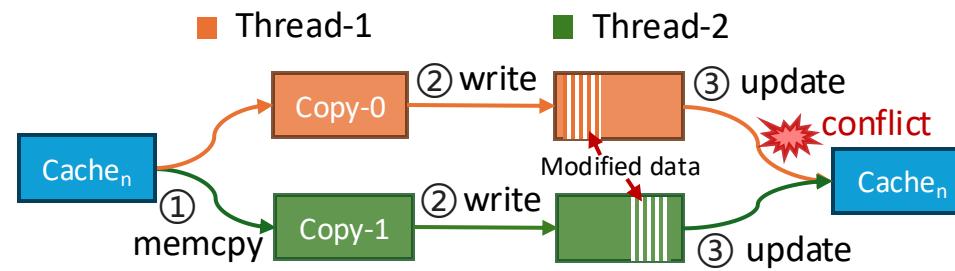
Integrate User-managed Buffers into AGILE

When multiple threads request the same data, data conflicts may occur.



Integrate User-managed Buffers into AGILE

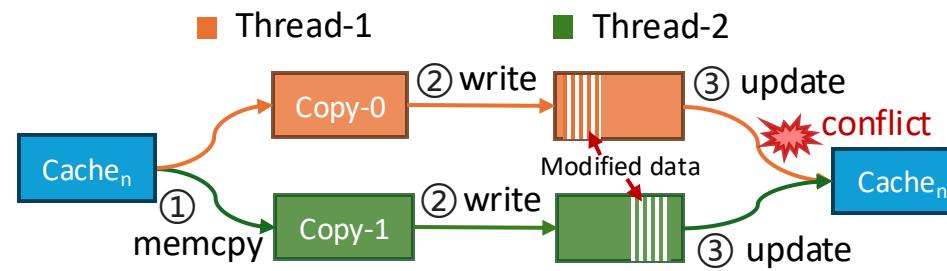
When multiple threads request the same data, data conflicts may occur.



AGILE provides a compile-time optional mechanism (share-table) to let user threads access the same user-managed buffer.

Integrate User-managed Buffers into AGILE

When multiple threads request the same data, data conflicts may occur.



AGILE provides a compile-time optional mechanism (share-table) to let user threads access the same user-managed buffer.

When the share-table is enabled at compile time:

- user-managed buffer (L1; managed by user; registered to share-table)
- Software cache (L2; managed by AGILE)
- SSDs (L3; managed by AGILE)

AGILE's API

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

Customizing software cache policy

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

Customizing software cache policy

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain; // Debug option for detecting deadlock  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

Customizing software cache policy

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

Debug option for detecting deadlock

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

Init AGILE &
Start AGILE
service

AGILE's API

Customizing software cache policy

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

Init AGILE &
Start AGILE
service

Start user kernel

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

AGILE's API

Customizing software cache policy

```
class GPUcache:public GPUcacheBase<GPUcache>{...};  
#define AGILE_CTRL AgileCtrl<GPUcache, ShareTable>  
__global__  
void kernel(AGILE_CTRL * ctrl, void * data){  
    ...  
    AgileLockChain chain;  
  
    // Method-1: AGILE prefetch  
    ctrl->prefetch(dev_idx, blk_idx, chain);  
  
    // Method-2: AGILE async_issue  
    AgileBufPtr buf(data + tid * ctrl->line_size);  
    ctrl->asyncRead(dev_idx, blk_idx, buf, chain);  
    buf.wait();  
    ctrl->asyncWrite(dev_idx, blk_idx, buf, chain);  
  
    // Method-3: AGILE array-like synchronous API  
    auto agileArr = ctrl->getArrayWrap<int>(chain);  
    int val = agileArr[dev_idx][idx];  
}
```

Init AGILE &
Start AGILE
service

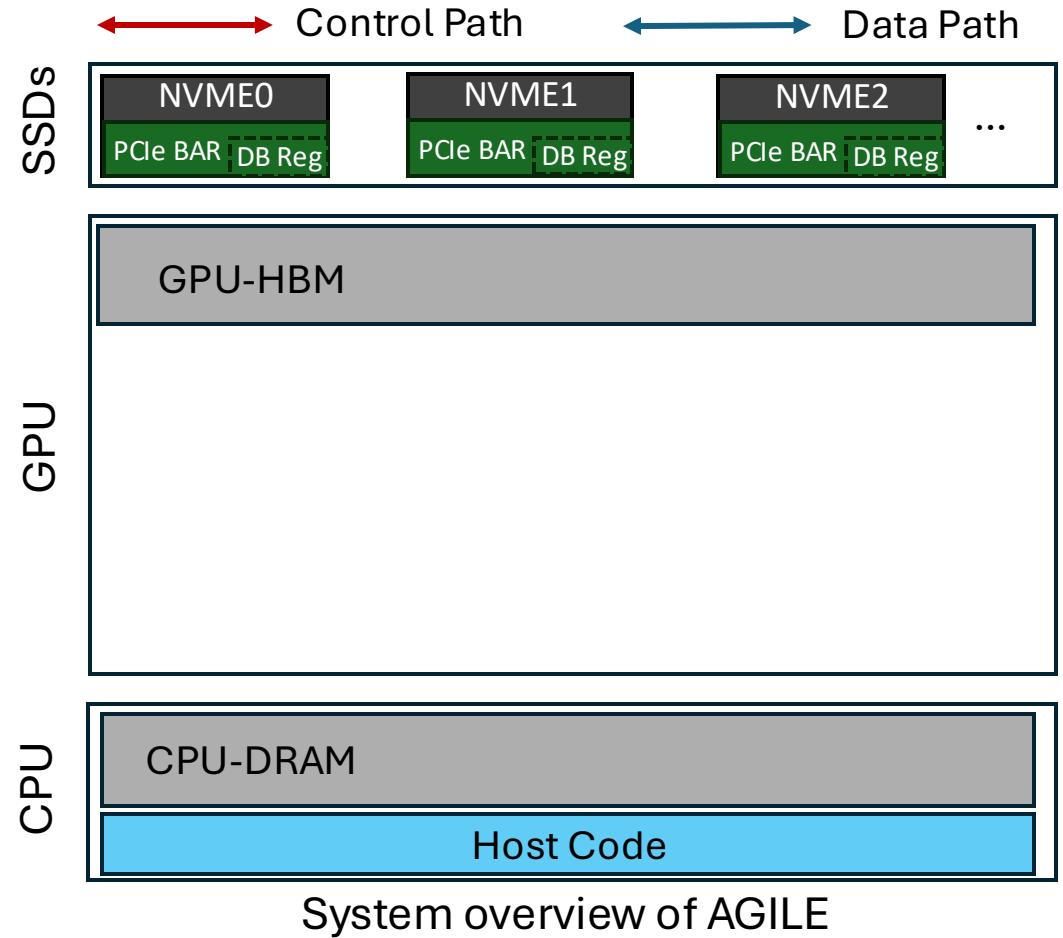
Start user kernel

Stop AGILE
service & close
SSDs

```
int main(int argc, char** argv){  
    // GPU Configurations  
    AGILE_HOST host(...);  
    // Policy Configurations  
    SHARE_TABLE_IMPL s_table(...);  
    GPU_CACHE_IMPL g_cache(...);  
    host.setGPUcache(g_cache);  
    host.setShareTable(s_table);  
    // Add and open target SSDs in the program  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.addNvmeDev("/dev/AGILE-xxx", ...);  
    host.initNvme();  
    // Initialize AGILE controller  
    host.initializeAgile(...);  
    // CUDA kernel parallelism configurations  
    host.configKernelParallelism(...);  
    host.queryOccupancy(kernel);  
    // Start the lightweight AGILE service  
    host.startAgile();  
    // Execute the CUDA kernel  
    host.runKernel(kernel, args...);  
    // Stop AGILE service  
    host.stopAgile();  
    // Close the opened SSDs  
    host.closeNvme();  
}
```

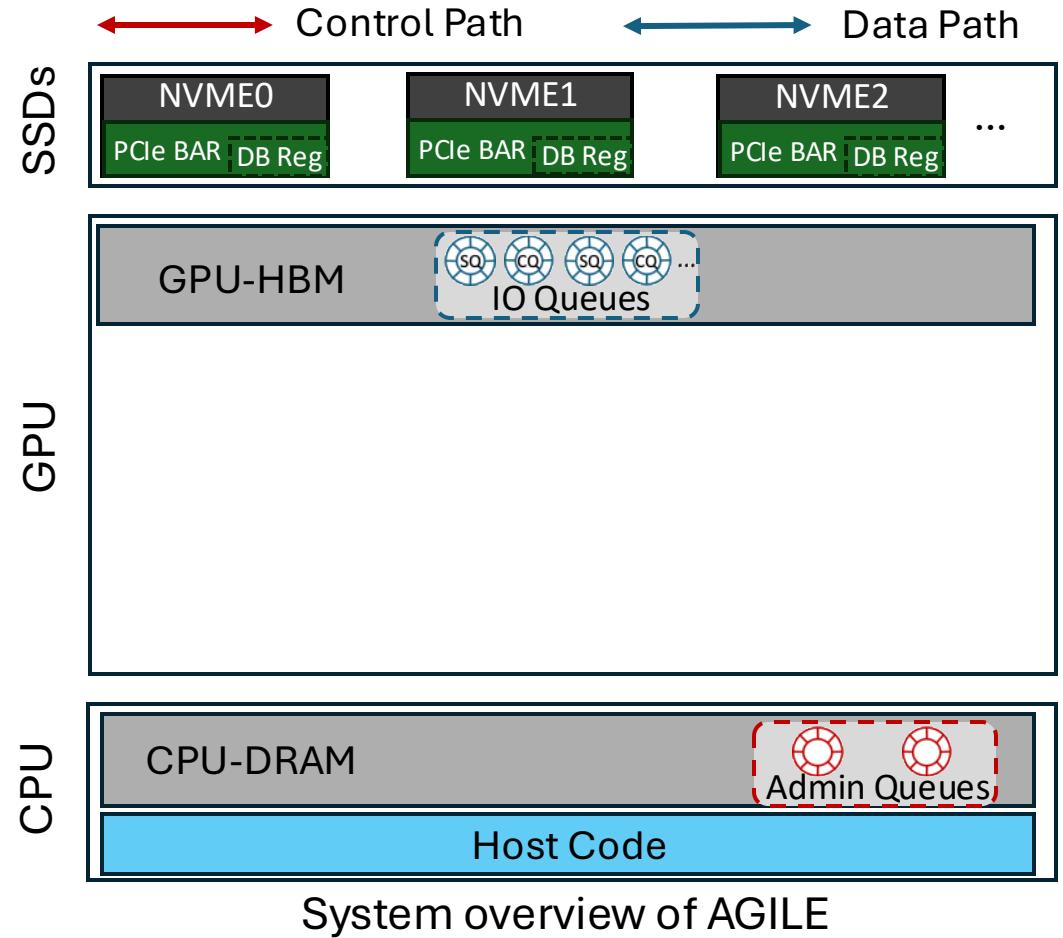
Summary of AGILE

Feature Highlights:



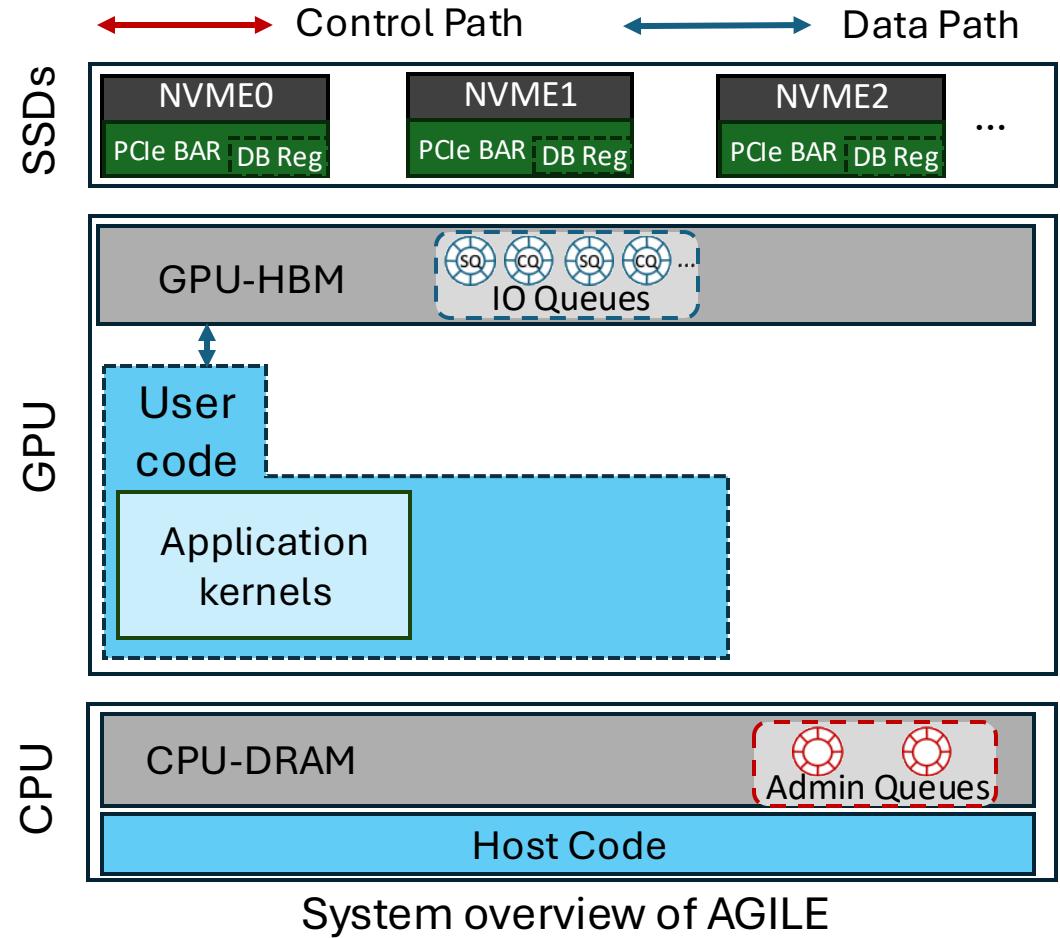
Summary of AGILE

Feature Highlights:



Summary of AGILE

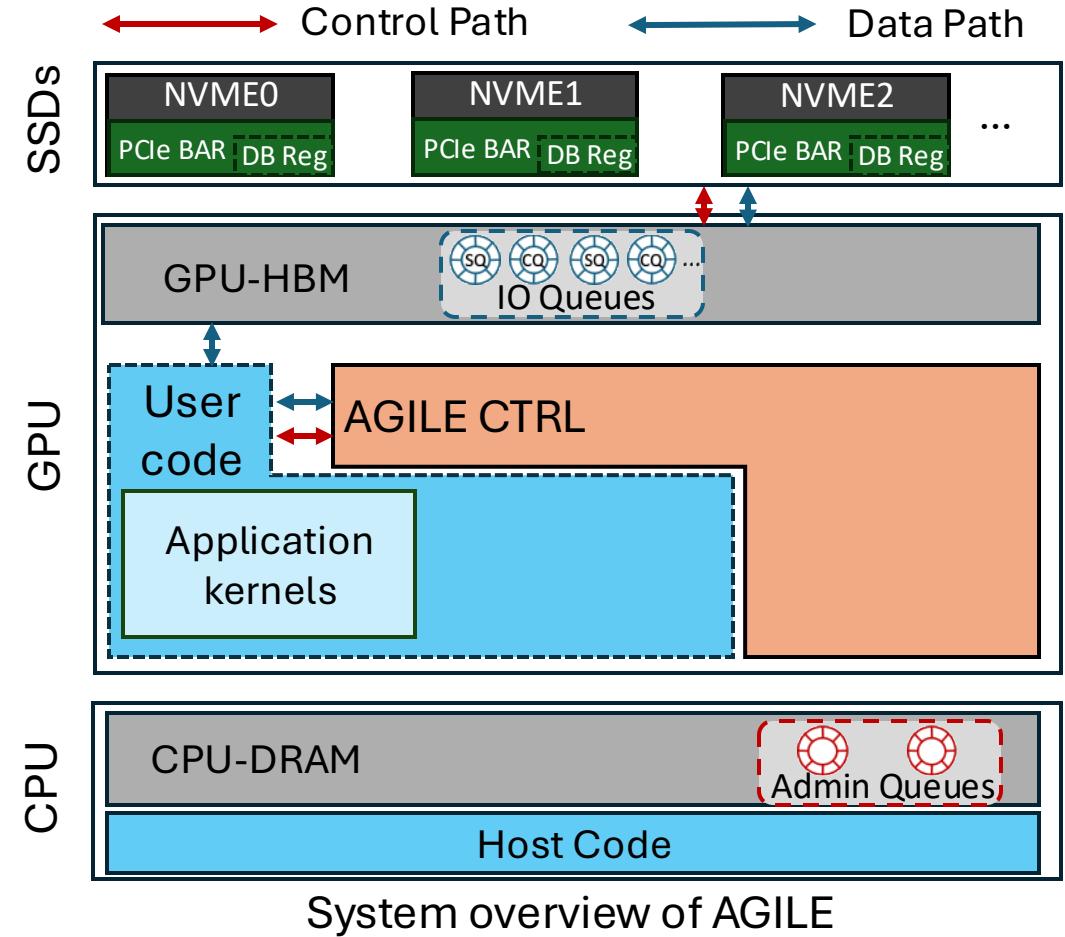
Feature Highlights:



Summary of AGILE

Feature Highlights:

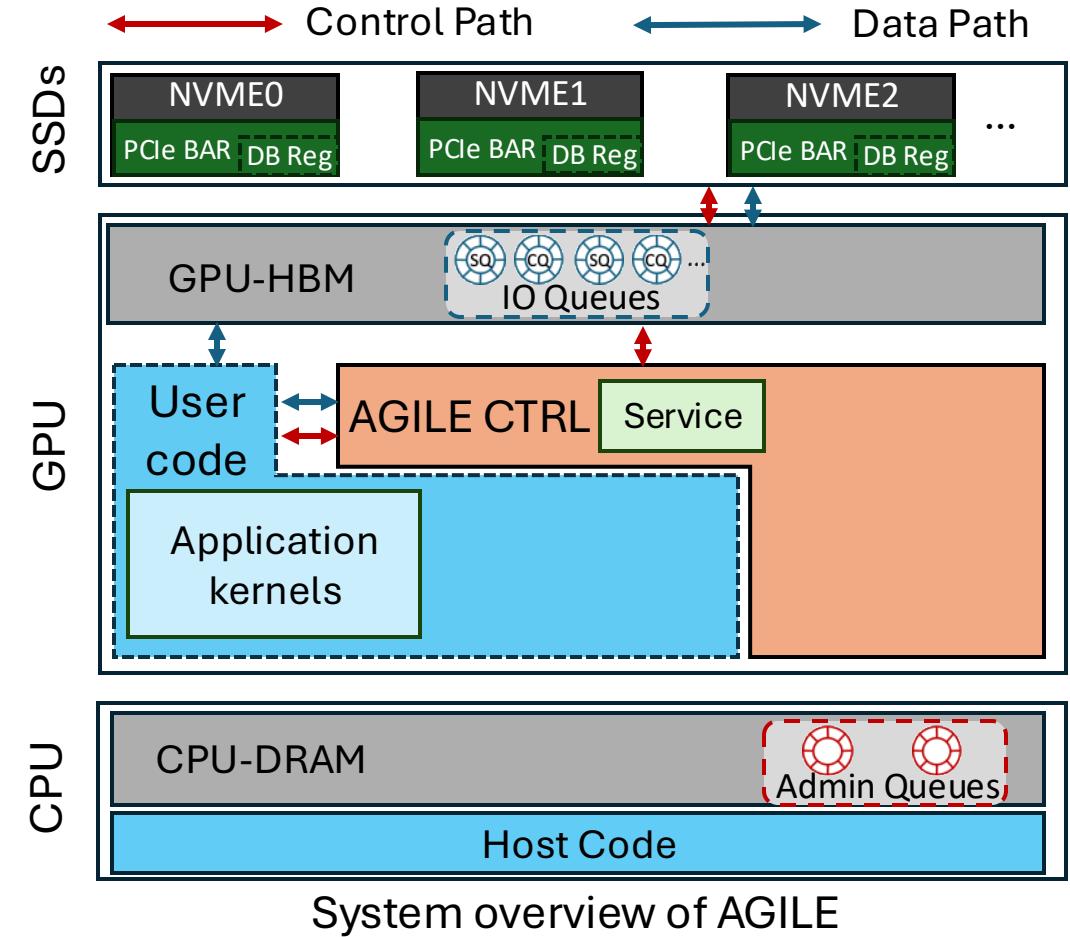
1. AGILE is the first async GPU-centric I/O model, allowing GPU threads to access SSDs asynchronously.



Summary of AGILE

Feature Highlights:

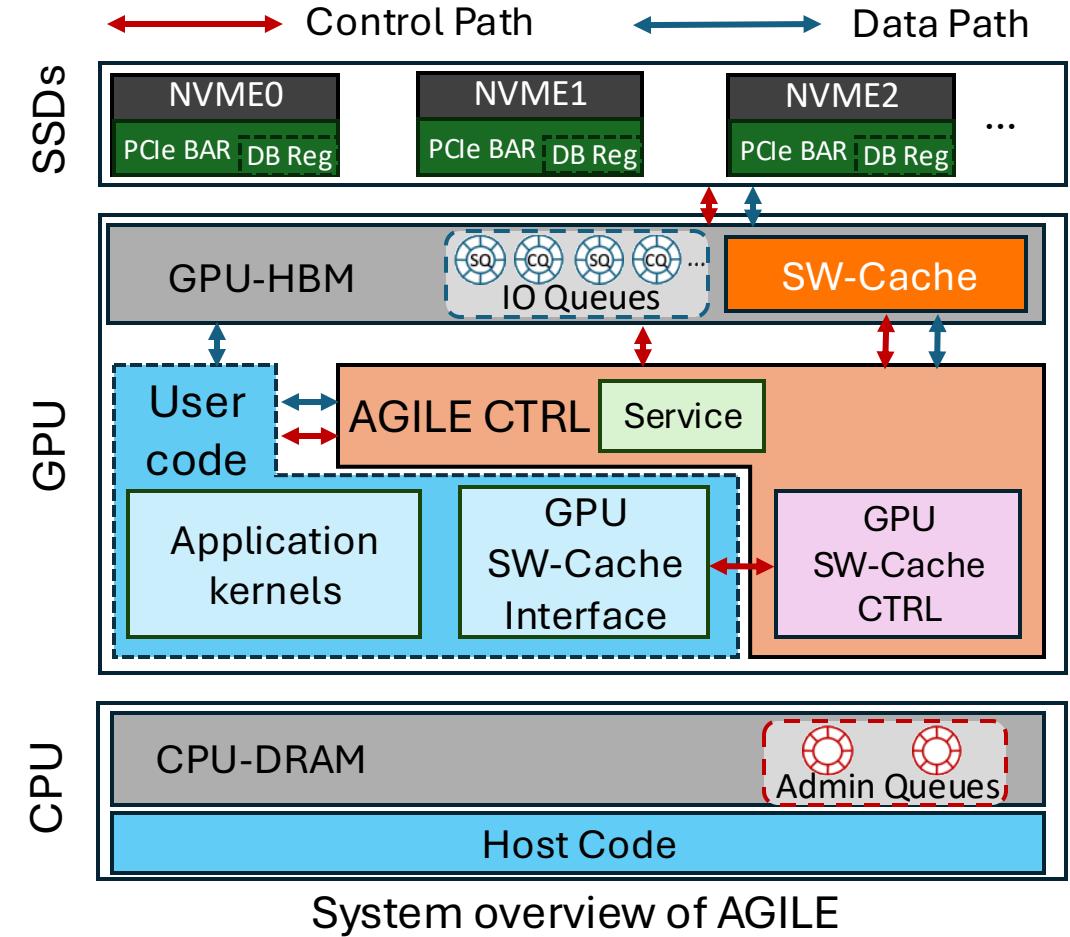
1. AGILE is the first async GPU-centric I/O model, allowing GPU threads to access SSDs asynchronously.
2. AGILE eliminates the deadlock risks in asynchronous NVMe I/O via AGILE polling service.



Summary of AGILE

Feature Highlights:

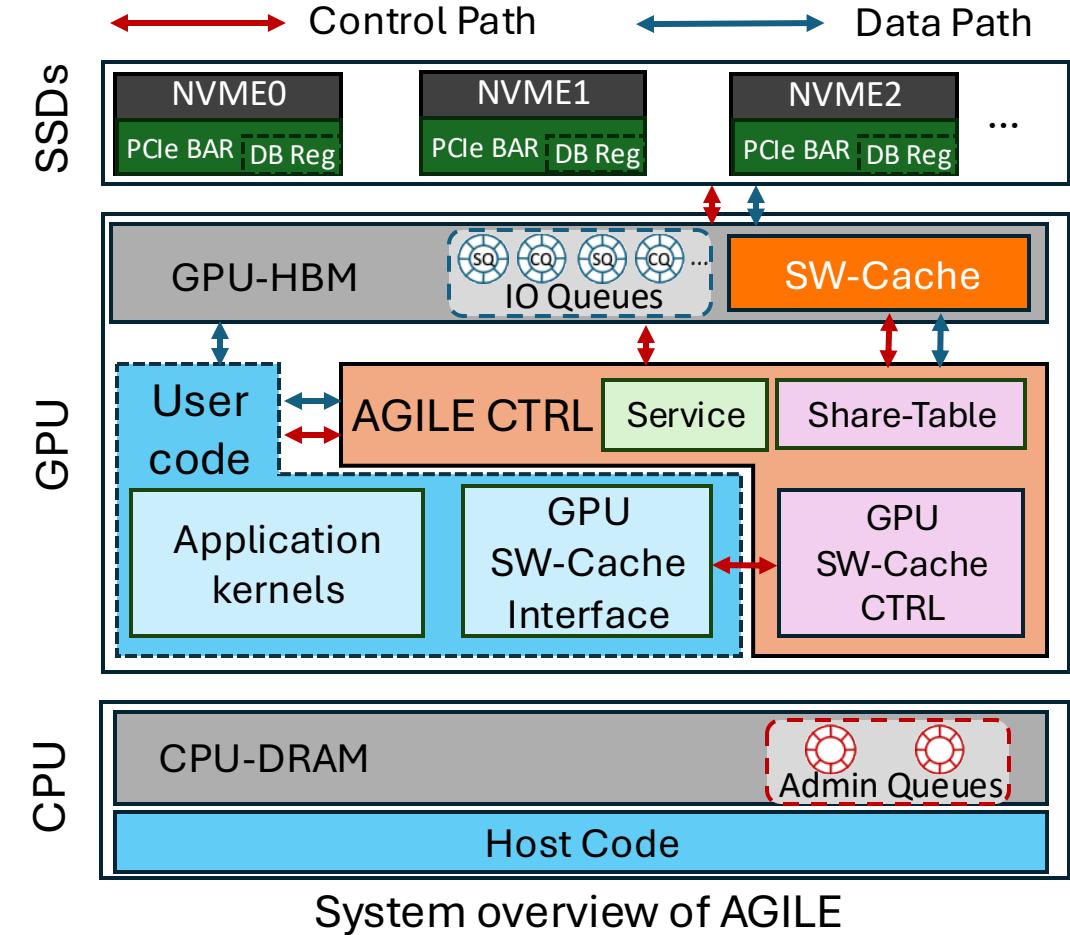
1. AGILE is the first async GPU-centric I/O model, allowing GPU threads to access SSDs asynchronously.
2. AGILE eliminates the deadlock risks in asynchronous NVMe I/O via AGILE polling service.
3. AGILE allows users to customize software cache policies.



Summary of AGILE

Feature Highlights:

1. AGILE is the first async GPU-centric I/O model, allowing GPU threads to access SSDs asynchronously.
2. AGILE eliminates the deadlock risks in asynchronous NVMe I/O via AGILE polling service.
3. AGILE allows users to customize software cache policies.
4. AGILE allows user-maintained buffers to be integrated into the software cache hierarchy.



Experimental Setup

- Dell R750 Server
 - Ubuntu 20.04 (Linux 5.4.0-200-generic)
- Nvidia RTX 5000 Ada GPU (PCIe Gen4 x16)
- NVMe SSDs
 - Dell Ent NVMe AGN MU AIC 1.6TB SSD (PCIe Gen4 x4)
 - Two Samsung 990 PRO 1TB SSDs (PCIe Gen4 x4)
- Software Cache Policy
 - Clock replacement algorithm^[1] (keep the same with BaM)

[1] F. J. Corbato. 1968. A Paging Experiment With The Multics System. Technical Report, Massachusetts Institute of Technology, Cambridge, Project MAC (1968).

Experimental Evaluation

1. Compute and communication overlap.

- 1 thread block issues 64 NVMe commands, and uses fetched data for compute.

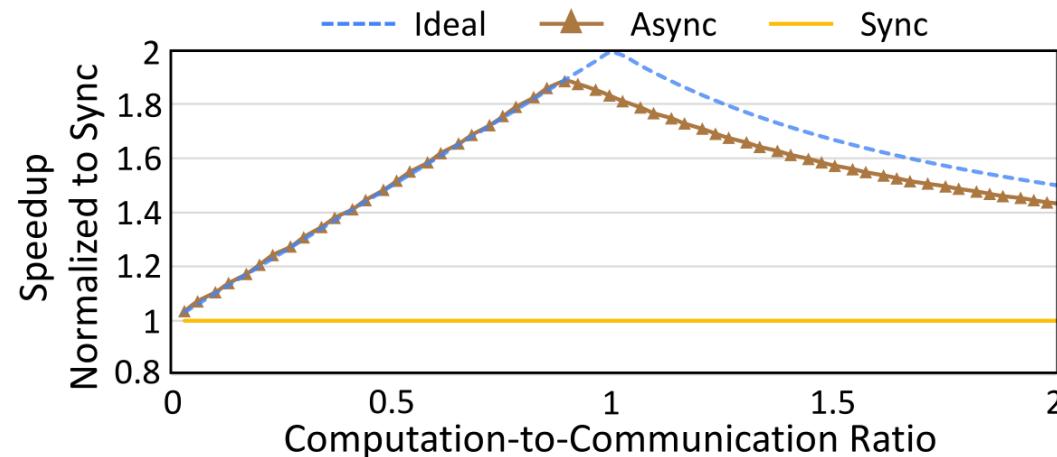


Figure 4: Speedup comparison of asynchronous I/O over synchronous I/O on workloads with different Computation-to-Communication Ratio (CTC).

Experimental Evaluation

1. Compute and communication overlap.

- 1 thread block issues 64 NVMe commands, and uses fetched data for compute.

➤ AGILE can hide slow communication with computation effectively.

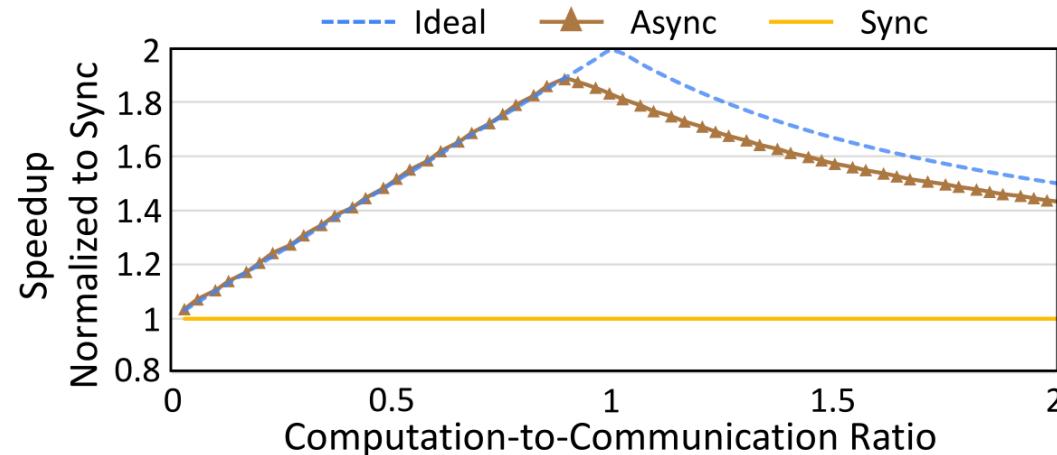


Figure 4: Speedup comparison of asynchronous I/O over synchronous I/O on workloads with different Computation-to-Communication Ratio (CTC).

Experimental Evaluation

2. Scalability in 4KB random read/write

- Access different SSDs in an interleaving fashion.

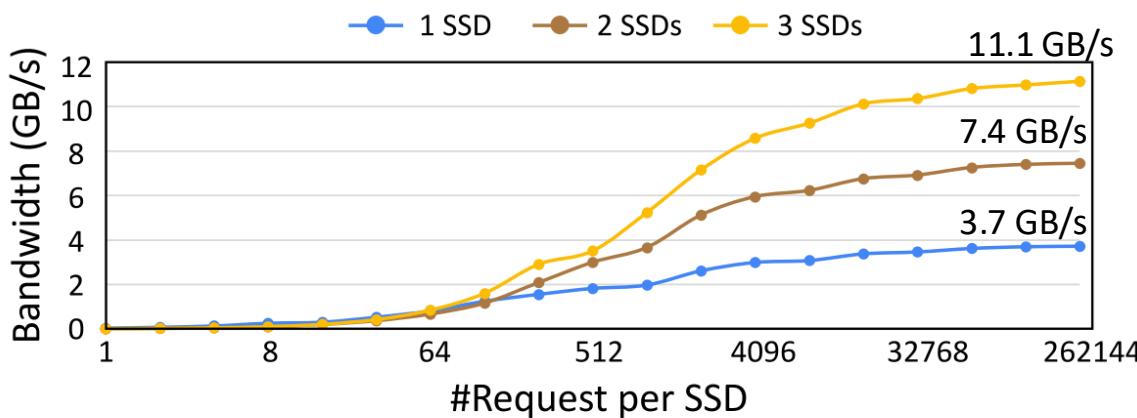


Figure 5: AGILE 4KB random read on multiple SSDs

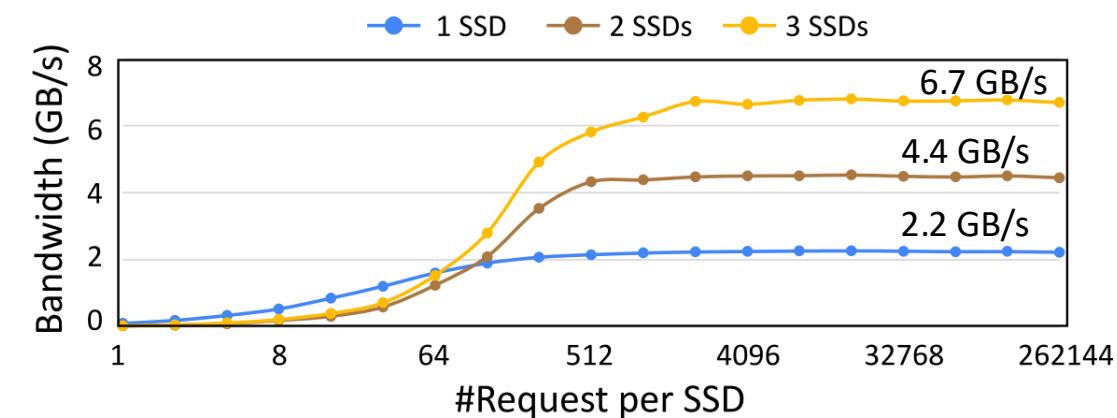


Figure 6: AGILE 4KB random write on multiple SSDs

Experimental Evaluation

2. Scalability in 4KB random read/write

- Access different SSDs in an interleaving fashion.

➤ AGILE shows good scalability.

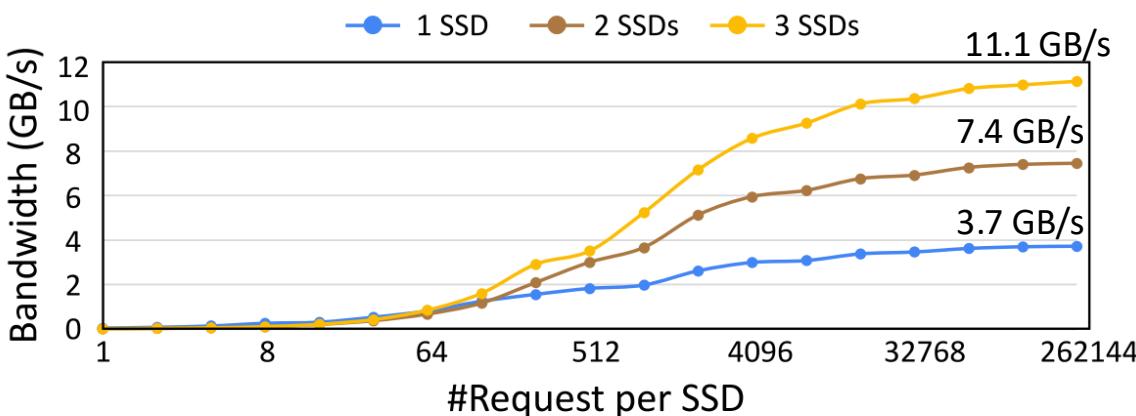


Figure 5: AGILE 4KB random read on multiple SSDs

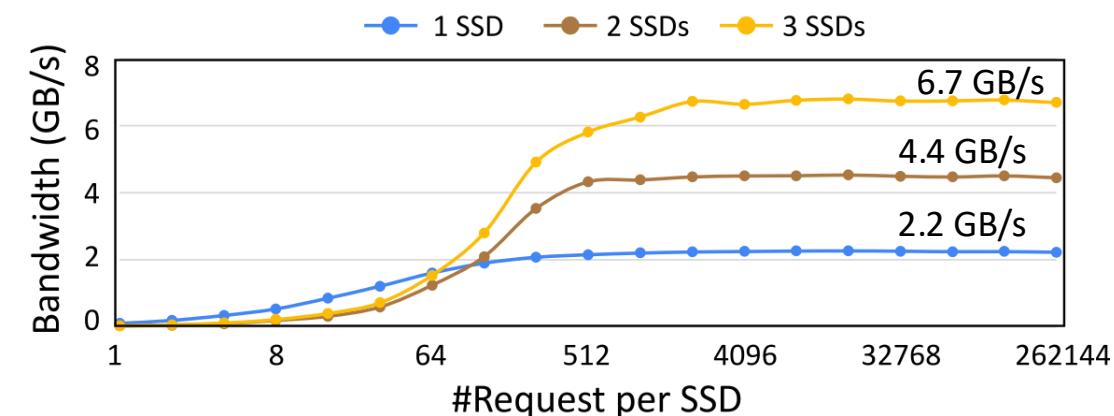


Figure 6: AGILE 4KB random write on multiple SSDs

Experimental Evaluation

3. Comparison with BaM on DLRM inference MicroBenchmark.

- The DLRM model is adopted from [1] with Criteo 1TB Click Logs dataset[2].
- cuBLAS is used for all matrix multiplications.
- BaM and AGILE are used for fetching data.
- AGILE is used in both synchronous mode and asynchronous mode

Experimental Evaluation

3. Comparison with BaM on DLRM inference MicroBenchmark.

- The DLRM model is adopted from [1] with Criteo 1TB Click Logs dataset[2].
 - cuBLAS is used for all matrix multiplications.
 - BaM and AGILE are used for fetching data.
 - AGILE is used in both synchronous mode and asynchronous mode
- Sweep key parameter: Batch Size & the Number of I/O Queues.
- Default parameters:
- Batch Size: 2048
 - #I/O queues: 128
 - Software cache size: 2 GB

Experimental Evaluation

3. Comparison with BaM on DLRM inference MicroBenchmark.

- The DLRM model is adopted from [1] with Criteo 1TB Click Logs dataset[2].
 - cuBLAS is used for all matrix multiplications.
 - BaM and AGILE are used for fetching data.
 - AGILE is used in both synchronous mode and asynchronous mode
- Sweep key parameter: Batch Size & the Number of I/O Queues.
- Default parameters:
- Batch Size: 2048
 - #I/O queues: 128
 - Software cache size: 2 GB

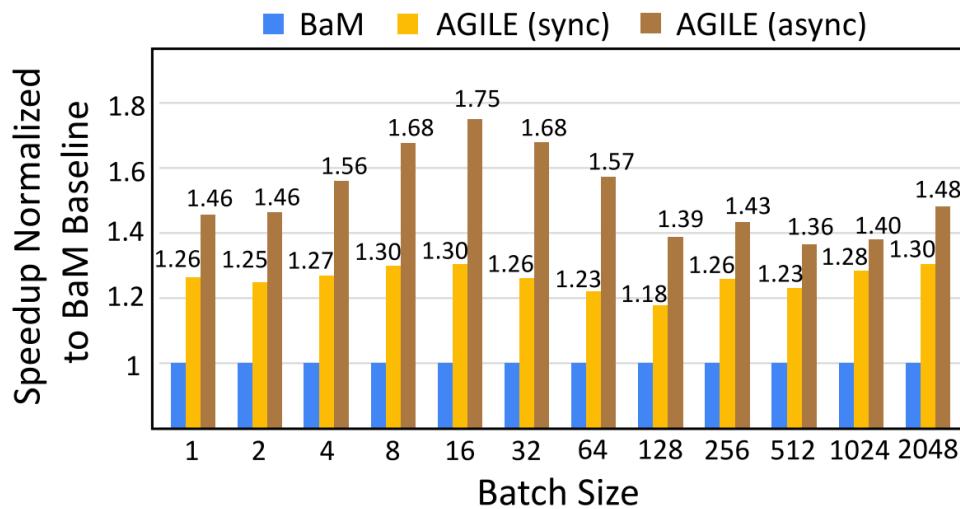


Figure 8: Speedup comparison of AGILE (async and sync modes) and BaM across varying batch sizes in DLRM inference.

Experimental Evaluation

3. Comparison with BaM on DLRM inference MicroBenchmark.

- The DLRM model is adopted from [1] with Criteo 1TB Click Logs dataset[2].
 - cuBLAS is used for all matrix multiplications.
 - BaM and AGILE are used for fetching data.
 - AGILE is used in both synchronous mode and asynchronous mode
- Sweep key parameter: Batch Size & the Number of I/O Queues.

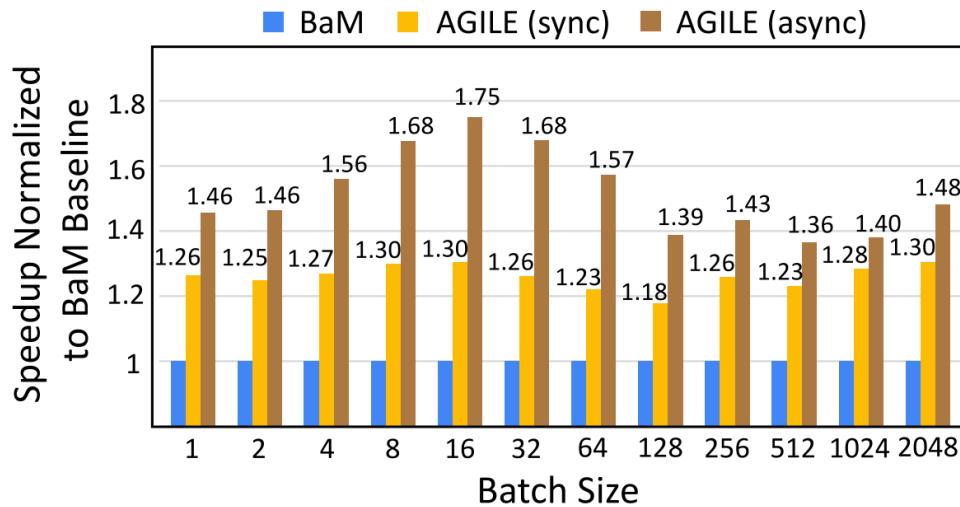


Figure 8: Speedup comparison of AGILE (async and sync modes) and BaM across varying batch sizes in DLRM inference.

[1] Naumov, Maxim, et al. "Deep learning recommendation model for personalization and recommendation systems." *arXiv preprint arXiv:1906.00091* (2019).

[2] Criteo AI Lab. 2025. Download Criteo 1TB Click Logs dataset - Criteo AI Lab. <https://ailab.criteo.com/download-criteo-1tb-click-logs-dataset/>

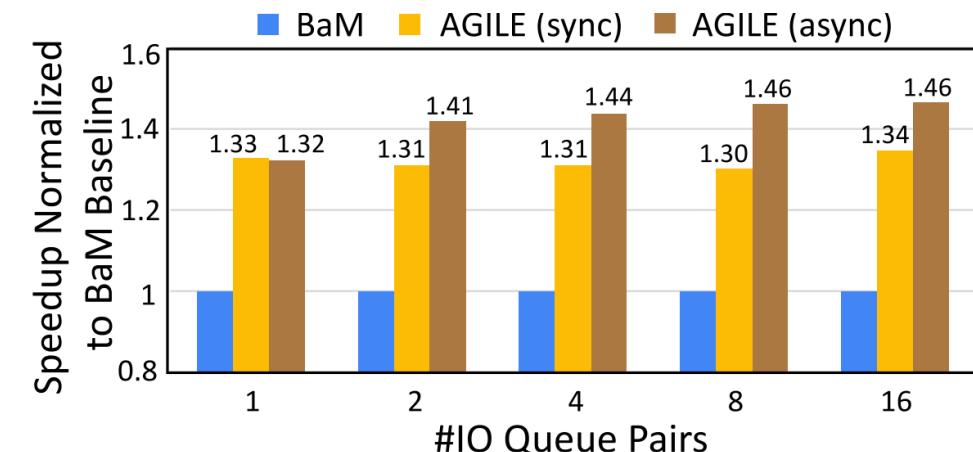


Figure 9: Speedup comparison of AGILE (async and sync modes) and BaM under varying numbers of I/O queue pairs in DLRM inference.

Experimental Evaluation

3. Comparison with BaM on DLRM inference MicroBenchmark.

- The DLRM model is adopted from [1] with Criteo 1TB Click Logs dataset[2].
 - cuBLAS is used for all matrix multiplications.
 - BaM and AGILE are used for fetching data.
 - AGILE is used in both synchronous mode and asynchronous mode
- Sweep key parameter: Software Cache Size

Default parameters:

- Batch Size: 2048
- #I/O queues: 128
- Software cache size: 2 GB

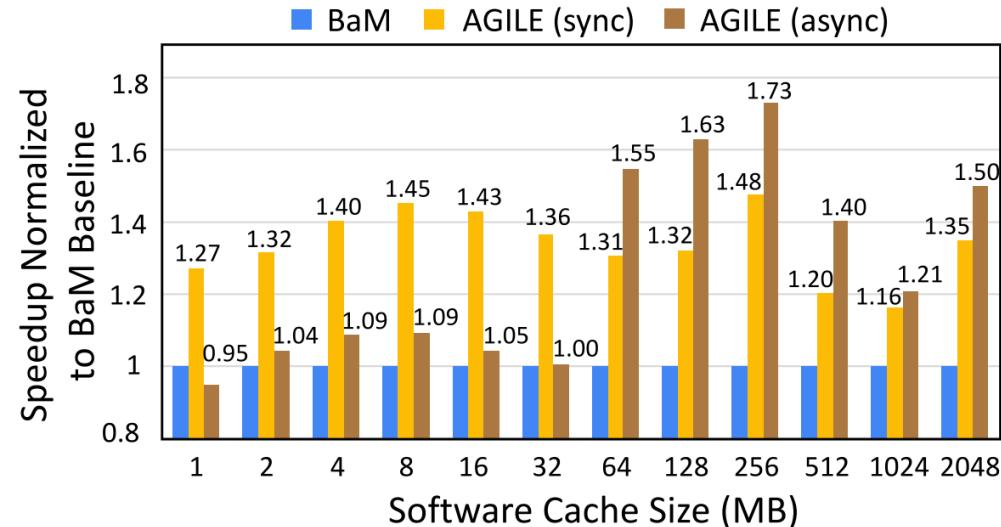


Figure 10: Speedup comparison of AGILE (async and sync modes) and BaM under varying software cache sizes in DLRM inference.

THANK YOU!

QUESTIONS?

zhuoping_yang@brown.edu
peipei_zhou@brown.edu

<https://peipeizhou-eecs.github.io/>

<https://github.com/arc-research-lab/AGILE>

AGILE is open source!



BROWN
UNIVERSITY

